

Validated Synthetic Patient Generation for Small Longitudinal Cohorts: Coagulation Dynamics Across Pregnancy

Jeffrey D. Varner^{1,*}, Maria Cristina Bravo², Carole McBride³,
Thomas Orfeo², Ira Bernstein³

¹Robert F. Smith School of Chemical and Biomolecular Engineering,
Cornell University, Ithaca, NY 14850

²Department of Biochemistry, The Robert Larner, M.D.
College of Medicine, University of Vermont Medical Center, Burlington, VT 05401

³Department of Obstetrics, Gynecology, and Reproductive Sciences,
The Robert Larner, M.D. College of Medicine,
University of Vermont Medical Center, Burlington, VT 05401

*Corresponding author: jdv27@cornell.edu

Abstract

Small longitudinal clinical cohorts, common in maternal health, rare diseases, and early-phase trials, limit computational modeling: too few patients to train reliable models, yet too costly and slow to expand through additional enrollment. We present multiplicity-weighted Stochastic Attention (SA), a generative framework based on modern Hopfield network theory that addresses this gap. SA embeds real patient profiles as memory patterns in a continuous energy landscape [whose induced distribution places one component on each stored patient, and generates novel synthetic patients by sampling it while preserving the geometry of the original cohort. We generated the reported cohort with Langevin dynamics. The same distribution is also available in closed form, and an exact analytic sampler reproduced the same fidelity and privacy results.](#) Per-pattern multiplicity weights enable targeted amplification of rare clinical subgroups at inference time without retraining. We applied SA to a longitudinal coagulation dataset from 23 pregnant patients spanning 72 biochemical features across 3 visits (pre-pregnancy baseline, first trimester, and third trimester), including rare subgroups such as polycystic ovary syndrome and preeclampsia. Across multiple validation tests, including an ordinary differential equation model of the coagulation cascade, [separately specified, blind to the generator, and calibrated on the same real cohort](#), synthetic patients generated by SA [closely matched their real counterparts under the statistical, structural, and mechanistic measures applied at this sample size.](#) A downstream utility test further showed that a mechanistic model calibrated entirely on synthetic patients predicted held-out real patient outcomes as well as one calibrated on real data. These results [demonstrate, as a proof-of-concept, that SA can produce synthetic cohorts useful for the modeling use-cases tested here \(mechanistic calibration and hypothesis generation\)](#) from very small longitudinal datasets, enabling data-augmented modeling in small-cohort settings.

Keywords: synthetic data generation, stochastic attention, modern Hopfield networks, multiplicity weighting, coagulation, pregnancy, longitudinal data, mechanistic validation

1 Introduction

In 2021, the United States maternal mortality rate reached 32.9 deaths per 100,000 live births, the highest level since 1965^{1,2}. Between 2018 and 2022, Non-Hispanic Black women died at 2.8 times the rate of White women, and disorders related to pregnancy, including hemorrhage and venous complications, were the leading underlying cause of death, accounting for 17.4% of 6,283 pregnancy-related deaths³. These figures highlight the role of the coagulation system in maternal outcomes. Blood coagulation is a complex enzymatic cascade in which a small tissue factor (TF) stimulus triggers a series of serine protease activation reactions that ultimately produce thrombin, the central effector of hemostasis^{4,5}. Thrombin cleaves fibrinogen into fibrin, activates platelets, and amplifies its own production through positive feedback via Factors V, VIII, and XI, while simultaneously initiating its own shutdown through the thrombomodulin–protein C anticoagulant pathway^{6,7}. Pregnancy alters this hemostatic balance, shifting toward a hypercoagulable state that prepares for delivery^{8–10}. Procoagulant factors including fibrinogen, Factor VIII, and von Willebrand factor increase substantially across gestation, while natural anticoagulants such as antithrombin and protein S decline. Two conditions further shift this balance in ways that are not fully understood. Polycystic ovary syndrome (PCOS), a hormonal and metabolic disorder affecting approximately 6.6% of reproductive-age women¹¹, has been associated with impaired fibrinolysis and a prothrombotic tendency¹². Preeclampsia (PE), a hypertensive disorder that complicates 2–8% of pregnancies worldwide¹³, is characterized by endothelial dysfunction, platelet activation, and altered thrombin generation¹⁴. Mathematical models of the coagulation cascade have provided mechanistic insight into patient-specific thrombin generation^{15,16}. These range from the original Hockin–Mann stoichiometric framework⁶ through extensions incorporating the thrombomodulin–protein C pathway⁷ to the Brummel-Ziedins 2012 (BZ2012) model with detailed activated protein C (APC)-mediated Factor Va inactivation kinetics¹⁷. However, applying these models to pregnancy cohorts requires longitudinal datasets capturing coagulation factors, inhibitors, and thrombin generation assay (TGA) measurements across gestation, data that is rare and expensive to collect.

The core barrier is sample size. Longitudinal pregnancy coagulation studies typically enroll tens of patients with complete follow-up across multiple trimesters, yielding datasets where the number of features p (coagulation factors, TGA parameters, viscoelastic measurements) exceeds the number of patients n . Rare complications compound this problem. Small longitudinal cohorts inevitably contain only a handful of patients with any given condition, far too few for condition-specific statistical analysis. This $n < p$ regime limits conventional approaches because covariance matrices are rank-deficient, cross-validation is unreliable, and overfitting is nearly inevitable¹⁸. Synthetic data generation addresses this by augmenting small real datasets with statistically faithful synthetic records, enabling downstream analyses (mechanistic modeling, clinical comparisons, hypothesis generation for rare complications) that would otherwise be infeasible^{19,20}. **An original motivation for this work was to support deep learning models in both directions: predicting TGA responses from coagulation factors and estimating factor profiles compatible with an observed TGA response.** However, existing methods face limitations in the small-cohort longitudinal setting. Sampling from a fitted multivariate normal (MVN) distribution is singular when $n < p$ and must be regularized²¹, introducing bias that distorts the joint distribution. MVN also assumes Gaussian marginals and linear correlations, and provides no mechanism for conditional generation: it cannot amplify a specific subpopulation while preserving its signatures. More expressive models such as generative adversarial networks (GANs) and variational autoencoders (VAEs)^{22–24}, including recent longitudinal extensions²⁰, require substantially larger training sets and are prone to mode collapse at small n ²⁵ (we confirm this empirically for Conditional Tabular GAN (CTGAN) at $n=23$ in this work). A generative method is needed that can operate directly on the geometry of a small

dataset without estimating a full covariance model.

Stochastic Attention (SA) is such a framework. Rather than fitting a full covariance model or training a neural generator, SA treats the stored patient profiles themselves as the model. Building on modern Hopfield network theory²⁶, each profile becomes a memory pattern in a continuous energy landscape. This energy induces a continuous probability distribution that is exactly a finite Gaussian mixture in the reduced space, with one nonzero-variance component centered on each stored pattern. Sampling from this distribution therefore produces continuous variations around the stored patterns rather than exact copies. Because sampling occurs in a principal component analysis (PCA)-reduced linear subspace anchored on the observed cohort, SA never forms a full $p \times p$ covariance matrix and so avoids both the rank deficiency that limits MVN and the mode collapse that limits GAN/VAE approaches when $n < p$. We recently introduced a multiplicity-weighted extension of the Hopfield energy that attaches a per-pattern weight to each stored profile²⁷, which provides exactly the lever that MVN lacks: continuous interpolation between unconditioned generation and targeted amplification of a rare subpopulation at inference time, without retraining. What remains open, and what this study addresses, is whether SA can faithfully generate *longitudinal* trajectories from a small cohort. In earlier proceedings work, we paired our mechanistic coagulation pipeline with per-visit MVN sampling to produce synthetic populations²⁸; that approach captured single-visit marginals but, by fitting each visit independently, generated patients whose trajectories across gestation were statistically incoherent. A generator that respects the joint structure across visits is what a small longitudinal pregnancy cohort actually demands.

In this study, we used multiplicity-weighted SA to generate complete longitudinal patient profiles from a small pregnancy coagulation cohort ($K=23$ patients, 72 features per visit, 3 visits) that includes rare clinical subgroups (PCOS, $n=3$; Developed PE, $n=5$). Our pipeline concatenates each patient’s multi-visit profile into a single vector, applies PCA for dimensionality reduction, and generates new patients via multiplicity-weighted Langevin dynamics on the Hopfield energy landscape, with a direction-magnitude decomposition that preserves the natural variance structure of continuous clinical measurements. We then asked whether the synthetic patients reproduced marginal feature distributions, the cross-visit covariance structure, condition-specific signatures of rare clinical subgroups, and the biological consistency obtained when the profiles were processed by a separately specified, generator-blind ordinary differential equation (ODE) model of the coagulation cascade that was calibrated on the same real cohort¹⁷. As a downstream-utility test, we further asked whether a mechanistic model calibrated entirely on synthetic patients could predict held-out real patient outcomes. As a proof-of-concept, the results show that SA recovers the statistical structure of the real cohort and matches its mechanistic plausibility at this sample size. If this generalizes, the implication is that the bottleneck for studying rare obstetric and pediatric conditions may be shifting from cohort size toward cohort fidelity: a few dozen carefully phenotyped longitudinal patients, augmented through SA, may be enough to support mechanistic and statistical analyses that have traditionally relied on much larger cohorts.

2 Results

We generated $N=100$ synthetic longitudinal patient profiles using SA and compared them against the $K=23$ real patients across four validation levels. This sample size provided approximately $4\times$ amplification while remaining in the regime where SA produces diverse, non-degenerate samples. As a baseline, we also generated $N=100$ patients from a regularized multivariate normal (MVN) distribution fitted to the same data.

2.1 Marginal Plausibility

We first assessed whether SA-generated patients reproduced per-feature summary statistics across all three visits. The pooled median mean relative error (MRE) across all 216 feature–visit entries was 1.2% (95% bootstrap CI: 1.0–1.6%), with 193 of 216 entries (89%) below 5% (Table S5), indicating that SA captured the central tendencies of the real population. Per-visit MREs for representative coagulation features are summarized in Table 2, where Factor II and Factor VIII fell below 2%, antithrombin below 3%, and fibrinogen below 1% across all three visits. We verified that synthetic patients were not memorized copies of real patients using two complementary metrics summarized in Table S9. The mean novelty score ($1 - \max_k \cos(\hat{\xi}, \hat{\mathbf{m}}_k)$) was 0.50 at $\beta = \beta^*$, indicating that generated samples were on average 50% away from their nearest stored pattern in angular distance. The median nearest-neighbor distance from each synthetic patient to its closest real patient was 6.14 standardized units, compared with 6.87 for the median real-to-real nearest-neighbor distance, a ratio of 0.89 indicating that synthetic patients were approximately as far from real patients as real patients were from each other.

We next examined whether known physiological relationships and longitudinal trajectories were preserved in the synthetic cohort. Pairwise scatter plots of coagulation factor levels against thrombin generation parameters showed that SA-generated patients occupied the same joint regions as real patients (Fig. 1). The expected inverse relationship between antithrombin and TF-initiated peak thrombin was preserved, as were the positive associations between Factor VIII and peak, prothrombin and endogenous thrombin potential (ETP), and the inverse relationship between antithrombin and ETP. Per-visit means and standard deviations for six key coagulation features confirmed that synthetic patients tracked the real longitudinal trajectories closely, with overlapping confidence bands at all three visits (Fig. 2). Fibrinogen, Factor VIII, and von Willebrand factor showed the expected monotonic increases from baseline through the third trimester in both cohorts, while Factor II increased modestly and antithrombin declined then partially recovered, consistent with the known hypercoagulable shift of pregnancy. Together, these results indicated that SA captured not only univariate distributions and pairwise dependencies but also the directionality and magnitude of longitudinal change. The 72 features also included viscoelastic (ROTEM) and fibrinolytic parameters, which were generated as part of the concatenated vector and showed comparable MREs to the coagulation factors (Table S5), but are not individually highlighted here because the mechanistic validation uses the BZ2012 thrombin generation model, which does not cover fibrinolysis or clot viscoelasticity. *These variables contributed through their loadings on the retained principal components to the memory patterns that defined the energy landscape. Their fidelity was evaluated statistically here, whereas their mechanistic fidelity was not evaluated.*

We also compared SA against two widely used deep generative methods for tabular data, Conditional Tabular GAN (CTGAN) and Tabular Variational Autoencoder (TVAE)²⁴ (Table S6). CTGAN failed completely, producing median MREs of $\approx 19\%$, an order of magnitude worse than SA, across all epoch counts tested (300–3,000), consistent with known GAN mode collapse on small datasets. TVAE achieved comparable marginal fidelity at high epoch counts (median MRE 1.8% at 3,000 epochs), but operated on per-visit rows ($n=90$) rather than concatenated longitudinal profiles ($n=23$), meaning it could not capture cross-visit covariance structure and had no mechanism for conditional subgroup generation.

To distinguish the contribution of PCA from that of the stochastic-attention generation process, we compared stochastic attention with four generators operating in the identical $d=18$ PCA subspace: a two-component diagonal-covariance Gaussian mixture, a kernel-density estimator, k -nearest-neighbor interpolation, and a Gaussian copula (Supplementary Table S13). The comparison indicated that the shared representation accounted for much of the observed fidelity but did not

guarantee performance. The nearest-neighbor samples remained close to stored profiles: their median novelty, defined as $1 - \max_k \cos(\hat{\xi}, \hat{\mathbf{m}}_k)$, was 0.05, and only 5% exceeded the operational 0.2 cutoff, corresponding to cosine similarity below 0.8 to every stored pattern. The kernel-density estimator had higher marginal error (1.9%) and lower pooled TF-only mechanistic cloud overlap (0.82): after pooling across patients, visits, and TGA features, 82% of the predicted-to-recorded TGA ratios for its synthetic patients fell within the real-data 5th–95th percentile range. By contrast, the mixture and copula were competitive with stochastic attention on marginal error (1.1% and 1.3%, respectively, versus 1.2%) and mechanistic overlap (0.92 and 0.93 versus 0.90). All 100 profiles generated by stochastic attention and by the copula exceeded the novelty cutoff, compared with 94% of mixture profiles. Stochastic attention nevertheless had the largest empirical cross-visit correlation error among the five methods (0.64, versus 0.42 for the mixture and copula). Because this metric compared each generated cohort with the same 23 patients used to construct it, it measured in-sample reproduction rather than population generalization. The separate known-ground-truth benchmark addressed the latter question: stochastic attention recovered the population covariance more accurately than the sample-covariance Gaussian generator in 35 of 39 low-rank Gaussian $n < p$ settings (Fig. 4A). Thus, the ablation attributed much of the marginal and mechanistic fidelity to the shared principal-component representation. Within this comparison, no generator dominated across the evaluated metrics.

2.2 Cross-Visit Covariance Structure

Having established that SA reproduces marginal distributions and pairwise relationships, we next asked whether it preserves the higher-order joint structure across visits. We compared the full 216×216 cross-visit correlation matrices for real, SA-generated, and MVN-generated populations. The real data exhibited a characteristic block structure with strong within-visit correlations along the diagonal and structured off-diagonal blocks reflecting cross-visit dependencies; for example, a patient’s Visit 1 Factor X level predicting their Visit 3 Factor X level. We found that SA preserved this block structure, including the off-diagonal cross-visit correlations (Fig. 3). The SA–Real residual matrix showed small, unstructured deviations distributed throughout, whereas the MVN–Real residual was notably smoother, reflecting the regularization-induced shrinkage of off-diagonal correlations toward zero. This difference was most apparent in the off-diagonal blocks, where MVN systematically underestimated cross-visit dependencies that SA retained. We note that Pearson correlation captures linear associations; Spearman rank correlations, used in the mechanistic validation, would additionally capture monotonic nonlinear dependencies.

We examined the eigenvalue spectrum of the sample covariance matrix to understand the structural origin of this difference (Fig. S2). The real data had rank 22, reflecting the $K=23$ patient constraint. SA and MVN make different trade-offs in handling this rank deficiency: SA truncates via PCA (zeroing variance beyond component 18), while MVN regularizes via Ledoit–Wolf shrinkage (providing a full-rank estimate by inflating eigenvalues in the 194 null dimensions). The consequence of MVN’s approach is that it introduces variance in dimensions where the data contain no signal, which manifested as increased dispersion in PCA projections. PCA projections by visit confirmed this dispersion gap (Supplementary Fig. S8). SA-generated patients occupied the same region of PCA space as real patients across all three visits, while MVN-generated patients showed substantially greater scatter, particularly at Visits 2 and 3.

2.3 Robustness to Sample Size, Dimensionality, and Nonlinearity

We evaluated robustness with three complementary tests (Methods, Sec. 5.5). The first test examined recovery of a known population covariance from a small, high-dimensional sample. We simulated low-rank Gaussian populations whose covariance was known before drawing any training data. For each combination of training size ($n=8, 15, 23, 40, 80$), dimensionality ($p=30, 120, 216$), and latent rank ($r=3, 5, 10$), we constructed the stochastic-attention memory matrix and estimated a sample-covariance Gaussian generator from the same training cohort, generated 100 profiles from each, and compared each generated covariance with the known population covariance using relative Frobenius error. This tested population recovery rather than reproduction of the training draw. Stochastic attention had lower error in 35 of the 39 rank-deficient $n < p$ settings (Fig. 4A). At the cohort-matched setting ($n=23, p=216$), averaged across the three latent ranks, the Gaussian generator’s error was about 1.5 times that of stochastic attention. Across the grid, the advantage was concentrated in $n < p$ settings, where the Gaussian sample covariance was rank deficient.

We next asked how the method performed when fewer real patients were represented in the memory matrix. We randomly selected 20, 15, 10, or 8 patients from the full cohort, repeating this process ten times at each size. For each subset, we constructed a new memory matrix and generated 100 synthetic profiles. We compared every synthetic cohort with the same complete set of 23 real patients, providing a common reference across all subset sizes. As n fell from 20 to 8, mean cross-visit correlation error rose from 0.63 to 1.11 and mean pooled marginal error rose from 2.0% to 4.9%, while mechanistic cloud overlap remained near 0.90 (Fig. 4B). The final test evaluated how the linear embedding performed when the population followed a curved rather than straight pattern. We used a noisy S-curve embedded in $p=30$ or 120 dimensions. We constructed the stochastic-attention memory matrix and estimated the Gaussian generator from $n=23$ samples, then measured both distance from the known curve and covariance error as curvature increased. The mean distance of stochastic-attention samples from the curve doubled relative to the straight-line case between curvatures 0.25 and 0.5. Its samples remained closer to the curve than the Gaussian samples, whereas the Gaussian generator recovered the population covariance more accurately. Thus, stochastic attention had lower covariance error than the Gaussian generator in most small, linear, low-rank settings, but the tests also identified deterioration with fewer patients and a clear limitation when the population structure was nonlinear.

2.4 Interpolation and Extrapolation

We then examined whether the sampler stayed within the real cohort or reached beyond it. We constructed the convex hull of the 23 real patients in the $d=18$ principal-component space and tested whether it contained each synthetic patient (Supplementary Table S14). This hull was the smallest convex region enclosing the observed patients. Every synthetic patient lay outside it. In a separate leave-one-out analysis, each real patient also lay outside the convex hull formed by the remaining 22 real patients. This indicated that every real patient formed an outer corner, or vertex, of the sparse, high-dimensional hull. This binary inside-or-outside result did not show whether synthetic patients lay just beyond the hull or much farther away. We therefore used distance to the hull to measure the degree of extrapolation. Our hypothesis was that unusually extrapolative generation would place synthetic patients farther from the hull than real patients treated as unseen. For a like-for-like comparison, the distances for both groups were calculated against hulls containing 22 real patients. The leave-one-out calculation already omitted the real patient being evaluated; for each synthetic patient, we omitted its nearest real patient. Median distance to the resulting hulls was 11.0 for synthetic patients and 11.9 for held-out real patients, and a descriptive Mann–Whitney test

did not detect a difference ($p=0.07$). Thus, the comparison provided no evidence that the synthetic profiles lay farther beyond the observed cohort than held-out real patients. This result did not establish equivalence, but it showed why binary hull membership alone was not an informative measure of extrapolation in this sparse, high-dimensional setting.

Because distance from the convex hull did not show whether a profile was biologically valid, we applied two additional checks. The first checked two biological constraints: quantities that should be positive could not be negative, and coagulation-factor or TGA values could not lie more than five real-cohort standard deviations from the mean. Fifty-four of 100 synthetic patients passed this check; the other 46 had a negative estradiol or progesterone value. The unconstrained multivariate-normal baseline had a similar failure rate, with negative hormone values in 43 of 100 patients. Neither generator enforced positivity. Hormone concentrations were closer to zero relative to their observed variation than the coagulation-factor levels, which explained why only the hormones crossed zero in these samples. The positive coagulation factors did not imply a general positivity guarantee. The second check used the BZ2012 model to determine whether the generated factor combinations produced thrombin-generation behavior consistent with the real cohort. Estradiol and progesterone were not inputs to this model, and all nine generated coagulation-factor inputs were positive, so all 100 synthetic patients were included. For each thrombin-generation feature, 83–92% of patients had a model-predicted-to-recorded ratio within the real cohort’s 5th–95th percentile range. Forty-one patients met this criterion for every feature and visit, and 24 passed both checks. These results showed that extrapolation beyond the observed cohort and biological plausibility were separate questions. Future implementations could enforce non-negativity through positive-variable transformations or constrained decoding.

2.5 Membership-Inference Risk

Each real patient’s 18-dimensional principal-component profile was one column of the memory matrix used during stochastic-attention sampling. We asked whether someone who already possessed a candidate de-identified assay profile could use the synthetic cohort to infer whether that profile was included in the matrix. This tested record inclusion, not personal identity; the source cohort contained no direct identifiers or links to named individuals. As an initial check, synthetic profiles were closer to their nearest real record than real records were to one another (median distance 13.8 versus 15.9). We tested membership more directly across 40 random splits. For each split, the memory matrix contained 15 patients, the other 8 were held out, and a new synthetic cohort was generated. A smaller distance from a candidate record to that cohort provided stronger evidence of inclusion (Supplementary Table S15). Included records were closer than held-out records (median 10.7 versus 15.4). The area under the receiver operating characteristic curve (AUC) measured how often the test ranked an included record above a held-out record; 0.5 indicated chance and 1.0 perfect discrimination. The mean AUC was 0.97, with perfect discrimination in 9 of 40 splits. Thus, the test usually inferred whether a candidate profile had been included, but could not identify a person or reconstruct an unknown patient’s measurements. To identify the source of this signal, we compared the numerical sampler with the target distribution. Metropolis-adjusted sampling left the result nearly unchanged (AUC 0.96). Lowering the inverse temperature across a 40-fold range increased the distance from synthetic to stored profiles, but it never reached the real-to-real baseline and cross-visit covariance fidelity degraded. Exact draws from the closed-form mixture did not use the Markov-chain sampler and reproduced the result (AUC 0.98; Supplementary Table S16, Fig. S9). Thus, at $K=23$, the membership signal was a property of the target distribution, not of the Markov-chain implementation. Synthetic generation therefore added no privacy guarantee beyond de-identification of the source cohort.

2.6 Conditional Generation of Rare Subgroups

The previous analyses used unconditioned SA, generating from the full cohort. A key clinical application, however, is generating patients from specific subpopulations that are too small to study independently. We therefore tested whether SA could amplify rare clinical subpopulations while preserving their condition-specific signatures. We defined three overlapping subgroups by pregnancy outcome and comorbidity: Uncomplicated ($n=18$), PCOS ($n=3$, cross-cutting; 1 patient also in the PE group), and Developed PE ($n=5$). The PCOS and Developed PE groups were too small for any independent statistical analysis. We used SA’s multiplicity-weighted sampling to generate condition-specific cohorts of 100 patients each by upweighting attention on the relevant stored patterns.

We compared condition-specific feature means between real and SA-generated patients across eight coagulation features that exhibited between-condition variation (Fig. 5). SA-generated cohorts preserved the condition-specific patterns. PCOS patients showed elevated Factor VIII and vWF relative to Healthy patients in both real and synthetic data, PE patients showed elevated α_2 -AP and Factor IX, and the between-condition rank ordering was maintained across all eight features. We then performed a [descriptive bootstrap Mann–Whitney diagnostic](#). For each feature and condition, we subsampled the synthetic cohort to match the real sample size (n_{real}), applied a Mann–Whitney U test, and repeated 1,000 times. We reported the fraction of replicates where $p > 0.05$, [meaning that this test did not detect a difference in that replicate. This power-limited fraction was descriptive and did not establish equivalence.](#) Across all 24 feature–condition pairs, 20 (83%) achieved $p > 0.05$ in $\geq 90\%$ of replicates, with a median no-difference-detected fraction of 98.6% (full per-pair results in Table S10). The four pairs that fell below 90% were high-variance features in the smallest subgroups (FIX and plasminogen in PCOS; FIX and plasminogen in Developed PE), where the real data exhibited large inter-patient variability. [No multiple-testing correction was applied because this diagnostic was not used to make simultaneous inferential claims; the per-pair values and the 90% threshold were descriptive reporting summaries.](#) PCA projections of the conditioned cohorts are provided in Figs. S3–S4. These results showed that SA’s multiplicity-weighted interpolation can amplify rare subgroups in a way that preserves their distinguishing clinical features, useful for [exploratory hypothesis generation and simulation-based workflow testing, but not for increasing the effective clinical sample size.](#) MVN has no mechanism for this: fitting a separate MVN to three PCOS patients across 216 features is mathematically impossible (rank 2 covariance), and post-hoc subsetting of unconditional MVN samples does not preserve subgroup-specific structure.

[The conditioned cohorts reproduced the observed subgroup signatures, but generating 100 profiles did not increase the amount of independent subgroup information. The PCOS and Developed PE signals still came from only three and five real patients; after accounting for the lower-weight background patterns, the corresponding participation ratios were 4.6 and 7.7 \(Table S2\), meaning that repeated draws were concentrated around only a few stored profiles. Treating the 100 synthetic profiles as independent observations would therefore understate uncertainty. Conditioning instead provided a denser set of profiles around the distributions suggested by those few patients, which could support exploratory hypothesis generation, workflow testing, and simulations for planning future studies. However, it did not add evidence about the conditions or replace recruitment of additional patients.](#)

2.7 Mechanistic Consistency

The preceding validations assessed statistical properties of the synthetic data. [We next tested whether factor combinations generated by stochastic attention produced biologically plausible out-](#)

puts when passed through a separately specified mechanistic model. The model was blind to the generation process but was calibrated on the same cohort. We used the Hockin–Mann BZ2012 model^{6,17}, a system of 58 ordinary differential equations with 64 rate constants that simulates thrombin generation from patient-specific coagulation factor inputs. We calibrated 5 of the 64 rate constants (prothrombinase, intrinsic and extrinsic tenase, protein C activation, and meizothrombin conversion; Table S4) on Visit 1 real patients at the population level and held the remaining 59 at published literature values. We then ran the calibrated model on all 23 real patients and 100 synthetic patients under both TF-only and TF+TM conditions (738 total simulations, 0 failures). For each patient, we computed five ODE-predicted thrombin generation features (lagtime, peak, time-to-peak, maximum rate, and ETP) and compared them against the corresponding values from the patient record. In these comparisons (Fig. 6, top row), the horizontal axis is the TGA value from the patient record, which exists for both real and synthetic patients, while the vertical axis is the value computed by the BZ2012 model from that patient’s factor inputs. Systematic deviations from the $y=x$ line reflect ODE model bias (e.g., lagtime is underpredicted at $\approx 0.77\times$), not a failure of the synthetic data; the key observation is that real and synthetic patients share the same pattern of model bias, occupying the same cloud with the same systematic offset.

We formalized this observation by computing the ODE-predicted/dataset ratio for each patient and comparing the resulting distributions between real and synthetic populations (Fig. 6, bottom row). These ratio distributions capture how the ODE model processes each patient’s factor levels. If SA had generated patients with biologically implausible factor combinations, the model would process them differently, producing a shifted or broadened ratio distribution. Instead, the distributions overlapped substantially, with cloud overlap (fraction of synthetic ratios within the 5th–95th percentile of real ratios) ranging from 0.83 for ETP to 0.92 for T.Peak under TF-only conditions, and two-sample Kolmogorov–Smirnov (KS) tests did not detect a difference between the real and synthetic ratio distributions at this sample size across all five TGA features ($D = 0.081$ – 0.123 , all $p > 0.30$; full diagnostics in Table S11). Under TF+TM conditions, the model systematically overcorrected the protein C anticoagulant pathway, producing peak and maximum rate predictions approximately $0.5\times$ measured values for both real and synthetic patients (Fig. S5, same layout as the main-text comparison), yet this shared systematic bias was itself informative. Cloud overlap remained high (0.88–0.92 across the five features), confirming that the model processed real and synthetic patients identically even under conditions where the calibration was imperfect.

The calibration also generalized across pregnancy timepoints. Although fitted only on Visit 1 real patients, the per-visit median predicted-to-measured ratios under TF-only conditions remained close to 1.0 at Visits 2 and 3 (Fig. S6), confirming that the calibrated rate constants captured a stable population-level property of the coagulation system rather than overfitting to the training visit. Spearman rank correlations between predicted and measured values were moderate for ETP ($\rho \approx 0.6$ – 0.8 for real, $\rho \approx 0.6$ – 0.65 for synthetic; Fig. S7) and weak for other features, consistent with a 5-parameter population-level calibration that was not designed to resolve inter-patient variability. The key finding was not patient-level prediction but population-level plausibility. Applied identically to both groups, the mechanistic model produced overlapping predicted-to-recorded ratio distributions for real and synthetic patients under both experimental conditions. Because the model was calibrated on this same cohort, this was a same-cohort consistency check rather than independent validation.

2.8 Downstream Utility: Mechanistic Model Calibration

The preceding results established that SA-generated patients are statistically and mechanistically plausible. To test whether this plausibility translates to practical utility, we calibrated a mechanistic

model entirely on synthetic patients and evaluated whether it could predict real patient outcomes. We calibrated the BZ2012 model (TF-only condition, 5 rate constants) on synthetic V1 patients ($N=100$) using the same bounded Nelder–Mead optimization applied to the real V1 calibration ($K=23$), with 11 random restarts to mitigate local-minimum sensitivity (Table S4 for parameter details). We then evaluated both calibrations on held-out real V2 and V3 patients that neither model had seen during training. The synth-calibrated model achieved comparable or slightly better performance across all five TGA features, with per-feature median relative errors 2–10% lower than the real-calibrated model (overall ratio $0.94\times$; Fig. 7; per-feature breakdown in Table S12). This improvement likely reflects the smoother loss landscape afforded by $N=100$ synthetic training patients compared with $K=23$ real ones, reducing overfitting during optimization. The two calibrations found different parameter values but produced highly correlated predictions across all five TGA features, confirming that the synthetic data captured the same underlying coagulation structure. We note that the synthetic training data is one step removed from the real data (real $\rightarrow$ SA $\rightarrow$ synthetic $\rightarrow$ calibration), so the synth-calibrated model is not fully independent of the real data; rather, it demonstrates that SA’s representation preserves sufficient biological structure for a mechanistic model to learn generalizable parameters.

The modest reduction in prediction error did not show that synthetic data were superior. All 100 synthetic profiles were derived from the same 23 patients, so they did not increase the independent evidence available about the broader patient population. Their value was computational: averaging the calibration objective over a denser sample from the represented distribution may have made the optimization more stable. The results therefore showed that synthetic profiles could support mechanistic-model calibration, exploratory comparisons of small subgroups, and factor-level hypotheses for testing in larger studies. However, confirming those hypotheses and establishing clinical utility would require additional real patients or independent experimental data.

3 Discussion

This study tested whether a geometry-preserving generative method could produce synthetic patients from a very small longitudinal cohort that were not merely statistically similar to real patients but consistent with the underlying coagulation biology. Our results provided evidence at each validation level that SA-generated patients met this standard. Individual feature distributions were reproduced with median MRE of 1.2%, which MVN also achieved by construction. SA additionally preserved the joint cross-visit covariance structure that MVN’s rank-deficient estimate could not represent, a difference made visible in the eigenvalue spectrum where MVN introduced spurious variance in 194 null dimensions while SA respected the low-rank geometry of the data. SA could amplify clinically meaningful rare subgroups, generating 100 synthetic PCOS patients from only 3 real ones, while preserving subgroup-specific signatures that MVN cannot capture. And a separately specified, generator-blind ODE model of the coagulation cascade, calibrated on the same real cohort, placed the synthetic profiles in the same factor-to-thrombin-generation envelope as the real profiles, with 83–92% cloud overlap and no difference detected by the Kolmogorov–Smirnov tests under these conditions ($p > 0.35$ for all five TGA features).

The ability to generate validated synthetic cohorts from very small longitudinal datasets has practical implications for maternal health research. Rare pregnancy complications such as PCOS and preeclampsia are difficult to study because assembling large, well-characterized longitudinal cohorts requires years of recruitment across multiple clinical sites. Our results suggest that SA can enable hypothesis generation, power analyses, and computational modeling studies for these understudied populations by generating synthetic cohorts that preserve both the statistical and

mechanistic properties of the real data. The conditional generation capability is particularly relevant. Rather than requiring researchers to collect additional rare patients, SA can provide denser simulated realizations of the distribution implied by an existing small subgroup for exploratory modeling and workflow testing. It does not increase the subgroup’s inferential sample size: that remains three patients for PCOS and five for Developed PE, and recruitment of additional real patients is still required for clinical inference.

A further intended use was to train deep learning models in both directions: predicting TGA responses from coagulation factors and estimating factor profiles compatible with an observed TGA response. Related modeling work has predicted thrombin dynamics from factor levels and searched for factor changes that move thrombin generation toward a target²⁹. We did not evaluate either task here. Synthetic profiles could support model development or initial training, but performance would still need to be tested on held-out or external real patients. Because more than one factor profile may produce a similar TGA response, an inverse model should return a distribution or set of compatible profiles rather than a single recovered profile.

Understanding why the method worked in this setting required separating the principal component representation from the sampler. The ablation showed that the shared 18-dimensional subspace supplied much of the marginal and mechanistic fidelity. Stochastic attention used the patient profiles in that subspace as its memory matrix and avoided estimating a full 216-dimensional covariance from 23 patients. The competitive mixture and copula baselines showed, however, that the benefit of this representation was not unique to stochastic attention. A natural concern is whether the β^* selection procedure, originally developed for protein sequence generation³⁰, transfers to clinical coagulation data. The entropy transition rule selected the point of greatest downward curvature in attention entropy, a geometric property of K patterns in d dimensions rather than a property of what the features represented. For this dataset it gave $\beta^* \approx 2.94$, with the random-pattern scale $\sqrt{d} \approx 4.24$ serving as a theoretical reference, and a sensitivity analysis confirmed that generation quality degraded gracefully across $\beta/\beta^* \in [0.1, 3.0]$ (Table S1). Similarly, varying the PCA variance retention threshold from 85% to 99% ($d_{\text{PCA}} = 13$ to 21) had minimal impact on generation quality. Median MRE remained between 1.2% and 1.8% across all thresholds (Table S7), confirming that the 95% threshold was not a critical design choice.

The multiplicity weighting for conditional generation introduced an additional consideration. Achieving 80% of finite-mixture component probability on the PCOS subgroup required $\rho \approx 26.7$, reducing the participation ratio to $K_{\text{eff}} \approx 4.6$ (multiplicity parameters for all three subgroups in Table S2), which indicated that mixture-component probability was concentrated around a small number of stored profiles across repeated draws. The subgroup-specific signal still rested on only three real PCOS profiles, and the lower-weight background profiles pulled generated means toward the population average; K_{eff} did not represent an independent patient count. Magnitudes were drawn from all 23 empirical PCA-space norms, so conditioning acted through the weighted directions rather than a subgroup-specific magnitude distribution. Related studies from this group applied the same geometry-based operating-point rule and multiplicity weighting to protein sequence generation and conditional steering^{27,30,31}. These studies showed that the sampling framework was not limited to coagulation data, although they did not replace validation in an independent clinical cohort.

An alternative precedent for validated longitudinal synthetic generation is the VAMBN-MT framework of Kühnel et al.²⁰, which extends a Variational Autoencoder Modular Bayesian Network with an LSTM time-encoder to better capture cross-visit dependencies. They evaluated VAMBN-MT on the DONALD nutritional cohort and on the Alzheimer’s Disease Neuroimaging Initiative (ADNI) cohort. Their study is instructive because it offers a regime contrast to ours. With $N=1,312$ DONALD participants and 33 longitudinal variables, their setting sits in the data-rich

regime for which VAMBN-MT was designed, where deep parametric models can be trained directly on the cohort. The published configuration combines expert-curated module assignments, Bayesian-network edge constraints, an LSTM-augmented HI-VAE per module, and 10,000 synthetic samples drawn from the trained model to reproduce a published age-trend regression. In their evaluation, the best variant reported a cross-visit correlation matrix relative Frobenius error of roughly 0.70, and downstream time trends weaker than the strongest signals in the real data were not reproduced at the sample sizes they tested, both honest indicators of how hard this generative task is even in the data-rich regime. On the other hand, our setting inverts both the number of patients and the number of features. With $K=23$ patients across 216 features we operate two orders of magnitude smaller and squarely in the $n < p$ regime, well below the cohort sizes for which VAMBN-MT was developed. The non-parametric design of SA addresses this regime in three ways that contrast with VAMBN-MT. First, PCA-derived linear geometry replaces the modular HI-VAE architecture as the mechanism for capturing cross-visit dependencies. Second, a generic mechanistic-model validation pattern, transferable to any domain with a calibrated ODE model, replaces the domain-specific dependency rules used in their evaluation, such as graduation-status monotonicity, age-time arithmetic identities, and macronutrient summation. Finally, multiplicity weighting provides an inference-time conditional-generation lever that the VAMBN-MT framework does not offer. These are complementary positions in the design space rather than a head-to-head comparison. VAMBN-MT addresses the moderate-cohort regime where deep generative models are feasible but expert-specified modular structure helps capture cross-visit dependencies, while SA addresses the small-cohort regime where parametric models cannot be fit and the geometry of the few real patients must itself carry the generative signal.

Beyond these statistical and geometric considerations, the mechanistic validation introduced an approach that applies beyond coagulation. Rather than relying solely on statistical comparisons between real and synthetic distributions, we tested whether synthetic patients produced biologically consistent outputs when processed by a separately specified computational model that was blind to the generator but calibrated on the same real cohort, and went further to show that a mechanistic model calibrated entirely on synthetic data predicted real patient outcomes as well as one calibrated on real data (overall ratio $0.94\times$ on held-out V2+V3 patients, with 11 random restarts per calibration). This approach is available whenever a validated mechanistic model exists for the domain of interest (pharmacokinetic models in drug development³², physiological models in critical care³³, tumor growth models³⁴, or metabolic models^{35,36}), and the requirement is not that the model be a perfect predictor but that synthetic patients fall within the same model-predicted envelope as real patients. In the descriptive bootstrap Mann–Whitney diagnostic of condition-specific features, 4 of 24 (17%) feature–condition pairs fell below the criterion of no detected difference in at least 90% of equal-size resamples. This power-limited diagnostic is not an equivalence test.

The mean relative errors for these features remained small (5–15%), and the lower non-detection fractions were consistent with compressed tails rather than shifted means. This compression could reflect both PCA truncation and the finite empirical distribution of 23 PCA-space norms. The clinical consequence is important. A synthetic cohort that under-represents severe states cannot be used to estimate how often those states occur or how severe they become. Although synthetic profiles extended beyond the real cohort’s convex hull, that result did not show that they reproduced the true tail distribution. We therefore do not recommend using these cohorts to estimate extreme risk without additional real data. The primary mechanistic cloud-overlap measure also focused on the real 5th–95th percentile range and therefore did not validate behavior in the tails. More generally, several findings reflected the limits of generating from only 23 source patients. Hull membership had little resolution, membership in the memory matrix remained detectable, and conditioned cohorts still relied on three PCOS or five Developed PE patients. The modest calibration

improvement was consistent with denser sampling of the same patient-derived distribution, not with added biological information. These findings supported simulation and model testing, but they did not increase the clinical evidence supplied by the original cohort.

This study had several limitations. The clinical analysis used one longitudinal cohort of 23 patients. The known-covariance simulations and the protein studies tested the method outside this single coagulation dataset, but neither replaced validation in an independent clinical cohort. The BZ2012 model was also calibrated on the same cohort used for generation. Applying one fixed, generator-blind model to both groups tested whether it treated real and synthetic profiles similarly; it was not independent biological validation. The calibration required a 50-fold reduction in intrinsic tenase k_{cat} and a 15-fold increase in extrinsic tenase k_{cat} relative to published values (Table S4), likely because the published measurements and the plasma-based assay used different experimental conditions. These fitted values should not be read as new biochemical estimates. The mechanistic test covered thrombin generation but not the fibrinolytic or viscoelastic features. We also did not test clinical outcomes. Independent clinical data, broader mechanistic models, and prospective outcome validation are therefore still needed.

Concatenating visits allowed the method to preserve cross-visit relationships, but it required every patient to have the same assays at the same aligned visits. The current pipeline could not directly accept irregular timing, missing assays or visits, or variable-length records. Imputation, padding, or a sequence-aware extension could support such data, but each would add assumptions that require validation in a small cohort. The PCA representation imposed a related limit because it captured linear structure. In the S-curve benchmark, generated profiles moved farther from the known curve as curvature increased. Stochastic attention remained closer to the curve than the Gaussian generator, whereas the Gaussian generator recovered covariance more accurately. Neither method was best on both measures in this benchmark. Nonlinear embeddings could better represent curved or branching relationships, but they would add fitting and decoding choices that are difficult to estimate from 23 patients. The linear representation was therefore a practical choice for this cohort, not evidence that clinical trajectories are generally linear.

Computational cost and behavior at larger cohort sizes are separate questions. The reported Langevin sampler required work proportional to TKd for each profile. The exact mixture sampler removed the T sequential updates: its component weights could be prepared once, after which each profile required a component draw, a d -dimensional Gaussian draw, and decoding. The principal-component model was estimated once, but inverse transformation and decoding were still performed for every generated profile. The simulation advantage over the sample-covariance Gaussian was concentrated in the $n < p$ settings. We did not test cohorts larger than 23 patients, so stability and membership risk as K grows remain open questions.

A separate limitation concerns membership privacy. The source cohort was fully de-identified, so neither the synthetic cohort nor our analysis identified a named participant. The narrower test assumed that someone already possessed a candidate assay profile and asked whether it had been included in the memory matrix. The test usually ranked included records above held-out records, with a mean AUC of 0.97. Metropolis-adjusted sampling, the tested inverse-temperature range, and exact sampling all retained this signal. The exact sampler showed why: at $K=23$, the target was a finite Gaussian mixture with one component centered on each stored profile. The signal therefore came from the target distribution rather than from numerical sampling error. The method was not differentially private, and patient-level synthetic records should not be treated as anonymized data. The intended uses, including mechanistic calibration and hypothesis generation, can be performed without releasing patient-level records.

4 Conclusion

We have demonstrated that multiplicity-weighted Stochastic Attention can generate synthetic longitudinal patient data from very small clinical cohorts ($K=23$) that pass four levels of validation: marginal plausibility, cross-visit covariance preservation, conditional subgroup amplification, and mechanistic consistency with a coagulation ODE model calibrated on the same cohort. A downstream utility test showed that a mechanistic model calibrated entirely on synthetic patients predicted real patient outcomes as well as one calibrated on real data. The SA framework preserves the geometry of the clinical data in regimes where parametric approaches such as multivariate normal sampling fail due to rank deficiency. An analytic sampler for the same latent target reproduced the same fidelity and privacy results, and offers a faster route for future use. Combined with a multi-level validation framework, this approach provides a proof-of-concept path for generating synthetic cohorts useful for the modeling use-cases demonstrated here to support computational research in rare conditions and small longitudinal studies.

5 Methods

5.1 Population Dataset

We analyzed a longitudinal dataset of coagulation, fibrinolysis, thrombin generation assay (TGA), and rotational thromboelastometry (ROTEM) measurements from women desiring a pregnancy. The study was approved by the Institutional Review Board at the University of Vermont. Written consent was obtained from all participants prior to enrollment. Blood collection, plasma processing, and assay protocols have been described previously^{37–40}. Briefly, citrated platelet-poor plasma was collected without tourniquet use after supine rest at three clinical visits: Visit 1 (pre-pregnancy baseline in the follicular phase), Visit 2 (end of first trimester), and Visit 3 (mid-third trimester). Plasma aliquots were stored at -80°C for subsequent analysis. Measurements of secondary analytes were conducted as follows: activity levels of coagulation factors were determined by Stago clotting assays, fibrinolysis proteome levels were measured by ELISA, and hormone levels were determined by the CLIA-certified laboratory at the University of Vermont Medical Center. TGA was performed using 5 pM relipidated tissue factor with fluorogenic substrate on a Synergy4 plate reader as described in McLean et al.³⁷. Clot formation and fibrinolysis dynamics were assessed via viscoelastometry using the ROTEM Delta Instrument (Werfen). Thawed plasma was mixed with or without 4 nM tPA (with respect to plasma volume), recalcified with 15 mM calcium chloride, and incubated for 3 minutes at 37°C . The reaction was initiated by adding the plasma mixture to ROTEM cups containing tissue factor at a 17:1 ratio; the final concentration of TF in the reaction was 5 pM.

The dataset contained 153 total records across approximately 50 subjects. After removing columns with $> 30\%$ missing values and retaining only subjects with complete data across all three visits, we obtained $K = 23$ complete longitudinal profiles across $n = 72$ assay measurements per visit. The 72 assay measurements spanned five categories: (i) hormones (estradiol, progesterone), (ii) coagulation factors and inhibitors (Factors II, V, VII, VIII, IX, X, XI, XII; antithrombin (AT), protein C (PC), tissue factor pathway inhibitor (TFPI), von Willebrand factor (vWF), fibrinogen, etc.), (iii) fibrinolytic markers (plasminogen, plasminogen activator inhibitor-1 (PAI-1), plasminogen activator inhibitor-2 (PAI-2), α_2 -antiplasmin (α_2 -AP), thrombin-activatable fibrinolysis inhibitor (TAFI), tissue plasminogen activator (tPA), D-dimer equivalents), (iv) thrombin generation parameters under four initiator conditions (TF at 5 pM, TF+thrombomodulin (TM), No TF, No TF+anti-FXIIa), and (v) ROTEM/viscoelastic parameters (clotting time (CT), maximum clot

firmness (MCF), alpha angle, lysis onset time (LOT), maximum lysis (ML), lysis time (LT), area under the curve (AUC) under TF Only and TF+tPA conditions).

Demographics and clinical characteristics of the resulting cohort, including age at enrollment, prepregnancy BMI, parity, gestational ages at each visit, and the cross-tabulation of enrollment cohort against pregnancy outcome, are summarized in Table 1. The 23 patients were drawn from three enrollment conditions: Healthy nulliparous women ($n=14$), women with a personal history of preeclampsia from a previous pregnancy (Prior PE, $n=6$), and women with polycystic ovarian syndrome (PCOS, $n=3$; 2 individuals were nulliparous). For conditional generation, we defined three overlapping subgroups based on pregnancy outcome and comorbidity rather than enrollment condition: Uncomplicated ($n=18$, all patients whose study pregnancy did not result in preeclampsia), PCOS ($n=3$, a cross-cutting comorbidity; 1 patient also developed PE), and Developed PE ($n=5$, patients who developed preeclampsia during the study pregnancy, drawn from 2 healthy-enrolled, 2 prior-PE-enrolled, and 1 PCOS-enrolled participants).

5.2 Multiplicity-Weighted Stochastic Attention

We generated synthetic patients using a multiplicity-weighted variant of Stochastic Attention (SA) rooted in modern Hopfield network theory²⁶. We stored K memory patterns as columns of a matrix $\hat{\mathbf{X}} \in \mathbb{R}^{d \times K}$ and defined a weighted Hopfield energy landscape:

$$E_r(\boldsymbol{\xi}) = \frac{1}{2} \|\boldsymbol{\xi}\|^2 - \frac{1}{\beta} \log \sum_{k=1}^K r_k \exp \left(\beta \mathbf{m}_k^\top \boldsymbol{\xi} \right) \quad (1)$$

where $\{\mathbf{m}_k\}_{k=1}^K$ are stored memory patterns, β is an inverse temperature parameter controlling the sharpness of retrieval, and $\{r_k\}_{k=1}^K$ are per-pattern multiplicity weights. When all $r_k = 1$, this reduces to the standard modern Hopfield energy; when $r_k = \rho > 1$ for a designated subset of patterns, the energy landscape is continuously deformed to favor that subset.

This energy defines a probability distribution that can be written exactly. Writing $\pi_r(\boldsymbol{\xi}) \propto \exp[-\beta E_r(\boldsymbol{\xi})]$ and completing the square in each exponent,

$$\exp \left[-\frac{\beta}{2} \|\boldsymbol{\xi}\|^2 + \beta \mathbf{m}_k^\top \boldsymbol{\xi} \right] = \exp \left(\frac{\beta}{2} \|\mathbf{m}_k\|^2 \right) \exp \left[-\frac{\beta}{2} \|\boldsymbol{\xi} - \mathbf{m}_k\|^2 \right], \quad (2)$$

so the resulting sampling distribution is exactly a finite isotropic Gaussian mixture,

$$\pi_r(\boldsymbol{\xi}) = \sum_{k=1}^K q_k \mathcal{N}(\mathbf{m}_k, \beta^{-1} \mathbf{I}), \quad q_k \propto r_k \exp \left(\frac{\beta}{2} \|\mathbf{m}_k\|^2 \right). \quad (3)$$

Each stored patient is the center of one component. The common component variance is β^{-1} , and the multiplicity weights enter the mixture probabilities directly. The memories used here are normalized to unit length, so the exponential factor is shared by every component and $q_k = r_k / \sum_j r_j$ exactly. We could therefore sample this distribution directly: first choose k from q and then draw $\boldsymbol{\xi} \sim \mathcal{N}(\mathbf{m}_k, \beta^{-1} \mathbf{I})$. This procedure requires no Markov chain. The cohort analyzed in this work was generated with the Langevin implementation given next, and that is what the reported results describe. We added the direct-sampling procedure afterward as a check, using the same normalization, magnitude rescaling, and decoding steps (Table S16).

We sampled from this landscape using an Unadjusted Langevin Algorithm (ULA):

$$\boldsymbol{\xi}_{t+1} = (1 - \alpha) \boldsymbol{\xi}_t + \alpha \hat{\mathbf{X}} \text{softmax} \left(\beta \hat{\mathbf{X}}^\top \boldsymbol{\xi}_t + \log \mathbf{r} \right) + \sqrt{\frac{2\alpha}{\beta}} \boldsymbol{\varepsilon}_t \quad (4)$$

where α is a step size, $\boldsymbol{\varepsilon}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, and $\log \mathbf{r}$ is the element-wise log of the multiplicity vector. The deterministic component drove $\boldsymbol{\xi}_t$ toward a multiplicity-weighted combination of stored patterns, while the stochastic component enabled exploration and novelty. The $\log r_k$ bias in the softmax logits provided a minimal, theoretically grounded mechanism for continuously interpolating between unconditioned generation ($\rho = 1$, all patterns equally weighted) and hard subset curation ($\rho \rightarrow \infty$, only designated patterns contribute). This is an inference-time operation that requires no change to the stored patterns.

The synthetic cohort analyzed in this work was produced by initializing each Langevin chain near a randomly selected stored pattern and retaining the final state after $T=2000$ steps. Because the starting point and chain length could affect the output, we compared four sampling procedures. These used the original initialization near a stored pattern, random initialization on the unit sphere, pooled samples from multiple chains after burn-in and thinning, and Metropolis-adjusted sampling. As a separate check, we also drew directly from Eq. (3), without using a Markov chain. The fidelity and membership-inference results were similar across all four procedures and the direct draws. This agreement indicated that the original implementation reproduced the evaluated behavior of the sampling distribution at β^* (Supplementary Fig. S9).

The inverse temperature β affects Hopfield retrieval and direct sampling in different ways. In the Hopfield update, small β spreads attention across many stored profiles and tends toward their weighted average, whereas large β concentrates attention on individual profiles. When sampling directly from Eq. (3), the stored profile is chosen according to q_k . For the unit-normalized profiles used here, these probabilities are independent of β and are set by the multiplicity weights. Instead, β controls the spread around the selected profile through the covariance $\beta^{-1}\mathbf{I}$. Small β therefore produces broader draws, whereas large β produces tighter draws. The attention-entropy transition used to define β^* marks the shift from attention shared across profiles to attention concentrated on individual profiles. We selected β^* from the transition in normalized attention entropy across a logarithmic sweep of β . Specifically, β^* was the grid point where the entropy curve bent downward most strongly, calculated as the most negative finite-difference second derivative with respect to $\log \beta$ (Supplementary Methods). This choice depended on the relative positions of the K profiles in d dimensions, not on what their features represented. For this dataset ($K=23$, $d_{\text{PCA}}=18$), it gave $\beta^* \approx 2.94$; the random-pattern prediction $\sqrt{d} \approx 4.24$ served as a theoretical reference. A sensitivity analysis across $\beta/\beta^* \in [0.1, 3.0]$ showed that the selected value balanced novelty and fidelity (Table S1). For conditioned generation, we repeated the entropy calculation with the multiplicity bias included, giving a ρ -dependent $\beta^*(\rho)$.

For unit-normalized memories, the component probabilities satisfy $q_k = r_k / \sum_j r_j$. We summarized how concentrated these probabilities were with the participation ratio $K_{\text{eff}} = 1 / \sum_k q_k^2 = (\sum_k r_k)^2 / \sum_k r_k^2$. This quantity is the effective number of equally weighted components needed to produce the same concentration. Equal weights gave $K_{\text{eff}} = 23$, whereas concentrating weight on three stored profiles moved it toward three. It showed whether repeated draws were spread across many stored profiles or concentrated near a few; it was not an independent clinical sample size. To make a designated subgroup account for a target fraction f_{target} of the component probability, we used $\rho = f_{\text{target}} K_{\text{bg}} / [K_{\text{des}}(1 - f_{\text{target}})]$, where K_{des} and K_{bg} are the designated and background pattern counts.

5.3 Pipeline for Continuous Longitudinal Data

Applying SA to continuous longitudinal clinical data required several adaptations beyond the standard Hopfield sampling framework. We first concatenated each patient’s $n=72$ assay measurements across all three visits into a single vector of dimension $d_{\text{concat}} = 3n = 216$, so that each stored mem-

ory pattern encoded a complete longitudinal profile and generated synthetic patients would have internally consistent multi-visit trajectories. We then standardized each feature to zero mean and unit variance and applied PCA retaining 95% of the variance, reducing dimensionality from 216 to $d_{\text{PCA}}=18$. This compression served two purposes: it removed collinearity among the 216 features and yielded a memory-to-dimension ratio of $K/d_{\text{PCA}} \approx 1.28$, placing SA in a favorable operating regime where the number of stored patterns exceeded the dimensionality of the representation.

Standard SA operates on unit-norm patterns, which is appropriate for discrete data (e.g., one-hot protein sequences) but destroys the anisotropic variance structure of continuous clinical measurements, collapsing dispersion across all principal components to a unit sphere. We addressed this with a direction-magnitude decomposition. Before normalization, we recorded each pattern’s PCA-space norm $s_k = \|\mathbf{m}_k\|$; we then ran SA on unit-norm patterns to obtain a directional sample $\hat{\xi}$, drew a magnitude s from the empirical distribution of $\{s_k\}$, and reconstructed the rescaled sample as $\xi_{\text{rescaled}} = s \cdot \hat{\xi}$. This preserved the directional structure learned by the Hopfield energy while restoring the natural scale of variation. The rescaled PCA vector was mapped back to the original feature space via inverse PCA and destandardization, then split into three 72-dimensional visit records.

To generate condition-specific cohorts (e.g., PCOS-only), we set the multiplicity $r_k = \rho$ for patterns belonging to the target subgroup and $r_k = 1$ for all others, and sampled using the weighted Langevin update (Eq. 4). The multiplicity ρ was computed to place a target fraction $f_{\text{target}} = 0.80$ of the mixture-component probability on the designated subset, while retaining 20% probability on background patterns, via $\rho = f_{\text{target}} \cdot K_{\text{bg}} / (K_{\text{des}} \cdot (1 - f_{\text{target}}))$. Multiplicity parameters for all three subgroups are listed in Table S2. For the PCOS subgroup ($K_{\text{des}}=3$), this required $\rho \approx 26.7$, reducing the participation ratio to $K_{\text{eff}} \approx 4.6$. **For both unconditioned and conditioned generation, magnitudes were drawn from the empirical distribution of all 23 PCA-space norms. Subgroup weighting affected directional sampling, not the distribution of magnitudes.**

The complete generation pipeline is summarized in Algorithm 1. The full set of SA hyperparameters used for all experiments is listed in Table S3, with key operating values $\alpha = 0.01$, $T = 2,000$ Langevin iterations per sample, and random seed 42 for reproducibility. A control experiment comparing ULA with the Metropolis-Adjusted Langevin Algorithm (MALA) confirmed that the discretization bias is negligible at this step size. Across 10 chains of 5,000 iterations the MALA acceptance rate was $99.9 \pm 0.03\%$, and mean post-burn-in energies and effective sample sizes were indistinguishable between the two samplers (Supplementary Table S8). The pipeline was implemented in Julia (v1.12) using the packages listed in the code repository. Generating 100 synthetic patients required approximately 0.1 seconds on a standard laptop (Apple M-series), comparable to MVN sampling (≈ 0.15 s), because SA operates in the 18-dimensional PCA space rather than the full 216-dimensional feature space. Each synthetic patient was generated by an independent Langevin chain (no shared state between patients); generation quality was stable across $T \in [500, 4000]$ iterations (median MRE 1.0–1.6%), confirming that $T=2,000$ was sufficient for convergence.

5.4 Baseline and Validation Framework

As a baseline, we generated synthetic patients by fitting a single multivariate normal (MVN) distribution to the same $K=23$ concatenated 216-dimensional patient profiles. Critically, both SA and MVN operated on the same concatenated representation. Each patient’s three visits were joined into one vector before generation, so that both methods had the opportunity to capture cross-visit dependencies. For SA, this structure was preserved through the Hopfield energy landscape operating on 18-dimensional PCA projections of the concatenated vectors. For MVN, the

Algorithm 1 Multiplicity-Weighted SA for Longitudinal Patient Generation

Require: Patient data $\{\mathbf{x}_k^{(v)}\}$ for $k = 1, \dots, K$ patients, $v = 1, 2, 3$ visits

Require: Multiplicity vector $\mathbf{r}$ (all ones for unconditioned; $r_k = \rho$ for designated subset)

Ensure: N synthetic longitudinal patient profiles

```

1: Concatenate:  $\mathbf{x}_k \leftarrow [\mathbf{x}_k^{(1)}; \mathbf{x}_k^{(2)}; \mathbf{x}_k^{(3)}] \in \mathbb{R}^{216}$ 
2: Standardize:  $\mathbf{z}_k \leftarrow (\mathbf{x}_k - \boldsymbol{\mu}) \oslash \boldsymbol{\sigma}$ 
3: PCA:  $\mathbf{m}_k \leftarrow \mathbf{W}^\top \mathbf{z}_k \in \mathbb{R}^{18}$  ▷ 95% variance retained
4: Record norms:  $s_k \leftarrow \|\mathbf{m}_k\|$ ; normalize  $\hat{\mathbf{m}}_k \leftarrow \mathbf{m}_k / s_k$ 
5: Select  $\beta^*$  where the attention-entropy curve bends downward most strongly, including  $\log \mathbf{r}$ 
   when conditioned
6: for  $i = 1, \dots, N$  do
7:   Select anchor  $k_0$  from all patterns, or from the designated subset when conditioned
8:    $\hat{\boldsymbol{\xi}}_0 \leftarrow \hat{\mathbf{m}}_{k_0} + 0.01\boldsymbol{\varepsilon}_0$ ; normalize  $\hat{\boldsymbol{\xi}}_0$ 
9:   for  $t = 0, \dots, T - 1$  do ▷ Langevin dynamics
10:     $\hat{\boldsymbol{\xi}}_{t+1} \leftarrow (1 - \alpha)\hat{\boldsymbol{\xi}}_t + \alpha \hat{\mathbf{M}} \text{softmax}(\beta^* \hat{\mathbf{M}}^\top \hat{\boldsymbol{\xi}}_t + \log \mathbf{r}) + \sqrt{2\alpha/\beta^*} \boldsymbol{\varepsilon}_t$ 
11:   end for
12:   Draw magnitude:  $s \sim \{s_k\}_{k=1}^K$ 
13:   Rescale:  $\mathbf{m}_{\text{synth}} \leftarrow s \cdot \hat{\boldsymbol{\xi}}_T / \|\hat{\boldsymbol{\xi}}_T\|$ 
14:   Inverse PCA + destandardize:  $\mathbf{x}_{\text{synth}} \leftarrow \boldsymbol{\sigma} \odot (\mathbf{W} \mathbf{m}_{\text{synth}} + \bar{\mathbf{z}}) + \boldsymbol{\mu}$ 
15:   Split into visit records:  $\mathbf{x}_{\text{synth}}^{(1)}, \mathbf{x}_{\text{synth}}^{(2)}, \mathbf{x}_{\text{synth}}^{(3)}$ 
16: end for

```

216 × 216 sample covariance matrix had rank at most 22 and was regularized using Ledoit–Wolf shrinkage²¹, which inflated eigenvalues in the 194 null dimensions. We then evaluated both SA and MVN synthetic patients through four progressively stringent validation levels. **Level 1 (Marginal Plausibility)** tested whether individual feature distributions matched, using per-feature means, standard deviations, and known physiological relationships. **Level 2 (Joint Structure)** tested whether the cross-visit covariance was preserved, by comparing full 216 × 216 correlation matrices, eigenvalue spectra, and PCA projections between real, SA, and MVN populations. **Level 3 (Conditional Structure)** tested whether rare subgroups could be amplified while preserving condition-specific signatures, by generating cohorts of 100 patients each from the Uncomplicated ($n=18$), PCOS ($n=3$), and Developed PE ($n=5$) subgroups using multiplicity-weighted sampling. **Level 4 (Mechanistic Consistency)** tested whether synthetic patients produced biologically plausible outputs under a separately specified, generator-blind mechanistic model calibrated on the same real cohort. We used the Hockin–Mann BZ2012 coagulation model^{6,17}, a system of 58 ODEs with 64 rate constants, in which 5 rate constants were calibrated on Visit 1 real patients at the population level and the remaining 59 were held at literature values. We ran the model on all real and synthetic patients under two conditions (TF-only and TF+TM) and compared predicted-versus-measured thrombin generation features. The primary metric was cloud overlap: the fraction of synthetic predicted/measured ratios falling within the 5th–95th percentile range of real ratios.

5.5 Robustness Tests

We used three tests, each designed to answer a separate question: whether generated profiles recovered a known population covariance, how performance changed as the clinical training set became smaller, and how performance changed as a data pattern became more nonlinear. First,

to test population covariance recovery, we generated training cohorts from a low-rank Gaussian factor model, $\mathbf{X} = \mathbf{Z}\mathbf{L}^\top + \boldsymbol{\epsilon}$, where $\mathbf{Z} \in \mathbb{R}^{n \times r}$ and $\mathbf{L} \in \mathbb{R}^{p \times r}$ had independent standard-normal entries, and $\boldsymbol{\epsilon}$ had independent Gaussian entries with variance 0.1. The population covariance $\boldsymbol{\Sigma}_{\text{pop}} = \mathbf{L}\mathbf{L}^\top + 0.1\mathbf{I}$ was therefore known independently of the training cohort. We tested every combination of $n \in \{8, 15, 23, 40, 80\}$, $p \in \{30, 120, 216\}$, and $r \in \{3, 5, 10\}$. In each cell, we constructed the stochastic-attention memory matrix from the training cohort and estimated a Gaussian generator from the same cohort; each produced 100 profiles. Stochastic attention used the same 95%-variance PCA, entropy-selected β^* , and generation settings as the clinical analysis. The Gaussian generator used the training mean and sample covariance, with only the diagonal jitter required for numerical sampling and no covariance shrinkage. We measured covariance recovery as $\|\hat{\boldsymbol{\Sigma}}_{\text{gen}} - \boldsymbol{\Sigma}_{\text{pop}}\|_F / \|\boldsymbol{\Sigma}_{\text{pop}}\|_F$. Figure 4A reports Gaussian error divided by stochastic-attention error for these simulated populations, averaged across the three latent ranks. A ratio above one indicated lower error for stochastic attention.

Second, we asked how performance changed as the clinical training set became smaller. We sampled without replacement from the 23 complete patients at $n \in \{20, 15, 10, 8\}$, using ten independently seeded subsets at each size. We constructed the stochastic-attention memory matrix from each subset, generated 100 profiles, and evaluated every generated cohort against the same full 23-patient reference rather than against its own training subset. The three outcomes were relative Frobenius error of the cross-visit correlation matrix, median marginal relative error pooled across visits and features, and TF-only mechanistic cloud overlap. Figure 4B shows the mean and standard deviation of the cross-visit error over the ten subsets.

Third, we asked how increasing curvature affected recovery of a nonlinear data pattern. We drew t uniformly from $[-1, 1]$ and defined the two S-curve coordinates as $x_1 = t$ and $x_2 = c \sin(\pi t)$. We mapped the curve into p dimensions with a random orthonormal embedding and added isotropic Gaussian noise with standard deviation 0.05. We used $n=23$ and $p=30$ or 120. We evaluated curvature $c=0, 0.25, 0.5, 1, 1.5$, and 2, with five independently seeded replicates per setting. Both generators produced 100 profiles per replicate. We measured the mean distance of generated profiles from the known curve within its two-dimensional embedding plane and the relative Frobenius error from the population covariance, estimated from a 50,000-sample reference draw. For stochastic attention, we marked the first nonzero curvature at which the mean distance from the curve reached twice its value for the straight-line case ($c=0$).

Ethics Statement

This study used de-identified secondary analysis measurements from samples collected longitudinally from women desiring a pregnancy. The primary prospective research study was approved by the Institutional Review Board at the University of Vermont. Written informed consent was obtained from all participants prior to enrollment.

Data Availability

The generated synthetic datasets and all intermediate validation outputs are available in the project repository at project GitHub repository. The real patient data were collected under IRB approval at the University of Vermont; de-identified data are available upon reasonable request to the corresponding author.

Code Availability

All code for synthetic patient generation, validation, figure reproduction, and mechanistic model calibration is available at project GitHub repository. The repository includes Julia source code, experiment scripts, and a Python baseline script for CTGAN/TVAE comparison.

Acknowledgments

This work was supported by NIH NHLBI R-33 HL 141787 (PIs Bernstein, Orfeo) and NIH NHLBI R01 HL 71944 (PI Bernstein).

Author Contributions

J.D.V. conceived the study, developed the SA pipeline and validation framework, performed computational analyses, and wrote the manuscript. M.C.B. and T.O. oversaw the secondary analysis measurements of coagulation and fibrinolysis markers, dynamic assays, and hormone measurements; provided domain expertise on coagulation biology and thrombin generation assays; and reviewed the manuscript. C.M. was involved in the enrollment of participants in the prospective research study and collection of clinical information of the participants. I.B. designed and oversaw the primary prospective research study of enrolled participants, provided clinical interpretation, and reviewed the manuscript. All authors approved the final version.

Competing Interests

The authors declare no competing interests.

References

- [1] Toy, S. U.S. maternal mortality hits highest level since 1965. *The Wall Street Journal* (2023). March 16, 2023.
- [2] Harris, E. US maternal mortality continues to worsen. *JAMA* **329**, 1248 (2023).
- [3] Chen, Y. *et al.* Pregnancy-related deaths in the US, 2018–2022. *JAMA Network Open* **8**, e254325 (2025).
- [4] Mann, K. G. Thrombin generation in hemorrhage control and vascular occlusion. *Circulation* **124**, 225–235 (2011).
- [5] Butenas, S., Orfeo, T. & Mann, K. G. Tissue factor in coagulation: Which? where? when? *Arteriosclerosis, Thrombosis, and Vascular Biology* **29**, 1989–1996 (2009).
- [6] Hockin, M. F., Jones, K. C., Everse, S. J. & Mann, K. G. A model for the stoichiometric regulation of blood coagulation. *Journal of Biological Chemistry* **277**, 18322–18333 (2002).
- [7] Bravo, M. C., Orfeo, T., Mann, K. G. & Everse, S. J. Modeling of human factor Va inactivation by activated protein C. *BMC Systems Biology* **6**, 45 (2012).

- [8] Hellgren, M. Hemostasis during normal pregnancy and puerperium. *Seminars in Thrombosis and Hemostasis* **29**, 125–130 (2003).
- [9] Grouzi, E., Pouliakis, A., Aktypi, A. *et al.* Pregnancy and thrombosis risk for women without a history of thrombotic events: a retrospective study of the real risks. *Thrombosis Journal* **20** (2022).
- [10] Bremme, K. A. Haemostatic changes in pregnancy. *Best Practice & Research Clinical Haematology* **16**, 153–168 (2003).
- [11] Azziz, R. *et al.* The prevalence and features of the polycystic ovary syndrome in an unselected population. *Journal of Clinical Endocrinology & Metabolism* **89**, 2745–2749 (2004).
- [12] Palomba, S., Santagni, S., Falbo, A. & La Sala, G. B. Complications and challenges associated with polycystic ovary syndrome: current perspectives. *International Journal of Women’s Health* **7**, 745–763 (2015).
- [13] Abalos, E. *et al.* Pre-eclampsia, eclampsia and adverse maternal and perinatal outcomes: a secondary analysis of the World Health Organization multicountry survey on maternal and newborn health. *BJOG: An International Journal of Obstetrics & Gynaecology* **121**, 14–24 (2014).
- [14] Dusse, L. M., Rios, D. R. A., Pinheiro, M. B., Cooper, A. J. & Lwaleed, B. A. Pre-eclampsia: relationship between coagulation, fibrinolysis and inflammation. *Clinica Chimica Acta* **412**, 17–21 (2011).
- [15] Luan, D., Zai, M. & Varner, J. D. Computationally derived points of fragility of a human cascade are consistent with current therapeutic strategies. *PLoS Computational Biology* **3**, e142 (2007).
- [16] Luan, D., Szlam, F., Tanaka, K. A., Barie, P. S. & Varner, J. D. Ensembles of uncertain mathematical models can identify network response to therapeutic interventions. *Molecular BioSystems* **6**, 2272–2286 (2010).
- [17] Brummel-Ziedins, K. E. *et al.* The prothrombotic phenotypes in familial protein c deficiency are differentiated by computational modeling of thrombin generation. *PLoS ONE* **7**, e44378 (2012).
- [18] Hastie, T., Tibshirani, R. & Friedman, J. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (Springer, 2009), 2nd edn.
- [19] Gonzales, A., Guruswamy, G. & Smith, S. R. Synthetic data in health care: A narrative review. *PLOS Digital Health* **2**, e0000082 (2023).
- [20] Kühnel, L. *et al.* Synthetic data generation for a longitudinal cohort study – evaluation, method extension and reproduction of published data analysis results. *Scientific Reports* **14**, 14412 (2024).
- [21] Ledoit, O. & Wolf, M. A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis* **88**, 365–411 (2004).
- [22] Goodfellow, I. *et al.* Generative adversarial nets. In *Advances in Neural Information Processing Systems*, vol. 27, 2672–2680 (2014).

- [23] Kingma, D. P. & Welling, M. Auto-encoding variational Bayes. In *International Conference on Learning Representations (ICLR)* (2014).
- [24] Xu, L., Skoularidou, M., Cuesta-Infante, A. & Veeramachaneni, K. Modeling tabular data using conditional GAN. In *Advances in Neural Information Processing Systems*, vol. 32, 7335–7345 (2019).
- [25] De Mitri, O., Wang, R. & Huber, M. F. Generative adversarial networks with limited data: A survey and benchmarking. *arXiv preprint arXiv:2504.05456* (2025).
- [26] Ramsauer, H. *et al.* Hopfield networks is all you need. In *International Conference on Learning Representations (ICLR)* (2021).
- [27] Varner, J. D. Conditioning protein generation via Hopfield pattern multiplicity. *arXiv preprint arXiv:2603.20115* (2026).
- [28] Bravo, M. *et al.* BSTModelKit.jl: Building biochemical systems theory models of individualized coagulation dynamics during pregnancy. JuliaCon 2023, <https://www.youtube.com/watch?v=IGtUqMLHBvQ> (2023). Conference proceedings talk.
- [29] Ghetmiri, D. E., Cohen, M. J. & Menezes, A. A. Personalized modulation of coagulation factors using a thrombin dynamics model to treat trauma-induced coagulopathy. *npj Systems Biology and Applications* **7**, 44 (2021).
- [30] Varner, J. D. Training-free generation of protein sequences from small family alignments via stochastic attention. *arXiv preprint arXiv:2603.14717* (2026).
- [31] Alswaidan, A. & Varner, J. D. Stochastic attention via Langevin dynamics on the modern Hopfield energy. *arXiv preprint arXiv:2603.06875* (2026).
- [32] Mould, D. R. & Upton, R. N. Basic concepts in population modeling, simulation, and model-based drug development—part 2: introduction to pharmacokinetic modeling methods. *CPT: Pharmacometrics & Systems Pharmacology* **2**, e38 (2013).
- [33] Chase, J. G. *et al.* Physiological modeling, tight glycemic control, and the ICU clinician: what are models and how can they affect practice? *Annals of Intensive Care* **1**, 11 (2011).
- [34] Ribba, B. *et al.* A tumor growth inhibition model for low-grade glioma treated with chemotherapy or radiotherapy. *Clinical Cancer Research* **18**, 5071–5080 (2012).
- [35] Yizhak, K., Chaneton, B., Gottlieb, E. & Ruppin, E. Modeling cancer metabolism on a genome scale. *Molecular Systems Biology* **11**, 817 (2015).
- [36] Shan, M., Dai, D., Vudem, A., Varner, J. D. & Stroock, A. D. Multi-scale computational study of the Warburg effect, reverse Warburg effect and glutamine addiction in solid tumors. *PLoS Computational Biology* **14**, e1006584 (2018).
- [37] McLean, K. C., Bernstein, I. M. & Brummel-Ziedins, K. Tissue factor dependent thrombin generation across pregnancy. *American Journal of Obstetrics and Gynecology* **207**, 135.e1–135.e6 (2012).
- [38] Hale, S. A. *et al.* Coagulation and fibrinolytic system protein profiles in women with normal pregnancies and pregnancies complicated by hypertension. *Pregnancy Hypertension* **2**, 152–157 (2012).

- 921 [39] Psinos, R. B. C., Morris, E. A., McBride, C. A. & Bernstein, I. M. Association of pre-
922 pregnancy subclinical insulin resistance with cardiac dysfunction in healthy nulliparous women.
923 *Pregnancy Hypertension* **26**, 11–16 (2021).
- 924 [40] Bernstein, I. M., Hale, S. A. & Badger, G. J. Differences in cardiovascular function comparing
925 prior preeclampsia with nulliparous controls. *Pregnancy Hypertension* **6**, 320–326 (2016).

Table 1: **Demographics and clinical characteristics of the 23 patients with complete longitudinal data across all three visits.** Values are reported as mean $\pm$ SD, median (range), or n (%) as appropriate.

Characteristic	Mean $\pm$ SD	Median (Range)	n (%)
Age at enrollment (yrs)	30.2 $\pm$ 5.1	29 (20–41)	
Race			
White			22 (96)
Not disclosed			1 (4)
Prepregnancy BMI	26.6 $\pm$ 5.3	25 (19.0–37.8)	
Parity		0 (0–2)	
Nulliparous			16 (70)
Parous			7 (30)
Enrollment condition			
Healthy nulliparous			14 (61)
PCOS			3 (13)
Prior PE			6 (26)
Cycle day at V1	9.2 $\pm$ 3.6	10 (3–14)	
Gestational age at V2 (days)	89.0 $\pm$ 5.5		
Gestational age at V3 (days)	217.0 $\pm$ 4.0		
<i>Enrollment cohort $\times$ pregnancy outcome</i>			
	Uncomplicated	Developed PE	Total
Healthy nulliparous	12	2	14
Prior PE	4	2	6
PCOS	2	1	3
Total	18	5	23

Table 2: Per-visit mean relative error (MRE) between real and SA-generated synthetic patients for representative coagulation features, computed as $|\bar{x}_{\text{synth}} - \bar{x}_{\text{real}}|/|\bar{x}_{\text{real}}|$.

Feature	Visit 1	Visit 2	Visit 3
Factor II (%)	0.018	0.017	0.007
Factor VIII (%)	0.018	0.008	0.005
AT (%)	0.025	0.001	0.008
Fibrinogen (mg/dL)	0.005	0.009	<0.001
TF Peak (nM)	0.025	0.002	0.017
TF ETP (nM·min)	0.013	0.008	0.023




Figure 1: **Physiological correlations are preserved in synthetic patients.** Scatter plots of coagulation factor levels versus thrombin generation parameters across all three visits. Filled markers are real patients; open markers are SA-generated synthetic patients. Colors encode visit: blue (V1/baseline), green (V2/first trimester), orange (V3/third trimester). Real and synthetic patients occupy the same joint regions across all four pairwise relationships, confirming that biologically meaningful dependencies are preserved: AT inversely correlates with TF-initiated peak thrombin (top left), FVIII positively correlates with peak (top right), and both FII (bottom left) and AT (bottom right) show the expected relationships with ETP. The visit-dependent shift in these relationships, reflecting the pregnancy-driven increase in procoagulant factors, is also preserved in the synthetic cohort.




Figure 2: **Pregnancy-driven longitudinal trajectories are reproduced.** Mean $\pm$ standard deviation bands for six key coagulation features across visits (BL = baseline, 1st = first trimester, 3rd = third trimester) for real (blue) and SA-generated synthetic (orange) patients. The characteristic pregnancy-driven increases in fibrinogen, Factor VIII, and vWF are captured, as are the stable-to-declining patterns in Factor II and AT.




Figure 3: **Cross-visit correlation structure.** Top row: full 216×216 Pearson correlation matrices for Real ($K=23$, left), SA ($N=100$, center), and MVN ($N=100$, right) populations. Each matrix is organized as a 3×3 grid of 72×72 blocks (delineated by black lines), where the diagonal blocks capture within-visit feature correlations (V1–V1, V2–V2, V3–V3) and the off-diagonal blocks capture cross-visit dependencies (e.g., V1–V3: does a patient’s baseline Factor VIII predict their third-trimester Factor VIII?). SA preserves both the within-visit and cross-visit block structure. Bottom row: pairwise residual matrices (element-wise subtraction), all plotted on the same ± 1 color scale as the top row. The SA–Real residual shows small, unstructured deviations. The MVN–Real residual is smoother, particularly in the off-diagonal blocks, reflecting the Ledoit–Wolf regularization that systematically shrinks cross-visit correlations toward zero. The SA–MVN residual confirms that the two methods differ most in these off-diagonal blocks, where SA retains longitudinal dependencies that MVN’s rank-deficient covariance estimate cannot represent. A version with amplified residual color scale (± 0.5) and Frobenius norms is provided in Fig. S1.




Figure 4: Population recovery and clinical-cohort subsampling. (A) For each low-rank Gaussian setting, the stochastic-attention memory matrix was constructed and a sample-covariance Gaussian generator was estimated from the same simulated training cohort; each generated 100 profiles. Color represented the ratio of their relative Frobenius errors against the known population covariance (Gaussian error divided by stochastic-attention error), averaged over latent ranks 3, 5, and 10. Values above one (blue) indicated lower error for stochastic attention. Because the target covariance was fixed before the training cohort was drawn, this measured population recovery rather than reproduction of the training data. (B) A stochastic-attention memory matrix was constructed from each of ten random subsets of the clinical cohort at $n=20, 15, 10$, and 8 , and each matrix generated 100 profiles. Points and bars represented the mean and standard deviation of the relative cross-visit correlation error against the same full 23-patient reference; the star represented generation using all 23 patients.




Figure 5: Condition-specific feature preservation. Grouped bar charts comparing real (dark bars) and SA-generated synthetic (light bars) patient means ($\pm$ SD error bars) for eight coagulation features, shown separately for each clinical subgroup: Uncomplicated ($n=18$, left, blue), PCOS ($n=3$, center, orange), and Developed PE ($n=5$, right, red). Each bar is the mean of that assay pooled across all three visits (V1–V3); factor activities (FVIII, vWF, FIX, FV, α_2 -AP, FVII, plasminogen) are expressed as % of normal pooled reference plasma, and fibrinogen (Fbgn) in mg/dL. Gray percentages indicate the mean relative error (MRE). Green annotations show a descriptive bootstrap Mann–Whitney diagnostic: the fraction of 1,000 bootstrap replicates (synthetic subsampled to n_{real}) in which $p > 0.05$, meaning that this test detected no difference at the available sample size. This is not an equivalence test. Features achieving $\geq 90\%$ are shown in green; $< 90\%$ in red. Across all 24 feature–condition pairs, 20 (83%) had no difference detected in $\geq 90\%$ of replicates (median: 98.6%). The four pairs below this threshold were FIX and plasminogen in the PCOS and Developed PE subgroups, where small subgroup sizes and large inter-patient variability reduced test power.




Figure 6: **Mechanistic validation (TF-only).** Top row: BZ2012 ODE-predicted TGA values (vertical axis) versus dataset TGA values (horizontal axis) for five thrombin generation parameters. The dataset values are the TGA measurements from each patient’s record, present for both real and synthetic patients. The ODE-predicted values are computed by running each patient’s coagulation factor levels through the 58-species BZ2012 model. The dashed line indicates where the ODE model would perfectly reproduce the dataset value ($y=x$); systematic deviations from this line (e.g., lagtime underprediction) reflect ODE model bias, not synthetic data failure. Filled markers are real patients, open markers are SA-generated synthetic patients, colored by visit: blue (V1/baseline, training set), green (V2/first trimester), orange (V3/third trimester). Real and synthetic patients occupy the same cloud with the same pattern of ODE model bias. Bottom row: ODE-predicted/dataset ratio distributions for real (blue) and synthetic (orange) patients. If SA had generated biologically implausible factor combinations, the ODE model would process them differently, producing divergent ratio distributions. Instead, the distributions overlap substantially. Each panel is annotated with cloud overlap (OL), the two-sample Kolmogorov–Smirnov statistic (D), and p -value. All five features show $p > 0.35$, confirming that the ODE model cannot distinguish real from synthetic patients.




Figure 7: Downstream utility: synth-calibrated vs. real-calibrated mechanistic model predictions on held-out real patients. Each panel compares the BZ2012 ODE predictions for one TGA feature when the model is calibrated on real V1 patients (horizontal axis) versus synthetic V1 patients (vertical axis), evaluated on the same held-out real V2 and V3 patients. Points near the $y=x$ line indicate that the two calibrations produce equivalent predictions. Green: V2 (first trimester); orange: V3 (third trimester). The synth-calibrated model matched or outperformed the real-calibrated model on all five features (overall ratio $0.94\times$), demonstrating that a mechanistic model trained entirely on synthetic data generalizes to real patients as well as one trained on real data. Both calibrations used 11 random restarts; all starts converged.

926 **Supplementary Information**

927 **Supplementary Tables**

Table S1: **Inverse temperature sensitivity analysis.** Generation under stochastic attention quality as a function of β/β^* , where $\beta^* \approx 2.94$ was selected at the attention-entropy transition in the 18-dimensional PCA space ($K=23$ stored patterns). Operating at $\beta = \beta^*$ (bold row) balances generation novelty against fidelity. Med MRE = median relative error of per-feature means; Novelty = $1 - \max_k \cos(\hat{\xi}, \hat{m}_k)$.

β/β^*	β	Noise scale	PC1 std ratio	Novelty	Med MRE (%)
0.1	0.29	0.261	0.538	0.542	0.7
0.3	0.88	0.151	0.550	0.536	0.7
0.5	1.47	0.117	0.562	0.528	0.7
1.0	2.94	0.082	0.589	0.502	0.6
1.5	4.41	0.067	0.613	0.474	0.7
2.0	5.88	0.058	0.638	0.446	0.7
3.0	8.82	0.048	0.689	0.403	0.8

Table S2: **Multiplicity weighting parameters for conditional generation.** For the unit-normalized memories used here, the multiplicity ratio ρ was computed to place a target fraction $f_{\text{target}} = 0.80$ of finite-mixture component probability on the designated patterns. K_{eff} is the participation ratio of those component probabilities and describes their concentration across repeated draws.

Condition	K_{sub}	K_{bg}	ρ	f_{eff}	K_{eff}
Uncomplicated	18	5	1.11	0.80	23.0
PCOS	3	20	26.67	0.80	4.6
Developed PE	5	18	14.40	0.80	7.7

Table S3: **SA generation hyperparameters.** The same values were used for unconditioned and conditioned generation except that β^* was recomputed with the multiplicity bias for conditioned cohorts.

Parameter	Value	Description
α	0.01	Langevin step size (see Table S8)
T	2,000	Langevin iterations per sample
β^*	2.94	Entropy-transition point (unconditioned)
PCA threshold	95%	Variance retained
d_{PCA}	18	PCA dimensionality
N	100	Synthetic patients per cohort
Random seed	42	For reproducibility
f_{target}	0.80	Designated-component probability (conditioned)

Table S4: **BZ2012 ODE model calibration.** Five rate constants were calibrated on Visit 1 real patients ($n=23$) using bounded Nelder–Mead optimization in log-scale-factor space. The cost function was the mean normalized relative squared error across five TGA features under TF-only conditions. Bounds: $\pm \log(100)$. Convergence: $f_{\text{tol}} = 10^{-6}$, 500 iterations, 11 random restarts.

Rate constant	Role	Default	Calibrated	Scale
prothrombinase k_{cat}	Thrombin burst	63.5 s^{-1}	21.94 s^{-1}	$0.35 \times$
intrinsic Xase k_{cat}	Amplification Xa	8.2 s^{-1}	0.169 s^{-1}	$0.021 \times$
extrinsic Xase k_{cat}	Initiation Xa	6.0 s^{-1}	93.2 s^{-1}	$15.5 \times$
PC activation k_{cat}	Protein C activation	0.41 s^{-1}	0.890 s^{-1}	$2.17 \times$
mIIa conversion k	mIIa $\rightarrow$ IIa	2.3×10^8	1.15×10^9	$5.01 \times$

Units for mIIa conversion: $\text{M}^{-1}\text{s}^{-1}$

Table S5: **Summary of mean relative errors across all 72 features.** Per-visit distribution of MREs between real and SA-generated synthetic patient means. The full feature-by-feature table ($72 \text{ features} \times 3 \text{ visits} = 216 \text{ entries}$) is provided in the code repository as `paper_summary_statistics.csv`.

Visit	Median	Mean	25th pct.	75th pct.	Maximum	MRE below 5%
V1 (baseline)	0.022	0.029	0.008	0.039	0.158	59/72 (81.9%)
V2 (1st tri)	0.009	0.021	0.004	0.022	0.224	65/72 (90.3%)
V3 (3rd tri)	0.011	0.016	0.005	0.024	0.062	69/72 (95.8%)
Pooled (all 3 visits)	0.012	0.022	0.005	0.028	0.224	193/216 (89.4%)

Pooled median MRE 95% bootstrap CI: (0.98%, 1.60%) over 10,000 resamples (seed 42).

Table S6: **Comparison with deep generative baselines.** CTGAN and TVAE were trained on the same 90 per-visit records (23 patients $\times$ 3 visits, 72 features) at three epoch counts. SA and MVN results from the main text are shown for reference. CTGAN fails at this sample size (median MRE $\approx$ 19%), consistent with GAN mode collapse on small datasets. TVAE achieves comparable marginal fidelity at high epoch counts but cannot capture cross-visit covariance or perform conditional generation.

Method	Epochs	Med MRE	Med KS	Corr MAE
SA	–	0.019	0.169	0.200
MVN	–	0.024	0.161	0.090
CTGAN	300	0.189	0.364	0.261
CTGAN	1,000	0.195	0.349	0.265
CTGAN	3,000	0.187	0.341	0.180
TVAE	300	0.037	0.174	0.152
TVAE	1,000	0.026	0.248	0.082
TVAE	3,000	0.018	0.224	0.092

Table S7: **PCA variance threshold sensitivity analysis.** SA generation quality is stable across thresholds from 85% to 99%. Median MRE remains between 1.0% and 1.7% across all thresholds. The 95% threshold used in the main text (bold) is not a critical design choice.

Threshold	d_{PCA}	K/d	β^*	Novelty	Med MRE	Corr MAE
85%	13	1.77	2.94	0.420	0.0098	0.124
90%	15	1.53	2.94	0.477	0.0160	0.126
95%	18	1.28	2.94	0.500	0.0124	0.143
97.5%	20	1.15	2.94	0.543	0.0120	0.135
99%	21	1.10	2.94	0.541	0.0172	0.138

Table S8: **ULA vs MALA sampling diagnostics.** Ten chains of 5,000 iterations each were run with both the Unadjusted Langevin Algorithm (ULA) and the Metropolis-Adjusted Langevin Algorithm (MALA) at step size $\alpha = 0.01$ and $\beta^* = 2.94$ on the 18-dimensional PCA memory ($K=23$ patterns). The MALA acceptance rate of 99.9% confirms that virtually every ULA proposal satisfies detailed balance, and the indistinguishable energy and effective sample size statistics confirm negligible discretization bias. [This comparison bounded the discretization error of the reported sampler. The direct reference for the target itself was the analytic sampler of Eq. \(3\), introduced below.](#)

Metric	ULA	MALA
Acceptance rate (%)	– (no rejection)	99.9 ± 0.03
Mean energy (post-burn-in)	1.932 ± 0.141	1.886 ± 0.231
τ_{int}	109.5 ± 49.5	106.7 ± 30.4
Effective sample size	41.8 ± 14.2	40.0 ± 10.2

Burn-in: 1,000 iterations discarded. τ_{int} : integrated autocorrelation time estimated with a windowed estimator (threshold 0.05). $\text{ESS} = (T - T_{\text{burn}})/\tau_{\text{int}}$.

Table S9: **Memorization diagnostics.** Two complementary checks that SA-generated patients are not memorized copies of the $K=23$ stored profiles. Novelty score is the angular distance between a synthetic sample and its nearest stored pattern at $\beta = \beta^*$. Nearest-neighbor distances are computed in the per-visit standardized 72-dimensional feature space.

Metric	Value
Mean novelty score, $1 - \max_k \cos(\hat{\xi}, \hat{m}_k)$	0.50
Median synth–real nearest-neighbor distance (std. units)	6.14
Median real–real nearest-neighbor distance (std. units)	6.87
Distance ratio (synth–real / real–real)	0.89

Table S10: **Descriptive bootstrap Mann–Whitney diagnostic per feature and condition.** For each of 24 feature–condition pairs, we subsampled the synthetic cohort to match the real sample size, applied a Mann–Whitney U test, and repeated 1,000 times. Frac. $p > 0.05$ is the fraction of replicates in which the test detected no difference at the available sample size. This power-limited diagnostic is not an equivalence test. 20 of 24 pairs achieve $p > 0.05$ in $\geq 90\%$ of replicates; the four failing pairs (italicized) all involve high-variance features in the smallest subgroups.

Condition	Feature	MRE	Frac. $p > 0.05$
Uncomplicated ($n=18$)	FVIII	0.057	0.952
Uncomplicated ($n=18$)	vWF	0.022	0.999
Uncomplicated ($n=18$)	FIX	0.018	0.998
Uncomplicated ($n=18$)	FV	0.006	0.996
Uncomplicated ($n=18$)	α_2 -AP	0.019	1.000
Uncomplicated ($n=18$)	FVII	0.057	0.941
Uncomplicated ($n=18$)	Fbgn	0.017	0.987
Uncomplicated ($n=18$)	Plgn	0.020	0.988
PCOS ($n=3$)	FVIII	0.109	0.985
PCOS ($n=3$)	vWF	0.150	0.959
PCOS ($n=3$)	<i>FIX</i>	<i>0.121</i>	<i>0.873</i>
PCOS ($n=3$)	FV	0.057	0.944
PCOS ($n=3$)	α_2 -AP	0.051	0.987
PCOS ($n=3$)	FVII	0.029	0.999
PCOS ($n=3$)	Fbgn	0.043	0.992
PCOS ($n=3$)	<i>Plgn</i>	<i>0.153</i>	<i>0.615</i>
Developed PE ($n=5$)	FVIII	0.097	0.932
Developed PE ($n=5$)	vWF	0.010	0.986
Developed PE ($n=5$)	<i>FIX</i>	<i>0.108</i>	<i>0.780</i>
Developed PE ($n=5$)	FV	0.006	0.987
Developed PE ($n=5$)	α_2 -AP	0.062	0.950
Developed PE ($n=5$)	FVII	0.045	0.993
Developed PE ($n=5$)	Fbgn	0.048	0.996
Developed PE ($n=5$)	<i>Plgn</i>	<i>0.093</i>	<i>0.814</i>

Median Frac. $p > 0.05$ across all 24 pairs: 0.986. Failing pairs (< 0.90 , italicized): PCOS–FIX, PCOS–Plgn, DPE–FIX, DPE–Plgn.

Table S11: **Mechanistic validation diagnostics per feature and condition.** Cloud overlap is the fraction of synthetic ODE-predicted/measured ratios falling within the 5th–95th percentile of the corresponding real distribution. KS D and KS p are from a two-sample Kolmogorov–Smirnov test on the same ratio distributions. Real $n=69$ (23 patients $\times$ 3 visits); synthetic $n=300$ (100 patients $\times$ 3 visits).

Condition	TGA Feature	Cloud overlap	KS D	KS p
TF-only	Lagtime	0.92	0.112	0.48
TF-only	Peak	0.91	0.081	0.86
TF-only	T.Peak	0.92	0.102	0.61
TF-only	Max Rate	0.91	0.083	0.84
TF-only	ETP	0.83	0.123	0.36
TF+TM	Lagtime	0.92	0.127	0.32
TF+TM	Peak	0.88	0.079	0.88
TF+TM	T.Peak	0.91	0.110	0.51
TF+TM	Max Rate	0.88	0.096	0.67
TF+TM	ETP	0.89	0.112	0.48

Under TF-only conditions, the cloud overlap range was 0.83–0.92. Under TF+TM conditions, the range was 0.88–0.92. Pooled across both conditions and all five features, the overlap was 0.902. KS D range across both conditions: 0.079–0.127, all $p > 0.30$.

Table S12: **Downstream utility — per-feature held-out V2/V3 performance.** Median relative error of the BZ2012 ODE-predicted TGA values against the held-out real V2 and V3 measurements, for the model calibrated on real V1 patients ($K=23$) and the model calibrated on synthetic V1 patients ($N=100$). The synth-calibrated model achieves slightly lower error on every feature. Synth/Real ratio < 1 favors the synth-calibrated model.

TGA Feature	Real-cal MRE	Synth-cal MRE	Synth/Real
Lagtime	0.165	0.162	0.98 $\times$
Peak	0.097	0.087	0.90 $\times$
T.Peak	0.330	0.306	0.93 $\times$
Max Rate	0.286	0.259	0.91 $\times$
ETP	0.197	0.183	0.93 $\times$
Overall median	0.201	0.189	0.94$\times$

Table S13: **PCA-space ablation.** Five generators operating in the identical $d=18$ PCA subspace and evaluated on the same harness. Median relative error (MRE) measures marginal fidelity. Mechanistic cloud overlap was the fraction of predicted-to-recorded TGA ratios for synthetic patients within the real-data 5th–95th percentile range after pooling across patients, visits, and TGA features under TF-only conditions. Median novelty is $1 - \max_k \cos(\hat{\xi}, \hat{\mathbf{m}}_k)$, and Frac. novel is the fraction above the operational 0.2 cutoff, equivalent to cosine similarity below 0.8 to every stored pattern. The cross-visit Frobenius error is the deviation of the synthetic correlation matrix from the real one. A lower value indicates closer in-sample reproduction and can therefore reward memorization. The comparison showed that the shared representation accounted for much of the marginal and mechanistic fidelity but did not guarantee performance.

Method	MRE (%)	Cross-visit Frob.	Mech. overlap	Median novelty	Frac. novel
SA	1.24	0.641	0.902	0.514	1.00
PCA+GMM	1.13	0.422	0.918	0.497	0.94
PCA+KDE	1.94	0.337	0.823	0.249	0.65
PCA+kNN	1.63	0.587	0.938	0.051	0.05
PCA+Copula	1.34	0.417	0.934	0.462	1.00

Table S14: **Convex-hull membership did not discriminate at $K=23$ in $d=18$.** Hull membership and distance for the 100 synthetic patients against the 23 real patients, and for each real patient against the other 22. The radial factor t^* is where the ray from the hull centroid through a point exits the hull, so $1/t^*$ is the factor by which a point overshoots the boundary and $t^* \geq 1$ would mean the point is inside. Because binary hull membership did not show how far a point lay beyond the boundary, we also compared distance to the hull. The real-patient calculation omitted the patient being evaluated. To make the synthetic-patient calculation equally stringent, we omitted each synthetic patient’s nearest real patient, leaving 22 real points in both comparisons. A descriptive Mann–Whitney test did not detect a difference in distance between synthetic and real patients ($p = 0.07$), while radial overshoot was smaller for the synthetic patients ($p < 10^{-4}$). Both tests treated the 100 synthetic patients as independent, which they were not, so the reported significance was optimistic. With only 23 patients in 18 dimensions, every real patient formed an outer corner and the hull occupied little of the surrounding space. The generator also normalized the stored patterns and generated directions to length one. This meant that a generated direction could remain inside the hull only by exactly matching a stored pattern.

The biological constraint included non-negativity of both hormones. Negative values occurred for estradiol in 19 patients and progesterone in 37 patients, with 10 overlapping; most were close to zero relative to the real-cohort scale, but the minima were -2109.8 and -75.8 , respectively. Positivity was not enforced for any variable. Relative to the pooled real-cohort distributions, zero was less than one standard deviation below the mean for both hormones but at least 2.6 standard deviations below the mean for each coagulation factor. This difference explained why only the hormones crossed zero in this sample; it did not guarantee positivity of coagulation factors in other draws. All nine coagulation-factor inputs to the BZ2012 model were positive, so all 100 synthetic patients were retained for the mechanistic check. Gaussian sampling in PCA space followed by linear decoding had unbounded support, and no clipping was applied. The unconstrained concatenated MVN baseline produced 58 negative hormone measurements across 43 of 100 patients, compared with 57 measurements across 46 of 100 for SA.

Quantity	Synthetic ($n=100$)	Real, leave-one-out ($n=23$)
Fraction inside hull	0.00	0.00
Median distance to hull	11.0	11.9
Median radial overshoot $1/t^*$	10.5	15.2
<i>Hull geometry of the real cohort</i>		
Median patient radius		13.8
Median hull extent, generic direction		2.0
Largest t^* attained on the unit sphere		0.231
<i>Plausibility of the out-of-hull synthetic patients</i>		
Biological constraints satisfied		54/100
Mechanistic band, per feature		83 to 92%
Mechanistic band, all features and visits		41/100
Both biological and strict mechanistic criteria		24/100

Table S15: **Membership inference recovered memory membership, not unseen patients.** Median nearest-neighbor distances in the standardized 216-dimensional space, and the test’s discrimination. For each split, we constructed the memory matrix from 15 patients and scored each real record by its distance to the resulting synthetic set, treating the 15 included patients as members and the 8 held-out patients as non-members. Every quantity came from the same protocol. We drew 40 partitions at random with the generation seed held fixed, so that only the partition varied, and pooled the distances across all partitions. Non-members sat at the same distance from the synthetic cohort as unrelated real patients sat from one another. The synthetic cohort carried no proximity signal about patients the generator never saw, and the discrimination came entirely from members sitting closer. The AUC measured how often the test ranked a member above a non-member; 0.5 indicated no discrimination and 1.0 indicated perfect discrimination.

Quantity	Value
<i>Median nearest-neighbor distance</i>	
Memory-matrix member to nearest synthetic	10.7
Held-out non-member to nearest synthetic	15.4
Real to nearest real, leave-one-out	15.9
Synthetic to nearest real	13.8
<i>Membership-inference AUC over 40 random splits</i>	
Mean	0.97
Standard deviation	0.02
Median	0.98
Minimum	0.90
Maximum	1.00
Splits reaching perfect separation	9/40
Splits with AUC above 0.9	39/40




Figure S1: **Cross-visit correlation structure (amplified residual scale)**. Same layout as the main-text correlation figure, but with the bottom-row residual panels plotted on a ± 0.5 color scale (vs. ± 1.0 in the main text) to reveal fine structure. Frobenius norms of each residual matrix are shown in the panel titles.

929 **Supplementary Methods: PCA-space ablation**

930 We compared the canonical stochastic-attention cohort with four baseline generators operating
 931 on the same unnormalized 18-dimensional PCA coordinates. Each baseline generated 100 profiles
 932 with a separate random-number stream initialized using seed 42. The Gaussian mixture used two
 933 components with diagonal covariance matrices, 100 expectation-maximization iterations, and a
 934 variance floor of 10^{-3} . The kernel-density sampler chose a stored PCA point uniformly and added
 935 Gaussian noise with the empirical PCA covariance and Scott bandwidth $h = K^{-1/(d+4)}$; $10^{-8}\mathbf{I}$
 936 was added for numerical factorization. The nearest-neighbor sampler chose a stored point, selected
 937 one of its three nearest neighbors uniformly, and interpolated between them with a coefficient
 938 drawn from $\text{Uniform}(0, 1)$. The Gaussian copula transformed empirical marginal ranks to normal
 939 scores, estimated their correlation matrix with $10^{-6}\mathbf{I}$ added for factorization, and mapped sampled
 940 scores back through the empirical marginal quantiles. All samples were passed through the same
 941 inverse-PCA and destandardization pipeline and evaluated with the same marginal, cross-visit,
 942 mechanistic, and novelty calculations.




Figure S2: **Covariance eigenvalue spectrum.** Log-scale eigenvalues of the 216×216 sample covariance for Real ($K=23$, rank 22), SA ($N=100$, rank 18), and MVN ($N=100$) populations. Dashed vertical lines mark the SA PCA truncation point (component 18) and the data rank limit (component 22). MVN’s Ledoit–Wolf regularization inflates eigenvalues beyond component 22 (green shaded region), introducing spurious variance in 194 dimensions where the data contain no signal.

Supplementary Methods: selecting the inverse-temperature operating point

The operating point β^* was located from the attention entropy. For a candidate β we computed the Shannon entropy of the attention weights $\text{softmax}(\beta \hat{\mathbf{X}}^\top \boldsymbol{\xi})$ at 20 probe points placed at stored patterns, averaged those entropies, and normalized by $\log K$. The normalized curve runs from 0 to 1. A value of 0 means one pattern carries all the weight, and a value of 1 means every pattern carries equal weight. Sweeping β logarithmically over 80 values from 0.1 to 1000 traces this curve from the diffuse regime into the concentrated one. We selected β^* at the point of greatest downward curvature, calculated as the most negative finite-difference second derivative of normalized entropy with respect to $\log \beta$. This marked the transition from attention shared across patterns toward attention localized on individual patterns. For multiplicity-weighted sampling we repeated the same computation with the $\log r_k$ bias included in the softmax logits. This gives a ρ -dependent $\beta^*(\rho)$, because the weighting changes how quickly attention concentrates as β grows. The random-pattern scale $\sqrt{d}$ served as the theoretical reference quoted in the main text.




Figure S3: **Conditioned generation: PCA projection (pooled)**. Large markers: real patients (blue = Uncomplicated, $n=18$; orange = PCOS, $n=3$; red = Developed PE, $n=5$). Small markers: synthetic patients ($N=100$ per condition). With only 3 PCOS and 5 Developed PE patients, the subgroups do not form clearly separable clusters in the first two principal components, but the synthetic cohorts concentrate around their respective real counterparts.




Figure S4: **Conditioned generation: PCA projection by visit**. Same as Fig. S3 but separated by visit (V1, V2, V3). Condition-specific clustering is maintained within each individual visit.




Figure S5: **Mechanistic validation (TF+TM condition)**. Same layout as main-text Fig. 6 but under TF+TM conditions. The model systematically underpredicts peak and maximum rate ($\approx 0.5\times$ measured) for both real and synthetic patients, reflecting overcorrection of the protein C anticoagulant pathway. Cloud overlap remains high at 88–92%, confirming that the model processes real and synthetic patients identically even under imperfect calibration conditions.




Figure S6: **Cross-visit calibration generalization**. Boxplots of the ODE-predicted/dataset ratio by visit for real patients. The model was calibrated on Visit 1 only. Top: TF-only conditions, where ratios near 1.0 at Visits 2 and 3 indicate that the calibration generalizes across pregnancy timepoints. Bottom: TF+TM conditions, where the calibration does not generalize as well, particularly for peak and maximum rate.




Figure S7: **Spearman rank correlations between ODE-predicted and dataset TGA features.** Left: TF-only. Right: TF+TM. ETP is the best-predicted feature ($\rho \approx 0.6$ – 0.8 for real patients under TF-only). Weak rank correlations for other features are expected given the 5-parameter population-level calibration.




Figure S8: **PCA projections by visit: SA vs. MVN.** Each panel shows the first two principal components (PC1: 27.7% variance, PC2: 15.4%) computed from standardized per-visit real patient data ($K=23$ patients per visit, 72 features). Dark markers indicate real patients; lighter markers indicate synthetic patients ($N=100$). Within each visit (column), the SA panel (top) and MVN panel (bottom) share identical axis ranges, so the relative dispersion of the two methods can be read against the same real data cloud. Top row: SA-generated patients (circles) cluster tightly around the real data cloud at all three visits. Bottom row: MVN-generated patients (diamonds) show greater dispersion, particularly at Visits 2 (first trimester) and 3 (third trimester), where MVN patients extend beyond the convex hull of the real data. This excess scatter reflects the spurious variance introduced by Ledoit–Wolf regularization of the rank-deficient covariance (Fig. S2). Colors encode visit: blue (V1, not pregnant), green (V2, first trimester), orange (V3, third trimester).

Sampling correctness and privacy diagnostics

We assessed a narrowly defined membership-inference question: if someone already possessed a candidate de-identified assay profile, could the synthetic cohort reveal whether that record had been included in the memory matrix? The source cohort contained no direct identifiers or link to named individuals, so this was not a test of personal identification. For every synthetic patient we computed the distance to the closest real record (DCR) in the standardized 216-dimensional space and compared the synthetic-to-real distribution against the leave-one-out real-to-real distribution. Synthetic patients sat slightly closer to the real cohort than the real patients sat to one another, with median DCR 13.8 against 15.9. We also ran a nearest-neighbor membership-inference test, which scored each real record by its proximity to the synthetic set over repeated random train/holdout splits. It separated the 15 memory-matrix members from the 8 held-out patients with a mean AUC of 0.97 (standard deviation 0.02 over 40 splits; 0.5 indicated no discrimination and 1.0 perfect discrimination). At the published operating point the generator therefore retained enough record-specific proximity for memory-matrix membership to be detected.

To attribute this behavior we asked whether it reflected the target distribution $\pi(\xi) \propto \exp(-\beta E)$ or merely the sampler used to draw from it. A mixing diagnostic showed that at β^* the target was multimodal, with one component centered on each of the $K=23$ stored patterns. It also showed that discretization bias in the Langevin updates was negligible, with Metropolis acceptance near 0.999 (Supplementary Fig. S10). We then constructed a four-rung ladder (Table S16, Fig. S9) that held β , the memory, and the decoding fixed and varied only the sampler. The rungs ran from the published anchor-initialized single-endpoint scheme, to sphere initialization, to burn-in-controlled thinned pooling, to Metropolis-adjusted sampling. Across the ladder the membership-inference AUC stayed near 0.96 to 0.97. The median DCR rose only from 14.0 to 14.3 and never reached the real-to-real baseline of 15.9, at a modest cost in marginal fidelity. Proper sampling of π did not remove the memorization. At $K=23$ the memorization was a property of the distribution.

The remaining lever was the inverse temperature. Lowering β spread probability mass away from the stored patterns and was expected to improve privacy, at the expense of the memorization that also underlay fidelity. We swept β across a 40-fold range around β^* (Fig. S9C). The sweep confirmed the trade-off and bounded it sharply. The median DCR rose monotonically as β fell, but it saturated near 14.3 and stayed below the real-to-real baseline of 15.9 even at one twentieth of β^* . Over the same range the cross-visit covariance error grew from 0.59 to 0.69. Marginal and mechanistic fidelity were insensitive to β , so the longitudinal covariance structure bore the cost of lowering it. No inverse temperature yielded a cohort that was both as private as the real patients were from one another and faithful to their joint structure.

Neither lever settled the question completely, because every scheme compared was a finite-time Markov chain. The target itself was available in closed form. Completing the square in Eq. (1) gives Eq. (3). For unit-normalized memories π_r is exactly the isotropic Gaussian mixture $\sum_k q_k \mathcal{N}(\mathbf{m}_k, \beta^{-1} \mathbf{I})$ with $q_k = r_k / \sum_j r_j$. It can therefore be sampled ancestrally, by drawing a component and adding isotropic noise. There is no chain, no initialization, no burn-in and no discretization error. We generated a protocol-matched exact cohort ($N=100$ at β^* , reusing the ladder’s seeded magnitude draw so that only the direction sampler differs) and evaluated it on the identical harness (Table S16). Every metric fell inside the range spanned by the four rungs: median relative error 1.68%, cross-visit Frobenius error 0.60, mechanistic overlap 0.91, median novelty 0.50, and median DCR 13.87. One hundred independently seeded exact cohorts placed these quantities in narrow intervals, with median relative error 1.40% (fifth to ninety-fifth percentile 0.97 to 1.91%) and median DCR 13.95 (13.61 to 14.39). The agreement did not rest on a favorable seed. Over the same 40 train/holdout splits used for the ladder endpoints, the membership-inference AUC under

exact sampling was 0.98, marginally above the published scheme’s 0.97. The memorization was therefore a property of π_r at $K=23$, where every stored patient was a component center. It was not a property of any approximation used to sample it.

Table S16: **Sampler-correctness comparison.** Each rung held β^* , the memory, and the decoding fixed and varied only the direction sampler. Fidelity metrics (median relative error, cross-visit correlation Frobenius error, mechanistic cloud overlap, median novelty) and privacy metrics (median distance to the closest real record, and repeated-split membership-inference AUC where evaluated) were reported per rung. The real-to-real DCR baseline was 15.9. The membership-inference test was run on the two endpoints of the ladder and on the analytic reference, over the same 40 train/holdout splits. *Exact* was not a rung. It drew ancestrally from Eq. (3) with no chain, and served as the analytic reference for the four sampled rungs. Over 100 independently seeded exact cohorts the median relative error was 1.40% (fifth to ninety-fifth percentile 0.97 to 1.91%) and the median DCR was 13.95 (13.61 to 14.39).

Rung	Sampler	MRE (%)	Cross-visit Frob.	Mech. overlap	Novelty	DCR	MIA AUC
V0	anchor, endpoint	1.31	0.65	0.89	0.51	13.97	0.97
V1	sphere init	1.53	0.57	0.89	0.50	13.96	n/a
V2	pooled ULA	2.02	0.59	0.88	0.50	14.15	n/a
V3	pooled MALA	1.89	0.62	0.92	0.52	14.33	0.96
Exact	analytic ancestral	1.68	0.60	0.91	0.50	13.87	0.98




Figure S9: **Sampling correctness and the memorization/privacy relationship.** (A) Across the V0 to V3 ladder the distance to the closest real record (DCR, left axis) rose only slightly and stayed below the real-to-real baseline, while the membership-inference AUC (right axis) remained near 0.96 to 0.97. (B) The fidelity cost of multi-chain pooling, as median relative error. (C) The inverse temperature swept across a 40-fold range around the operating point β^* . The DCR (privacy, left axis) rose as β fell but never reached the real-to-real baseline. The cross-visit covariance error (right axis) grew, and the mechanistic overlap stayed flat.




Figure S10: **Mixing diagnostic at β^* .** Energy traces, the between-chain $\hat{R}$ statistic, and the anchor-distance trajectory for sphere-initialized Langevin chains. The target distribution at β^* was multimodal, with one component centered on each stored pattern. Metropolis acceptance near 0.999 indicated negligible discretization bias.