

# Validated Synthetic Patient Generation for Small Longitudinal Cohorts: Coagulation Dynamics Across Pregnancy

Jeffrey D. Varner<sup>1,\*</sup>, Maria Cristina Bravo<sup>2</sup>, Carole McBride<sup>3</sup>, Thomas Orfeo<sup>2</sup>, Ira Bernstein<sup>3</sup>

<sup>1</sup>Robert F. Smith School of Chemical and Biomolecular Engineering,  
Cornell University, Ithaca, NY 14850

<sup>2</sup>Department of Biochemistry, The Robert Larner, M.D. College of Medicine,  
University of Vermont Medical Center, Burlington, VT 05401

<sup>3</sup>Department of Obstetrics, Gynecology, and Reproductive Sciences,  
The Robert Larner, M.D. College of Medicine,  
University of Vermont Medical Center, Burlington, VT 05401

\*Corresponding author: jdv27@cornell.edu

## Abstract

Small longitudinal clinical cohorts, common in maternal health, rare diseases, and early-phase trials, limit computational modeling: too few patients to train reliable models, yet too costly and slow to expand through additional enrollment. We present multiplicity-weighted Stochastic Attention (SA), a generative framework based on modern Hopfield network theory that addresses this gap. SA embeds real patient profiles as memory patterns in a continuous energy landscape whose induced distribution places one component on each stored patient, and generates novel synthetic patients by sampling it while preserving the geometry of the original cohort. The reported cohort was generated with Langevin dynamics; an exact analytic sampler for the same distribution, available in closed form, reproduced its evaluated behavior. Per-pattern multiplicity weights enable targeted amplification of rare clinical subgroups at inference time without retraining. We applied SA to a longitudinal coagulation dataset from 23 pregnant patients spanning 72 biochemical features across 3 visits (pre-pregnancy baseline, first trimester, and third trimester), including rare subgroups such as polycystic ovary syndrome and preeclampsia. Across multiple validation tests, including an ordinary differential equation model of the coagulation cascade, synthetic patients generated by SA were not distinguishable from their real counterparts by the statistical, structural, and mechanistic measures applied at this sample size. A downstream utility test further showed that a mechanistic model calibrated entirely on synthetic patients predicted held-out real patient outcomes as well as one calibrated on real data. These results demonstrate, as a proof-of-concept, that SA can produce synthetic cohorts useful for the modeling use-cases tested here (mechanistic calibration and hypothesis generation) from very small longitudinal datasets, enabling data-augmented modeling in small-cohort settings.

**Keywords:** synthetic data generation, stochastic attention, modern Hopfield networks, multiplicity weighting, coagulation, pregnancy, longitudinal data, mechanistic validation

# 1 Introduction

In 2021, the United States maternal mortality rate reached 32.9 deaths per 100,000 live births, the highest level since 1965<sup>1,2</sup>. Between 2018 and 2022, Non-Hispanic Black women died at 2.8 times the rate of White women, and disorders related to pregnancy, including hemorrhage and venous complications, were the leading underlying cause of death, accounting for 17.4% of 6,283 pregnancy-related deaths<sup>3</sup>. These figures highlight the role of the coagulation system in maternal outcomes. Blood coagulation is a complex enzymatic cascade in which a small tissue factor (TF) stimulus triggers a series of serine protease activation reactions that ultimately produce thrombin, the central effector of hemostasis<sup>4,5</sup>. Thrombin cleaves fibrinogen into fibrin, activates platelets, and amplifies its own production through positive feedback via Factors V, VIII, and XI, while simultaneously initiating its own shutdown through the thrombomodulin–protein C anticoagulant pathway<sup>6,7</sup>. Pregnancy alters this hemostatic balance, shifting toward a hypercoagulable state that prepares for delivery<sup>8–10</sup>. Procoagulant factors including fibrinogen, Factor VIII, and von Willebrand factor increase substantially across gestation, while natural anticoagulants such as antithrombin and protein S decline. Two conditions further shift this balance in ways that are not fully understood. Polycystic ovary syndrome (PCOS), a hormonal and metabolic disorder affecting approximately 6.6% of reproductive-age women<sup>11</sup>, has been associated with impaired fibrinolysis and a prothrombotic tendency<sup>12</sup>. Preeclampsia (PE), a hypertensive disorder that complicates 2–8% of pregnancies worldwide<sup>13</sup>, is characterized by endothelial dysfunction, platelet activation, and altered thrombin generation<sup>14</sup>. Mathematical models of the coagulation cascade have provided mechanistic insight into patient-specific thrombin generation<sup>15,16</sup>. These range from the original Hockin–Mann stoichiometric framework<sup>6</sup> through extensions incorporating the thrombomodulin–protein C pathway<sup>7</sup> to the Brummel-Ziedins 2012 (BZ2012) model with detailed activated protein C (APC)-mediated Factor Va inactivation kinetics<sup>17</sup>. However, applying these models to pregnancy cohorts requires longitudinal datasets capturing coagulation factors, inhibitors, and thrombin generation assay (TGA) measurements across gestation, data that is rare and expensive to collect.

The core barrier is sample size. Longitudinal pregnancy coagulation studies typically enroll tens of patients with complete follow-up across multiple trimesters, yielding datasets where the number of features  $p$  (coagulation factors, TGA parameters, viscoelastic measurements) exceeds the number of patients  $n$ . Rare complications compound this problem. Small longitudinal cohorts inevitably contain only a handful of patients with any given condition, far too few for condition-specific statistical analysis. This  $n < p$  regime limits conventional approaches because covariance matrices are rank-deficient, cross-validation is unreliable, and overfitting is nearly inevitable<sup>18</sup>. Synthetic data generation addresses this by augmenting small real datasets with statistically faithful synthetic records, enabling downstream analyses (mechanistic modeling, clinical comparisons, hypothesis generation for rare complications) that would otherwise be infeasible<sup>19,20</sup>. However, existing methods face limitations in the small-cohort longitudinal setting. Sampling from a fitted multivariate normal (MVN) distribution is singular when  $n < p$  and must be regularized<sup>21</sup>, introducing bias that distorts the joint distribution. MVN also assumes Gaussian marginals and linear correlations, and provides no mechanism for conditional generation: it cannot amplify a specific subpopulation while preserving its signatures. More expressive models such as generative adversarial networks (GANs) and variational autoencoders (VAEs)<sup>22–24</sup>, including recent longitudinal extensions<sup>20</sup>, require substantially larger training sets and are prone to mode collapse at small  $n$ <sup>25</sup> (we confirm this empirically for Conditional Tabular GAN (CTGAN) at  $n=23$  in this work). A generative method is needed that can operate directly on the geometry of a small dataset without [estimating a full covariance model](#).

Stochastic Attention (SA) is such a framework. [Rather than fitting a full covariance model or](#)

training a neural generator, SA treats the stored patient profiles themselves as the model. Building on modern Hopfield network theory<sup>26</sup>, each profile becomes a memory pattern in a continuous energy landscape, and the attention mechanism makes that landscape a distribution centered on the stored patterns. Sampling it yields new profiles rather than copies, because the memory-centered components overlap in the reduced space. Because sampling occurs in a principal component analysis (PCA)-reduced linear subspace anchored on the observed cohort, SA never forms a full  $p \times p$  covariance matrix and so avoids both the rank deficiency that limits MVN and the mode collapse that limits GAN/VAE approaches when  $n < p$ . We recently introduced a multiplicity-weighted extension of the Hopfield energy that attaches a per-pattern weight to each stored profile<sup>27</sup>, which provides exactly the lever that MVN lacks: continuous interpolation between unconditioned generation and targeted amplification of a rare subpopulation at inference time, without retraining. What remains open, and what this study addresses, is whether SA can faithfully generate *longitudinal* trajectories from a small cohort. In earlier proceedings work, we paired our mechanistic coagulation pipeline with per-visit MVN sampling to produce synthetic populations<sup>28</sup>; that approach captured single-visit marginals but, by fitting each visit independently, generated patients whose trajectories across gestation were statistically incoherent. A generator that respects the joint structure across visits is what a small longitudinal pregnancy cohort actually demands.

In this study, we used multiplicity-weighted SA to generate complete longitudinal patient profiles from a small pregnancy coagulation cohort ( $K=23$  patients, 72 features per visit, 3 visits) that includes rare clinical subgroups (PCOS,  $n=3$ ; Developed PE,  $n=5$ ). Our pipeline concatenates each patient’s multi-visit profile into a single vector, applies PCA for dimensionality reduction, and generates new patients via multiplicity-weighted Langevin dynamics on the Hopfield energy landscape, with a direction-magnitude decomposition that preserves the natural variance structure of continuous clinical measurements. We then asked whether the synthetic patients reproduced marginal feature distributions, the cross-visit covariance structure, condition-specific signatures of rare clinical subgroups, and the biological plausibility expected by an independent ordinary differential equation (ODE) model of the coagulation cascade<sup>17</sup>. As a downstream-utility test, we further asked whether a mechanistic model calibrated entirely on synthetic patients could predict held-out real patient outcomes. The results, presented in the next section, show that SA can, as a proof-of-concept, recover both the statistical structure and the mechanistic plausibility of the real cohort at this sample size. If this generalizes, the implication is that the bottleneck for studying rare obstetric and pediatric conditions may be shifting from cohort size toward cohort fidelity: a few dozen carefully phenotyped longitudinal patients, augmented through SA, may be enough to support mechanistic and statistical analyses that have traditionally relied on much larger cohorts.

## 2 Results

We generated  $N=100$  synthetic longitudinal patient profiles using SA and compared them against the  $K=23$  real patients across four validation levels. This sample size provided approximately  $4\times$  amplification while remaining in the regime where SA produces diverse, non-degenerate samples. As a baseline, we also generated  $N=100$  patients from a regularized multivariate normal (MVN) distribution fitted to the same data.

### 2.1 Marginal Plausibility

We first assessed whether SA-generated patients reproduced per-feature summary statistics across all three visits. The pooled median mean relative error (MRE) across all 216 feature–visit entries was 1.2% (95% bootstrap CI: 1.0–1.6%), with 193 of 216 entries (89%) below 5% (Table S5),

indicating that SA captured the central tendencies of the real population. Per-visit MREs for representative coagulation features are summarized in Table 2, where Factor II and Factor VIII fell below 2%, antithrombin below 3%, and fibrinogen below 1% across all three visits. We verified that synthetic patients were not memorized copies of real patients using two complementary metrics summarized in Table S9. The mean novelty score ( $1 - \max_k \cos(\hat{\xi}, \hat{\mathbf{m}}_k)$ ) was 0.50 at  $\beta = \beta^*$ , indicating that generated samples were on average 50% away from their nearest stored pattern in angular distance. The median nearest-neighbor distance from each synthetic patient to its closest real patient was 6.14 standardized units, compared with 6.87 for the median real-to-real nearest-neighbor distance, a ratio of 0.89 indicating that synthetic patients were approximately as far from real patients as real patients were from each other.

We next examined whether known physiological relationships and longitudinal trajectories were preserved in the synthetic cohort. Pairwise scatter plots of coagulation factor levels against thrombin generation parameters showed that SA-generated patients occupied the same joint regions as real patients (Fig. 1). The expected inverse relationship between antithrombin and TF-initiated peak thrombin was preserved, as were the positive associations between Factor VIII and peak, prothrombin and endogenous thrombin potential (ETP), and the inverse relationship between antithrombin and ETP. Per-visit means and standard deviations for six key coagulation features confirmed that synthetic patients tracked the real longitudinal trajectories closely, with overlapping confidence bands at all three visits (Fig. 2). Fibrinogen, Factor VIII, and von Willebrand factor showed the expected monotonic increases from baseline through the third trimester in both cohorts, while Factor II increased modestly and antithrombin declined then partially recovered, consistent with the known hypercoagulable shift of pregnancy. Together, these results indicated that SA captured not only univariate distributions and pairwise dependencies but also the directionality and magnitude of longitudinal change. The 72 features also included viscoelastic (ROTEM) and fibrinolytic parameters, which were generated as part of the concatenated vector and showed comparable MREs to the coagulation factors (Table S5), but are not individually highlighted here because the mechanistic validation uses the BZ2012 thrombin generation model, which does not cover fibrinolysis or clot viscoelasticity. *These features nonetheless load onto the retained principal components and so shape the energy landscape from which every synthetic patient is drawn, so their statistical fidelity matters even though they are not mechanistically checked; a fibrinolysis model under development is the route to closing that gap.*

We also compared SA against two widely used deep generative methods for tabular data, Conditional Tabular GAN (CTGAN) and Tabular Variational Autoencoder (TVAE)<sup>24</sup> (Table S6). CTGAN failed completely, producing median MREs of  $\approx 19\%$ , an order of magnitude worse than SA, across all epoch counts tested (300–3,000), consistent with known GAN mode collapse on small datasets. TVAE achieved comparable marginal fidelity at high epoch counts (median MRE 1.8% at 3,000 epochs), but operated on per-visit rows ( $n=90$ ) rather than concatenated longitudinal profiles ( $n=23$ ), meaning it could not capture cross-visit covariance structure and had no mechanism for conditional subgroup generation.

*To ask how much of this fidelity comes from the shared principal-component subspace rather than from the stochastic-attention dynamics, we compared SA against four generators fit in the identical  $d=18$  PCA subspace, a Gaussian mixture, a kernel-density estimator,  $k$ -nearest-neighbor interpolation, and a Gaussian copula (Supplementary Table S13). The subspace proved necessary but not sufficient. Two of the four baselines failed despite sharing it, with  $k$ -nearest-neighbor interpolation collapsing novelty (median 0.05, essentially returning the training patients) and the kernel-density estimator degrading both marginal fidelity (1.9% median relative error) and mechanistic plausibility (0.82 cloud overlap). The mixture and copula baselines cleared both bars and were competitive with SA on marginal error (1.1% and 1.3% against SA’s 1.2%) and mechanistic*

overlap (0.92 and 0.93 against 0.90), with SA and the copula the only generators for which every synthetic patient exceeded the novelty threshold. Much of the marginal and mechanistic fidelity is therefore supplied by the subspace and is available to any generator that samples it sensibly, which we take to be the substantive answer to whether the representation or the dynamics carries the result.

The tie did not extend to the remaining axes. SA showed the largest empirical cross-visit correlation error of the five methods (0.64 against 0.42 for the mixture and copula), but this quantity measures reproduction of the correlation matrix of the same 23 patients every generator was fit to, which at this sample size is a noisy estimate of the population structure rather than a test of generalization. On the simulation benchmark, where the population covariance is known by construction, SA recovered that structure more accurately than the multivariate normal baseline across the  $n < p$  regime (Fig. 4A). SA also fixes its operating point from the attention-entropy inflection rather than by selecting a bandwidth or component count on 23 patients, and it admits the inference-time multiplicity steering used below for conditional subgroup generation, which the subspace baselines do not provide directly.

## 2.2 Cross-Visit Covariance Structure

Having established that SA reproduces marginal distributions and pairwise relationships, we next asked whether it preserves the higher-order joint structure across visits. We compared the full  $216 \times 216$  cross-visit correlation matrices for real, SA-generated, and MVN-generated populations. The real data exhibited a characteristic block structure with strong within-visit correlations along the diagonal and structured off-diagonal blocks reflecting cross-visit dependencies; for example, a patient’s Visit 1 Factor X level predicting their Visit 3 Factor X level. We found that SA preserved this block structure, including the off-diagonal cross-visit correlations (Fig. 3). The SA–Real residual matrix showed small, unstructured deviations distributed throughout, whereas the MVN–Real residual was notably smoother, reflecting the regularization-induced shrinkage of off-diagonal correlations toward zero. This difference was most apparent in the off-diagonal blocks, where MVN systematically underestimated cross-visit dependencies that SA retained. We note that Pearson correlation captures linear associations; Spearman rank correlations, used in the mechanistic validation, would additionally capture monotonic nonlinear dependencies.

We examined the eigenvalue spectrum of the sample covariance matrix to understand the structural origin of this difference (Fig. S2). The real data had rank 22, reflecting the  $K=23$  patient constraint. SA and MVN make different trade-offs in handling this rank deficiency: SA truncates via PCA (zeroing variance beyond component 18), while MVN regularizes via Ledoit–Wolf shrinkage (providing a full-rank estimate by inflating eigenvalues in the 194 null dimensions). The consequence of MVN’s approach is that it introduces variance in dimensions where the data contain no signal, which manifested as increased dispersion in PCA projections. PCA projections by visit confirmed this dispersion gap (Supplementary Fig. S8). SA-generated patients occupied the same region of PCA space as real patients across all three visits, while MVN-generated patients showed substantially greater scatter, particularly at Visits 2 and 3.

## 2.3 Robustness Across Data Regimes

To test whether SA recovers population structure rather than reproducing the training cohort, we evaluated it on synthetic ground truth for which the true covariance is known and independent of any particular draw. On a low-rank Gaussian generative model swept across sample size  $n$ , dimensionality  $p$ , and latent rank, SA recovered the population covariance more accurately than a



multivariate Gaussian throughout the rank-deficient  $n < p$  regime (Fig. 4A): at the real cohort’s shape ( $n=23$ ,  $p=216$ ) the Gaussian’s relative covariance error was about 1.5 times SA’s, and SA was more accurate in 35 of the 39 Gaussian  $n < p$  cells. Because the target covariance is fixed in advance, this is genuine generalization rather than memorization, a distinction the empirical cross-visit comparison could not make. The advantage narrowed as  $n$  approached and exceeded  $p$ , where the Gaussian is well specified, so the benefit is specific to the small-sample regime this cohort occupies.

On the real cohort, subsampling the 23 patients to as few as eight and regenerating produced graceful degradation rather than collapse: recovery of the full-cohort cross-visit correlation structure worsened smoothly as  $n$  fell (Fig. 4B), the pooled marginal error rose from about 2% at  $n=20$  to about 5% at  $n=8$ , and mechanistic plausibility stayed near 0.90 across all subset sizes. On a deliberately nonlinear ground truth, an S-curve manifold, SA’s linear principal-component subspace began to place generated points off the true manifold as curvature increased, with the off-manifold residual crossing twice its linear-baseline value at a curvature of roughly 0.25 to 0.5; on this nonlinear target SA remained closer to the manifold than the multivariate Gaussian, whereas the Gaussian recovered the gross covariance better, so neither method is uniformly preferable and the linear embedding is best understood as a deliberate and bounded small-cohort choice.

## 2.4 Interpolation and Extrapolation

To ask whether the sampler stays within the real cohort or reaches beyond it, we tested convex-hull membership of the synthetic patients against the 23 real patients in the  $d=18$  principal-component space (Supplementary Table S14). Every synthetic patient fell outside the hull. So did every real patient with respect to the other 22, and when the comparison was matched for severity by removing each point’s nearest vertex on both sides, so that real and synthetic patients faced equally depleted hulls, the distance to the hull was not distinguishable between them. The hull criterion did not separate synthetic patients from real ones at this sample size.

The reason is geometric rather than statistical. With 23 points in 18 dimensions the hull is nearly a simplex, reaching the patients’ typical radius only along the 23 vertex directions themselves and extending a small fraction of that distance in a generic direction, so almost any point that does not sit on a stored pattern lies outside it. The sampler compounds this because it normalizes the stored patterns to unit length and returns each sample to the unit sphere, and a unit vector inside the convex hull of unit-norm patterns must coincide with one of those patterns, which makes hull membership structurally unavailable to this generator rather than merely uncommon. We therefore read the result as a limit on what the hull criterion can resolve at  $K=23$  rather than as evidence of runaway extrapolation, and assessed the generated points against the plausibility criteria directly, where most satisfied the biological constraints and per-feature mechanistic ratios stayed inside the real cohort’s percentile band for the large majority of patients, although the strict criterion requiring every feature and visit to fall inside the band simultaneously was met by fewer than half.

## 2.5 Re-identification Risk

Because the energy landscape places a basin at each stored patient, we asked how much the generated cohort reveals about who was in it. For every synthetic patient we measured the distance to the closest real record and compared it against the leave-one-out distance between real patients, and we ran a nearest-neighbor membership-inference attack that retraining the generator on a subset of the cohort and scores each real record by its proximity to the resulting synthetic set (Supple-

mentary Table S15). Synthetic patients sat slightly closer to the real cohort than real patients sat to one another, and the attack separated training members from held-out patients with a mean AUC of 0.97 across 40 random splits, reaching perfect separation on roughly a fifth of them. At the published operating point the generator therefore reproduces its training cohort closely enough for membership to be recovered.

Decomposing the attack shows what it does and does not recover. Held-out patients sat at essentially the generic distance between two unrelated patients from the synthetic cohort, matching the leave-one-out distance among the real patients themselves, whereas training members sat considerably closer. The attack therefore recovers which patients were stored in the memory rather than information about patients the generator never saw, and the separation follows from the same small-cohort geometry that left the hull criterion uninformative. The mechanism is visible in the form of the target itself: the distribution the energy induces is a mixture with one component centered on each stored patient (Eq. (3)), so every draw descends from a single stored patient and carries proximity to that patient into the decoded profile. Nothing in the construction averages membership across several patients. The attack also assumes an adversary who already holds a patient’s record and asks whether it was in the training cohort, so it bounds membership disclosure rather than showing that patient-level values can be reconstructed from the synthetic cohort alone. Whether the leak reflects the target distribution or the sampler used to draw from it is taken up separately, where Metropolis-adjusted sampling and a wide sweep of the inverse temperature both leave it essentially unchanged, and exact ancestral draws from Eq. (3), which involve no chain at all, reproduce it undiminished at an AUC of 0.98 (Supplementary Table S16, Fig. S9). The signal belongs to the target distribution at this cohort size, not to the sampler.

## 2.6 Conditional Generation of Rare Subgroups

The previous analyses used unconditioned SA, generating from the full cohort. A key clinical application, however, is generating patients from specific subpopulations that are too small to study independently. We therefore tested whether SA could amplify rare clinical subpopulations while preserving their condition-specific signatures. We defined three overlapping subgroups by pregnancy outcome and comorbidity: Uncomplicated ( $n=18$ ), PCOS ( $n=3$ , cross-cutting; 1 patient also in the PE group), and Developed PE ( $n=5$ ). The PCOS and Developed PE groups were too small for any independent statistical analysis. We used SA’s multiplicity-weighted sampling to generate condition-specific cohorts of 100 patients each by upweighting attention on the relevant stored patterns.

We compared condition-specific feature means between real and SA-generated patients across eight coagulation features that exhibited between-condition variation (Fig. 5). SA-generated cohorts preserved the condition-specific patterns. PCOS patients showed elevated Factor VIII and vWF relative to Healthy patients in both real and synthetic data, PE patients showed elevated  $\alpha_2$ -AP and Factor IX, and the between-condition rank ordering was maintained across all eight features. To quantify equivalence, we performed a bootstrap Mann–Whitney test. For each feature and condition, we subsampled the synthetic cohort to match the real sample size ( $n_{\text{real}}$ ), applied a Mann–Whitney U test, and repeated 1,000 times. We reported the fraction of replicates where  $p > 0.05$  (i.e., real and synthetic were statistically indistinguishable). Across all 24 feature–condition pairs, 20 (83%) achieved  $p > 0.05$  in  $\geq 90\%$  of replicates, with a median non-significance fraction of 98.6% (full per-pair results in Table S10). The four features that fell below 90% were high-variance features in the smallest subgroups (FIX and plasminogen in PCOS; FIX and plasminogen in Developed PE), where the real data exhibited large inter-patient variability. No multiple-testing correction was applied to the 24 simultaneous comparisons; the failing pairs

had consistently low non-significance fractions (62–87%) rather than borderline values, so correction would not change the classification. PCA projections of the conditioned cohorts are provided in Figs. S3–S4. These results showed that SA’s multiplicity-weighted interpolation can amplify rare subgroups in a way that preserves their distinguishing clinical features, useful for hypothesis generation and power analysis in small patient populations. MVN has no mechanism for this: fitting a separate MVN to three PCOS patients across 216 features is mathematically impossible (rank 2 covariance), and post-hoc subsetting of unconditional MVN samples does not preserve subgroup-specific structure.

The amplification is bounded by how few patterns underlie it. At the target effective fraction the participation ratio was 4.6 for the PCOS cohort and 7.7 for the Developed PE cohort (Table S2), so each cohort of 100 generated patients is a recombination of three and five real individuals respectively. Those cohorts carry no more information about the conditions than the underlying patients do, and any downstream analysis treating them as 100 independent observations will understate uncertainty. What conditioning provides is a denser sample of the distribution implied by those few patients, which supports hypothesis generation and power calculation rather than additional evidence about the conditions themselves.

## 2.7 Mechanistic Consistency

The preceding validations assessed statistical properties of the synthetic data. We next tested whether SA-generated patients produce biologically plausible outputs when their coagulation factor levels are fed through a mechanistic model that knows nothing about the SA generation process, though it was calibrated on the same cohort. We used the Hockin–Mann BZ2012 model<sup>6,17</sup>, a system of 58 ordinary differential equations with 64 rate constants that simulates thrombin generation from patient-specific coagulation factor inputs. We calibrated 5 of the 64 rate constants (prothrombinase, intrinsic and extrinsic tenase, protein C activation, and meizothrombin conversion; Table S4) on Visit 1 real patients at the population level and held the remaining 59 at published literature values. We then ran the calibrated model on all 23 real patients and 100 synthetic patients under both TF-only and TF+TM conditions (738 total simulations, 0 failures). For each patient, we computed five ODE-predicted thrombin generation features (lagtime, peak, time-to-peak, maximum rate, and ETP) and compared them against the corresponding values from the patient record. In these comparisons (Fig. 6, top row), the horizontal axis is the TGA value from the patient record, which exists for both real and synthetic patients, while the vertical axis is the value computed by the BZ2012 model from that patient’s factor inputs. Systematic deviations from the  $y=x$  line reflect ODE model bias (e.g., lagtime is underpredicted at  $\approx 0.77\times$ ), not a failure of the synthetic data; the key observation is that real and synthetic patients share the same pattern of model bias, occupying the same cloud with the same systematic offset.

We formalized this observation by computing the ODE-predicted/dataset ratio for each patient and comparing the resulting distributions between real and synthetic populations (Fig. 6, bottom row). These ratio distributions capture how the ODE model processes each patient’s factor levels. If SA had generated patients with biologically implausible factor combinations, the model would process them differently, producing a shifted or broadened ratio distribution. Instead, the distributions overlapped substantially, with cloud overlap (fraction of synthetic ratios within the 5th–95th percentile of real ratios) ranging from 0.83 for ETP to 0.92 for T.Peak under TF-only conditions, and two-sample Kolmogorov–Smirnov (KS) tests did not detect a difference between the real and synthetic ratio distributions at this sample size across all five TGA features ( $D = 0.081$ – $0.123$ , all  $p > 0.30$ ; full diagnostics in Table S11). Under TF+TM conditions, the model systematically overcorrected the protein C anticoagulant pathway, producing peak and maximum rate predictions



approximately  $0.5\times$  measured values for both real and synthetic patients (Fig. S5, same layout as the main-text comparison), yet this shared systematic bias was itself informative. Cloud overlap remained high (0.88–0.92 across the five features), confirming that the model processed real and synthetic patients identically even under conditions where the calibration was imperfect.

The calibration also generalized across pregnancy timepoints. Although fitted only on Visit 1 real patients, the per-visit median predicted-to-measured ratios under TF-only conditions remained close to 1.0 at Visits 2 and 3 (Fig. S6), confirming that the calibrated rate constants captured a stable population-level property of the coagulation system rather than overfitting to the training visit. Spearman rank correlations between predicted and measured values were moderate for ETP ( $\rho \approx 0.6$ –0.8 for real,  $\rho \approx 0.6$ –0.65 for synthetic; Fig. S7) and weak for other features, consistent with a 5-parameter population-level calibration that was not designed to resolve inter-patient variability. The key finding was not patient-level prediction but population-level plausibility. A mechanistic model of the coagulation cascade, structurally independent of the generator but calibrated on the same cohort, placed SA-generated patients within the same biologically plausible envelope as real patients, and could not distinguish them under either experimental condition.

## 2.8 Downstream Utility: Mechanistic Model Calibration

The preceding results established that SA-generated patients are statistically and mechanistically plausible. To test whether this plausibility translates to practical utility, we calibrated a mechanistic model entirely on synthetic patients and evaluated whether it could predict real patient outcomes. We calibrated the BZ2012 model (TF-only condition, 5 rate constants) on synthetic V1 patients ( $N=100$ ) using the same bounded Nelder–Mead optimization applied to the real V1 calibration ( $K=23$ ), with 11 random restarts to mitigate local-minimum sensitivity (Table S4 for parameter details). We then evaluated both calibrations on held-out real V2 and V3 patients that neither model had seen during training. The synth-calibrated model achieved comparable or slightly better performance across all five TGA features, with per-feature median relative errors 2–10% lower than the real-calibrated model (overall ratio  $0.94\times$ ; Fig. 7; per-feature breakdown in Table S12). This improvement likely reflects the smoother loss landscape afforded by  $N=100$  synthetic training patients compared with  $K=23$  real ones, reducing overfitting during optimization. The two calibrations found different parameter values but produced highly correlated predictions across all five TGA features, confirming that the synthetic data captured the same underlying coagulation structure. We note that the synthetic training data is one step removed from the real data (real  $\rightarrow$  SA  $\rightarrow$  synthetic  $\rightarrow$  calibration), so the synth-calibrated model is not fully independent of the real data; rather, it demonstrates that SA’s representation preserves sufficient biological structure for a mechanistic model to learn generalizable parameters.

The direction of that comparison should not be overread. Every piece of information in the synthetic training set derives from the same 23 patients, so the modest advantage is a variance-reduction effect rather than evidence of new biological content. Fitting five rate constants against 100 samples drawn from a distribution estimated on 23 patients smooths the objective and dilutes the influence of any single patient, which would be expected of any generator producing a denser sample of the same distribution. The result establishes that the synthetic cohort is a usable substitute for calibration, not a better one.

Read alongside the conditioned cohorts, this bounds what the method offers in practice. It enables analyses that the original sample size forecloses, calibrating a mechanistic model without spending the real cohort on training and comparing subgroups holding three and five patients, and it points to factor-level differences a larger study could test. It does not substitute for additional patients, and none of these results speak to clinical deployment, which would require prospective

### 3 Discussion

This study tested whether a geometry-preserving generative method could produce synthetic patients from a very small longitudinal cohort that were not merely statistically similar to real patients but consistent with the underlying coagulation biology. Our results provided evidence at each validation level that SA-generated patients met this standard. Individual feature distributions were reproduced with median MRE of 1.2%, which MVN also achieved by construction. SA additionally preserved the joint cross-visit covariance structure that MVN’s rank-deficient estimate could not represent, a difference made visible in the eigenvalue spectrum where MVN introduced spurious variance in 194 null dimensions while SA respected the low-rank geometry of the data. SA could amplify clinically meaningful rare subgroups, generating 100 synthetic PCOS patients from only 3 real ones, while preserving subgroup-specific signatures that MVN cannot capture. And an independent ODE model of the coagulation cascade, calibrated exclusively on real patients, confirmed that the factor-to-thrombin-generation mapping in synthetic patients was biologically plausible, with 83–92% cloud overlap and Kolmogorov–Smirnov tests unable to distinguish the two populations ( $p > 0.35$  for all five TGA features).

The ability to generate validated synthetic cohorts from very small longitudinal datasets has practical implications for maternal health research. Rare pregnancy complications such as PCOS and preeclampsia are difficult to study because assembling large, well-characterized longitudinal cohorts requires years of recruitment across multiple clinical sites. Our results suggest that SA can enable hypothesis generation, power analyses, and computational modeling studies for these understudied populations by generating synthetic cohorts that preserve both the statistical and mechanistic properties of the real data. The conditional generation capability is particularly relevant. Rather than requiring researchers to collect additional rare patients, SA can amplify existing small subgroups into cohorts large enough for meaningful analysis.

Understanding why SA succeeds where parametric methods fail requires examining the geometric structure of the problem. SA generates samples by interpolating between stored patient profiles via the Hopfield energy landscape operating in a PCA-reduced linear subspace, rather than fitting a parametric distribution that must regularize or truncate high-dimensional structure. In the  $n < p$  regime that characterizes most small clinical studies, any parametric distribution faces a fundamental identifiability problem. There are more parameters to estimate than observations to estimate them from. SA sidesteps this by operating in a PCA-reduced space where the memory-to-dimension ratio is favorable ( $K/d_{\text{PCA}} \approx 1.28$ ) and by preserving the full directional structure of the data through the attention mechanism. A natural concern is whether the  $\beta^*$  selection procedure, originally developed for protein sequence generation<sup>29</sup>, transfers to clinical coagulation data. The entropy inflection method identifies the phase transition in the Hopfield energy landscape, which is a geometric property of  $K$  patterns in  $d$  dimensions, not a property of what the features represent; for our dataset the empirical inflection ( $\beta^* \approx 2.94$ ) was close to the theoretical prediction ( $\sqrt{d} \approx 4.24$ ), and a sensitivity analysis confirmed that generation quality degraded gracefully across  $\beta/\beta^* \in [0.1, 3.0]$  (Table S1). Similarly, varying the PCA variance retention threshold from 85% to 99% ( $d_{\text{PCA}} = 13$  to 21) had minimal impact on generation quality. Median MRE remained between 1.2% and 1.8% across all thresholds (Table S7), confirming that the 95% threshold was not a critical design choice. The multiplicity weighting for conditional generation introduced an additional consideration. Achieving 80% attention on the PCOS subgroup required  $\rho \approx 26.7$ , reducing the effective pattern count to  $K_{\text{eff}} \approx 4.6$  (multiplicity parameters for all three subgroups in Table S2),

which explains the larger MREs observed for PCOS features because the energy landscape was effectively shaped by fewer than 5 independent patterns, pulling means toward the population average. The direction-magnitude decomposition compounds this limitation. For PCOS, magnitudes are drawn from only 3 empirical norms, providing minimal diversity in the scale of generated samples. More broadly, the SA framework and its multiplicity-weighted conditioning have been demonstrated in companion studies from this group on discrete protein sequence generation from small family alignments<sup>29</sup>, on steering generation toward functional subsets such as binding peptides<sup>27</sup>, and on the theoretical foundations connecting modern Hopfield networks to attention-based generation<sup>30</sup>. Across these domains, the critical temperature prediction ( $\beta^* \approx \sqrt{d}$ ) and the multiplicity-weighted conditioning mechanism transfer without modification, suggesting that the geometric properties exploited here are not specific to clinical coagulation data but are general features of the Hopfield energy landscape operating on small pattern sets.

An alternative precedent for validated longitudinal synthetic generation is the VAMBN-MT framework of Kühnel et al.<sup>20</sup>, which extends a Variational Autoencoder Modular Bayesian Network with an LSTM time-encoder to better capture cross-visit dependencies. They evaluated VAMBN-MT on the DONALD nutritional cohort and on the Alzheimer’s Disease Neuroimaging Initiative (ADNI) cohort. Their study is instructive because it offers a regime contrast to ours. With  $N=1,312$  DONALD participants and 33 longitudinal variables, their setting sits in the data-rich regime for which VAMBN-MT was designed, where deep parametric models can be trained directly on the cohort. The published configuration combines expert-curated module assignments, Bayesian-network edge constraints, an LSTM-augmented HI-VAE per module, and 10,000 synthetic samples drawn from the trained model to reproduce a published age-trend regression. In their evaluation, the best variant reported a cross-visit correlation matrix relative Frobenius error of roughly 0.70, and downstream time trends weaker than the strongest signals in the real data were not reproduced at the sample sizes they tested, both honest indicators of how hard this generative task is even in the data-rich regime. On the other hand, our setting inverts both the number of patients and the number of features. With  $K=23$  patients across 216 features we operate two orders of magnitude smaller and squarely in the  $n < p$  regime, well below the cohort sizes for which VAMBN-MT was developed. The non-parametric design of SA addresses this regime in three ways that contrast with VAMBN-MT. First, PCA-derived linear geometry replaces the modular HI-VAE architecture as the mechanism for capturing cross-visit dependencies. Second, a generic mechanistic-model validation pattern, transferable to any domain with a calibrated ODE model, replaces the domain-specific dependency rules used in their evaluation, such as graduation-status monotonicity, age-time arithmetic identities, and macronutrient summation. Finally, multiplicity weighting provides an inference-time conditional-generation lever that the VAMBN-MT framework does not offer. These are complementary positions in the design space rather than a head-to-head comparison. VAMBN-MT addresses the moderate-cohort regime where deep generative models are feasible but expert-specified modular structure helps capture cross-visit dependencies, while SA addresses the small-cohort regime where parametric models cannot be fit and the geometry of the few real patients must itself carry the generative signal.

Beyond these statistical and geometric considerations, the mechanistic validation introduced an approach that applies beyond coagulation. Rather than relying solely on statistical comparisons between real and synthetic distributions, we tested whether synthetic patients produced biologically consistent outputs when processed by an independent computational model, and went further to show that a mechanistic model calibrated entirely on synthetic data predicted real patient outcomes as well as one calibrated on real data (overall ratio  $0.94\times$  on held-out V2+V3 patients, with 11 random restarts per calibration). This approach is available whenever a validated mechanistic model exists for the domain of interest (pharmacokinetic models in drug development<sup>31</sup>, physiological

models in critical care<sup>32</sup>, tumor growth models<sup>33</sup>, or metabolic models<sup>34,35</sup>), and the requirement is not that the model be a perfect predictor but that synthetic patients fall within the same model-predicted envelope as real patients. The bootstrap Mann–Whitney analysis of condition-specific features revealed that 5 of 24 (21%) feature–condition pairs were statistically distinguishable at equal sample sizes, but the mean relative errors for these features remained small (5–15%), and the detected differences likely reflected SA’s tendency to compress distributional tails rather than shift means. This tail compression arises from two sources: the PCA dimensionality reduction, which cannot fully represent higher-order moments of every feature, and the direction-magnitude decomposition, which draws magnitudes from a finite empirical distribution of  $K=23$  norms. This is a trade-off. SA sacrifices some tail fidelity in exchange for preserving the manifold geometry that parametric methods cannot capture in the  $n < p$  regime.

The clinical reading of that trade-off deserves weight, because extreme states are often the phenomena of interest. If a generated cohort under-represents the severe end of a distribution, estimates of how often a severe phenotype occurs, or of how extreme it can become, will be biased toward the center. The extrapolation analysis shows that the generator does place patients outside the convex hull of the real cohort, so it is not confined to the observed range, but producing out-of-range points is not the same as reproducing the frequency of extreme ones, and the two should not be conflated. We therefore recommend against using SA cohorts to estimate extreme-quantile risk without real anchoring data. The mechanistic checks reported here are comparatively insensitive to this limitation, since they compare distributions within a fifth to ninety-fifth percentile band and so are by construction a test of the central mapping from factor levels to thrombin generation rather than of the tails.

Several of the findings above are best read as one observation rather than as a list. With 23 patients in a 216-dimensional measurement space the cohort is sparse enough that familiar diagnostics stop discriminating. Convex-hull membership rejects every synthetic patient, but it rejects every real patient with respect to the others as well, so it reports the resolution of the criterion rather than a property of the generator. The stored patterns sit far enough apart that their basins do not overlap, which is why a membership-inference attack succeeds and why conditioning on three PCOS patients leaves an effective pattern count below five. The advantage of the synthetic cohort in downstream calibration is a variance-reduction effect obtained by sampling more densely from a distribution estimated on those same patients. And the shared principal-component subspace supplies much of the marginal and mechanistic fidelity, which is why several baselines fit in that subspace match SA on those axes. Each of these follows from the sample size rather than from a separate defect of the method.

This study had several limitations. Our study used a single longitudinal dataset of  $K=23$  patients; further evaluation on independent datasets and across different clinical domains is needed to establish generalizability. We addressed this in two partial ways rather than with a second clinical cohort. The simulation benchmark supplies a ground truth that no single dataset can, since the population covariance is known in advance and generalization can be separated from reproduction of the sample. And the same energy landscape, operating-point rule, and multiplicity conditioning have been applied to discrete protein sequence generation from small family alignments and to steering generation toward functional subsets, which is evidence that the mechanism exploited here is geometric rather than specific to coagulation. Neither substitutes for replication in an independent clinical cohort. The PCA dimensionality reduction assumed linear structure in the feature space; nonlinear manifold learning methods may better capture complex clinical relationships, though at the cost of reduced interpretability. The mechanistic validation under TF+TM conditions revealed systematic ODE model bias, though the observation that real and synthetic patients shared the same bias pattern was itself informative. The BZ2012 calibration required substantial departures

from published rate constants, notably a 50-fold reduction in intrinsic tenase  $k_{\text{cat}}$  and a 15-fold increase in extrinsic tenase  $k_{\text{cat}}$  (Table S4). These deviations likely reflect differences between the *in vitro* conditions under which the literature rate constants were measured (purified component systems) and the plasma-based TGA used in this study, where phospholipid surfaces, factor concentrations, and inhibitor profiles differ from the model assumptions<sup>6,17</sup>. The purpose of the mechanistic validation was not to obtain biochemically accurate rate constants but to test whether the ODE model processed real and synthetic patients equivalently, a comparison that is invariant to the absolute parameter values. Because the ODE was calibrated on this same cohort, the mechanistic agreement demonstrates consistency with a model fitted to the same biological system rather than validation against a fully independent ground truth; its evidential value rests on the model being blind to whether a profile is real or synthetic and applying one fixed mapping to both. The mechanistic validation was limited to thrombin generation (BZ2012) and did not cover the fibrinolytic or viscoelastic features included in the dataset; extending this validation to a mechanistic model of fibrinolysis, currently under development, would provide additional evidence for these feature categories. Finally, we did not validate synthetic patients against clinical outcomes; our validation was restricted to statistical and mechanistic plausibility, and prospective validation against real clinical decision-making remains an important future direction.

The concatenation that makes cross-visit structure available also constrains the data the method accepts. Stacking three visits into a single vector assumes that every patient has every visit and that visits are aligned across patients, which holds in this cohort by design but not in most observational data. Irregular timing would place physiologically different states in the same coordinate, so the principal components would mix visit structure with between-patient variation. Missing assays or missing visits have no natural representation in a fixed-length vector and would require imputation before the embedding, which at this cohort size would itself be a substantial source of uncertainty. Variable-length records fall outside the formulation altogether and would need either padding to a common length or a sequence-aware extension in which attention operates over visits as well as over patients. These are scope boundaries of the present formulation rather than properties of the energy landscape, which is indifferent to what the coordinates mean.

Cost and behavior at larger cohort sizes follow from how the target is sampled. In the Langevin implementation used here, each step evaluates the attention weights over the stored patterns and forms their weighted combination, so the work per step is proportional to the product of the pattern count and the retained dimension, and the cost of one sample is that product multiplied by the number of steps. The closed form removes the sequential factor entirely for future use: drawing a component and adding isotropic noise costs one selection and one Gaussian draw per sample, independent of any step count, and the remaining work is the principal-component decomposition and the decoding, which are shared across the cohort. The method is linear in both the number of patients and the reduced dimension either way, and at very high feature counts the practical bound is the principal-component decomposition rather than the sampling. Statistically the picture runs the other way. The simulation benchmark shows that SA’s advantage over a fitted Gaussian is largest when the sample is smallest and narrows as the sample size approaches the dimension, where the Gaussian becomes well specified, so the method is positioned for the small-cohort regime rather than as a general-purpose generator. We expect the memorization documented above to ease as the cohort grows, because basins that currently stand apart would begin to overlap and generation would interpolate across several patients rather than remaining near one, but this cohort admits only subsampling downward and we did not test that expectation.

The linear subspace deserves the same scrutiny. Principal components capture linear structure, so a manifold with strong curvature, such as a branching trajectory or a threshold effect, is represented only approximately, and that approximation error is one of the sources of the tail



compression discussed above. The simulation benchmark bounds where this begins to matter. On a deliberately curved ground truth the off-manifold residual reached twice its linear-baseline value at a curvature of roughly 0.25 to 0.5, and beyond that point SA nonetheless stayed closer to the manifold than a fitted Gaussian, which in turn recovered the gross covariance better, so neither method dominates once the target stops being linear. Nonlinear embeddings such as kernel principal components, diffusion maps, or autoencoders would represent curvature more faithfully, at three costs. They give up components expressed directly in assay space and the interpretability that follows, they give up the closed-form relation between the critical temperature and the dimension that fixes the operating point here, and they introduce fitting choices that are themselves unstable to estimate from 23 patients. Linearity is therefore a deliberate choice matched to the cohort size rather than an oversight, and the subspace ablation indicates which behaviors follow from the representation, since generators sharing that subspace match SA on marginal and mechanistic fidelity while differing sharply in how far they move from the stored patients.

A distinct limitation concerns privacy. Because the generator concentrates probability mass near each of the  $K=23$  stored patterns, synthetic patients lie close to individual real records, and a nearest-neighbor membership-inference attack separates cohort members from held-out patients with high accuracy. We find that this behavior is intrinsic to the target distribution at this cohort size rather than an artifact of the sampler or the operating temperature. Metropolis-adjusted sampling, which removes the discretization bias of the unadjusted scheme, leaves the signal essentially unchanged, and lowering the inverse temperature improves the distance to the closest real record only marginally, never to the level separating real patients from one another, and at a measurable cost to cross-visit covariance fidelity. The attribution is settled by the closed form of the target itself. Because the latent distribution is a Gaussian mixture whose components are centered on the stored patterns, it can be drawn from analytically rather than by any chain, and exact draws carried through the same decoding pipeline reproduce the membership-inference signal undiminished. Proximity to individual real records is therefore a property of the distribution the method defines at this cohort size, not of how that distribution is approximated. We therefore do not present the method as privacy-preserving. It is not differentially private, and the practical risk is bounded instead by the fully de-identified nature of the source cohort and by intended uses, such as mechanistic calibration and hypothesis generation, that do not require releasing patient-level records. The burn-in and thinning used in the sampling diagnostic were derived once on the full cohort and reused for the subset analyses, so the diagnostic characterizes the operating regime rather than certifying each subset independently.

## 4 Conclusion

We have demonstrated that multiplicity-weighted Stochastic Attention can generate synthetic longitudinal patient data from very small clinical cohorts ( $K=23$ ) that pass four levels of validation: marginal plausibility, cross-visit covariance preservation, conditional subgroup amplification, and mechanistic consistency with a coagulation ODE model calibrated on the same cohort. A downstream utility test showed that a mechanistic model calibrated entirely on synthetic patients predicted real patient outcomes as well as one calibrated on real data. The SA framework preserves the geometry of the clinical data in regimes where parametric approaches such as multivariate normal sampling fail due to rank deficiency. An analytic sampler for the same latent target reproduced the reported implementation’s evaluated behavior and offers a faster route for future use. Combined with a multi-level validation framework, this approach provides a proof-of-concept path for generating synthetic cohorts useful for the modeling use-cases demonstrated here to support

computational research in rare conditions and small longitudinal studies.

## 5 Methods

### 5.1 Population Dataset

We analyzed a longitudinal dataset of coagulation, fibrinolysis, thrombin generation assay (TGA), and rotational thromboelastometry (ROTEM) measurements from women desiring a pregnancy. The study was approved by the Institutional Review Board at the University of Vermont. Written consent was obtained from all participants prior to enrollment. Blood collection, plasma processing, and assay protocols have been described previously<sup>36–39</sup>. Briefly, citrated platelet-poor plasma was collected without tourniquet use after supine rest at three clinical visits: Visit 1 (pre-pregnancy baseline in the follicular phase), Visit 2 (end of first trimester), and Visit 3 (mid-third trimester). Plasma aliquots were stored at  $-80^{\circ}\text{C}$  for subsequent analysis. Measurements of secondary analytes were conducted as follows: activity levels of coagulation factors were determined by Stago clotting assays, fibrinolysis proteome levels were measured by ELISA, and hormone levels were determined by the CLIA-certified laboratory at the University of Vermont Medical Center. TGA was performed using 5 pM relipidated tissue factor with fluorogenic substrate on a Synergy4 plate reader as described in McLean et al.<sup>36</sup>. Clot formation and fibrinolysis dynamics were assessed via viscoelastometry using the ROTEM Delta Instrument (Werfen). Thawed plasma was mixed with or without 4 nM tPA (with respect to plasma volume), recalcified with 15 mM calcium chloride, and incubated for 3 minutes at  $37^{\circ}\text{C}$ . The reaction was initiated by adding the plasma mixture to ROTEM cups containing tissue factor at a 17:1 ratio; the final concentration of TF in the reaction was 5 pM.

The dataset contained 153 total records across approximately 50 subjects. After removing columns with  $> 30\%$  missing values and retaining only subjects with complete data across all three visits, we obtained  $K = 23$  complete longitudinal profiles across  $n = 72$  assay measurements per visit. The 72 assay measurements spanned five categories: (i) hormones (estradiol, progesterone), (ii) coagulation factors and inhibitors (Factors II, V, VII, VIII, IX, X, XI, XII; antithrombin (AT), protein C (PC), tissue factor pathway inhibitor (TFPI), von Willebrand factor (vWF), fibrinogen, etc.), (iii) fibrinolytic markers (plasminogen, plasminogen activator inhibitor-1 (PAI-1), plasminogen activator inhibitor-2 (PAI-2),  $\alpha_2$ -antiplasmin ( $\alpha_2$ -AP), thrombin-activatable fibrinolysis inhibitor (TAFI), tissue plasminogen activator (tPA), D-dimer equivalents), (iv) thrombin generation parameters under four initiator conditions (TF at 5 pM, TF+thrombomodulin (TM), No TF, No TF+anti-FXIa), and (v) ROTEM/viscoelastic parameters (clotting time (CT), maximum clot firmness (MCF), alpha angle, lysis onset time (LOT), maximum lysis (ML), lysis time (LT), area under the curve (AUC) under TF Only and TF+tPA conditions).

Demographics and clinical characteristics of the resulting cohort, including age at enrollment, prepregnancy BMI, parity, gestational ages at each visit, and the cross-tabulation of enrollment cohort against pregnancy outcome, are summarized in Table 1. The 23 patients were drawn from three enrollment conditions: Healthy nulliparous women ( $n=14$ ), women with a personal history of preeclampsia from a previous pregnancy (Prior PE,  $n=6$ ), and women with polycystic ovarian syndrome (PCOS,  $n=3$ ; 2 individuals were nulliparous). For conditional generation, we defined three overlapping subgroups based on pregnancy outcome and comorbidity rather than enrollment condition: Uncomplicated ( $n=18$ , all patients whose study pregnancy did not result in preeclampsia), PCOS ( $n=3$ , a cross-cutting comorbidity; 1 patient also developed PE), and Developed PE ( $n=5$ , patients who developed preeclampsia during the study pregnancy, drawn from 2 healthy-enrolled, 2 prior-PE-enrolled, and 1 PCOS-enrolled participants).

## 5.2 Multiplicity-Weighted Stochastic Attention

We generated synthetic patients using a multiplicity-weighted variant of Stochastic Attention (SA) rooted in modern Hopfield network theory<sup>26</sup>. We stored  $K$  memory patterns as columns of a matrix  $\hat{\mathbf{X}} \in \mathbb{R}^{d \times K}$  and defined a weighted Hopfield energy landscape:

$$E_r(\boldsymbol{\xi}) = \frac{1}{2} \|\boldsymbol{\xi}\|^2 - \frac{1}{\beta} \log \sum_{k=1}^K r_k \exp\left(\beta \mathbf{m}_k^\top \boldsymbol{\xi}\right) \quad (1)$$

where  $\{\mathbf{m}_k\}_{k=1}^K$  are stored memory patterns,  $\beta$  is an inverse temperature parameter controlling the sharpness of retrieval, and  $\{r_k\}_{k=1}^K$  are per-pattern multiplicity weights. When all  $r_k = 1$ , this reduces to the standard modern Hopfield energy; when  $r_k = \rho > 1$  for a designated subset of patterns, the energy landscape is continuously deformed to favor that subset.

This energy determines its induced distribution in closed form. Writing  $\pi_r(\boldsymbol{\xi}) \propto \exp[-\beta E_r(\boldsymbol{\xi})]$  and completing the square in each exponent,

$$\exp\left[-\frac{\beta}{2} \|\boldsymbol{\xi}\|^2 + \beta \mathbf{m}_k^\top \boldsymbol{\xi}\right] = \exp\left(\frac{\beta}{2} \|\mathbf{m}_k\|^2\right) \exp\left[-\frac{\beta}{2} \|\boldsymbol{\xi} - \mathbf{m}_k\|^2\right], \quad (2)$$

so the target is exactly a finite isotropic Gaussian mixture,

$$\pi_r(\boldsymbol{\xi}) = \sum_{k=1}^K q_k \mathcal{N}(\mathbf{m}_k, \beta^{-1} \mathbf{I}), \quad q_k \propto r_k \exp\left(\frac{\beta}{2} \|\mathbf{m}_k\|^2\right). \quad (3)$$

Each stored patient is the center of one component,  $\beta^{-1}$  is the common component variance, and the multiplicity weights enter the mixture probabilities directly. Because the memories used here are normalized to unit length, the exponential factor is shared by every component and  $q_k = r_k / \sum_j r_j$  exactly. The target can therefore be drawn from ancestrally, by sampling  $k$  from  $q$  and then  $\boldsymbol{\xi} \sim \mathcal{N}(\mathbf{m}_k, \beta^{-1} \mathbf{I})$ , with no Markov chain. The cohort analyzed in this work was generated with the Langevin implementation given next, which is what the reported results describe. The analytic sampler was added afterward as a reference, using the identical normalization, magnitude and decoding pipeline (Supplementary Section 5.4).

We sampled from this landscape using an Unadjusted Langevin Algorithm (ULA):

$$\boldsymbol{\xi}_{t+1} = (1 - \alpha) \boldsymbol{\xi}_t + \alpha \hat{\mathbf{X}} \text{softmax}\left(\beta \hat{\mathbf{X}}^\top \boldsymbol{\xi}_t + \log \mathbf{r}\right) + \sqrt{\frac{2\alpha}{\beta}} \boldsymbol{\varepsilon}_t \quad (4)$$

where  $\alpha$  is a step size,  $\boldsymbol{\varepsilon}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ , and  $\log \mathbf{r}$  is the element-wise log of the multiplicity vector. The deterministic component drove  $\boldsymbol{\xi}_t$  toward a multiplicity-weighted combination of stored patterns, while the stochastic component enabled exploration and novelty. The  $\log r_k$  bias in the softmax logits provided a minimal, theoretically grounded mechanism for continuously interpolating between unconditioned generation ( $\rho = 1$ , all patterns equally weighted) and hard subset curation ( $\rho \rightarrow \infty$ , only designated patterns contribute). This is an inference-time operation that requires no retraining of the energy landscape.

The synthetic cohort analyzed in this work was produced by initializing each Langevin chain near a stored memory pattern and retaining the endpoint of a single trajectory of  $T=2000$  steps. Because that choice could in principle shape the generator independently of the target distribution, we characterized it directly with a mixing diagnostic and a four-rung sampler ladder that compares the published scheme against sphere-initialized, burn-in-controlled, thinned and pooled Metropolis-adjusted sampling, and against exact ancestral draws from Eq. (3). The schemes agree closely in

fidelity and in the memorization behavior relevant to privacy, and the analytic reference agrees with them as well, indicating that the published scheme faithfully represents the target law at  $\beta^*$  (Supplementary Fig. S9).

The inverse temperature  $\beta$  controlled the trade-off between retrieval (large  $\beta$ , converging to nearest memory) and generation (small  $\beta$ , uniform mixing). Put plainly,  $\beta$  sets how sharply the sampler commits to one stored patient. When  $\beta$  is small the attention weights are nearly uniform, every stored patient contributes about equally, and samples collapse toward the cohort average. When  $\beta$  is large a single patient dominates and samples reproduce that individual. The useful regime is the transition between the two, where stored patients stop blurring into one average and begin to act as distinct attractors, and  $\beta^*$  is our estimate of where that transition occurs. We identified the critical  $\beta^*$  at the phase transition by locating the inflection of the normalized attention entropy across a logarithmic sweep of  $\beta$  (the entropy definition, sweep parameters and finite-difference procedure are given in the Supplementary Methods). This procedure is domain-agnostic. The phase transition is a geometric property of  $K$  patterns in  $d$  dimensions, not a property of the features they represent. For our dataset ( $K=23$ ,  $d_{\text{PCA}}=18$ ), the empirical inflection yielded  $\beta^* \approx 2.94$ , close to the theoretical prediction  $\beta^* = \sqrt{d} \approx 4.24$ . A sensitivity analysis across  $\beta/\beta^* \in [0.1, 3.0]$  confirmed that the operating point balanced generation novelty against fidelity (Table S1). For weighted sampling, the entropy calculation incorporated the multiplicity bias, yielding a  $\rho$ -dependent  $\beta^*(\rho)$ .

The effective number of patterns contributing to the dynamics was quantified by the participation ratio  $K_{\text{eff}} = (\sum_k r_k)^2 / \sum_k r_k^2$ , which interpolated smoothly between the total pattern count (unconditioned) and the designated subset size (strongly conditioned). This answers a concrete question. If attention is spread unevenly across the 23 stored patients, how many patients is the generator effectively drawing on? Equal weights give 23, and weights concentrated on three patients give a value near three. Reporting  $K_{\text{eff}}$  alongside a conditioned cohort therefore states how much independent material that cohort rests on, which is the quantity that bounds what can be concluded from it. For a target effective fraction  $f_{\text{target}}$  of probability mass on the designated subset, the required multiplicity was  $\rho = f_{\text{target}} \cdot K_{\text{bg}} / (K_{\text{des}} \cdot (1 - f_{\text{target}}))$ , where  $K_{\text{des}}$  and  $K_{\text{bg}}$  are the designated and background pattern counts, respectively.

### 5.3 Pipeline for Continuous Longitudinal Data

Applying SA to continuous longitudinal clinical data required several adaptations beyond the standard Hopfield sampling framework. We first concatenated each patient’s  $n=72$  assay measurements across all three visits into a single vector of dimension  $d_{\text{concat}} = 3n = 216$ , so that each stored memory pattern encoded a complete longitudinal profile and generated synthetic patients would have internally consistent multi-visit trajectories. We then standardized each feature to zero mean and unit variance and applied PCA retaining 95% of the variance, reducing dimensionality from 216 to  $d_{\text{PCA}}=18$ . This compression served two purposes: it removed collinearity among the 216 features and yielded a memory-to-dimension ratio of  $K/d_{\text{PCA}} \approx 1.28$ , placing SA in a favorable operating regime where the number of stored patterns exceeded the dimensionality of the representation.

Standard SA operates on unit-norm patterns, which is appropriate for discrete data (e.g., one-hot protein sequences) but destroys the anisotropic variance structure of continuous clinical measurements, collapsing dispersion across all principal components to a unit sphere. We addressed this with a direction-magnitude decomposition. Before normalization, we recorded each pattern’s PCA-space norm  $r_k = \|\mathbf{m}_k\|$ ; we then ran SA on unit-norm patterns to obtain a directional sample  $\hat{\xi}$ , drew a magnitude  $r$  from the empirical distribution of  $\{r_k\}$ , and reconstructed the rescaled sample as  $\xi_{\text{rescaled}} = r \cdot \hat{\xi}$ . This preserved the directional structure learned by the Hopfield energy

while restoring the natural scale of variation. The rescaled PCA vector was mapped back to the original feature space via inverse PCA and destandardization, then split into three 72-dimensional visit records.

To generate condition-specific cohorts (e.g., PCOS-only), we set the multiplicity  $r_k = \rho$  for patterns belonging to the target subgroup and  $r_k = 1$  for all others, and sampled using the weighted Langevin update (Eq. 4). The multiplicity  $\rho$  was computed to achieve a target effective fraction  $f_{\text{target}} = 0.80$  of softmax attention on the designated subset, chosen to strongly bias generation toward the target subgroup while retaining 20% background influence for interpolation diversity, via  $\rho = f_{\text{target}} \cdot K_{\text{bg}} / (K_{\text{des}} \cdot (1 - f_{\text{target}}))$ . Multiplicity parameters for all three subgroups are listed in Table S2. For the PCOS subgroup ( $K_{\text{des}}=3$ ), this required  $\rho \approx 26.7$ , reducing the effective pattern count to  $K_{\text{eff}} \approx 4.6$ . Magnitudes for the direction-magnitude decomposition were drawn from the condition-specific subset of norms, ensuring that the scale of variation matched the target subpopulation rather than the full cohort.

The complete generation pipeline is summarized in Algorithm 1. The full set of SA hyperparameters used for all experiments is listed in Table S3, with key operating values  $\alpha = 0.01$ ,  $T = 2,000$  Langevin iterations per sample, and random seed 42 for reproducibility. A control experiment comparing ULA with the Metropolis-Adjusted Langevin Algorithm (MALA) confirmed that the discretization bias is negligible at this step size. Across 10 chains of 5,000 iterations the MALA acceptance rate was  $99.9 \pm 0.03\%$ , and mean post-burn-in energies and effective sample sizes were indistinguishable between the two samplers (Supplementary Table S8). The pipeline was implemented in Julia (v1.12) using the packages listed in the code repository. Generating 100 synthetic patients required approximately 0.1 seconds on a standard laptop (Apple M-series), comparable to MVN sampling ( $\approx 0.15$  s), because SA operates in the 18-dimensional PCA space rather than the full 216-dimensional feature space. Each synthetic patient was generated by an independent Langevin chain (no shared state between patients); generation quality was stable across  $T \in [500, 4000]$  iterations (median MRE 1.0–1.6%), confirming that  $T=2,000$  was sufficient for convergence.

## 5.4 Baseline and Validation Framework

As a baseline, we generated synthetic patients by fitting a single multivariate normal (MVN) distribution to the same  $K=23$  concatenated 216-dimensional patient profiles. Critically, both SA and MVN operated on the same concatenated representation. Each patient’s three visits were joined into one vector before generation, so that both methods had the opportunity to capture cross-visit dependencies. For SA, this structure was preserved through the Hopfield energy landscape operating on 18-dimensional PCA projections of the concatenated vectors. For MVN, the  $216 \times 216$  sample covariance matrix had rank at most 22 and was regularized using Ledoit–Wolf shrinkage<sup>21</sup>, which inflated eigenvalues in the 194 null dimensions. We then evaluated both SA and MVN synthetic patients through four progressively stringent validation levels. **Level 1 (Marginal Plausibility)** tested whether individual feature distributions matched, using per-feature means, standard deviations, and known physiological relationships. **Level 2 (Joint Structure)** tested whether the cross-visit covariance was preserved, by comparing full  $216 \times 216$  correlation matrices, eigenvalue spectra, and PCA projections between real, SA, and MVN populations. **Level 3 (Conditional Structure)** tested whether rare subgroups could be amplified while preserving condition-specific signatures, by generating cohorts of 100 patients each from the Uncomplicated ( $n=18$ ), PCOS ( $n=3$ ), and Developed PE ( $n=5$ ) subgroups using multiplicity-weighted sampling. **Level 4 (Mechanistic Consistency)** tested whether synthetic patients produced biologically plausible outputs under an independent mechanistic model. We used the Hockin–Mann BZ2012



---

**Algorithm 1** Multiplicity-Weighted SA for Longitudinal Patient Generation

---

**Require:** Patient data  $\{\mathbf{x}_k^{(v)}\}$  for  $k = 1, \dots, K$  patients,  $v = 1, 2, 3$  visits

**Require:** Multiplicity vector  $\mathbf{r}$  (all ones for unconditioned;  $r_k = \rho$  for designated subset)

**Ensure:**  $N$  synthetic longitudinal patient profiles

```
1: Concatenate:  $\mathbf{x}_k \leftarrow [\mathbf{x}_k^{(1)}; \mathbf{x}_k^{(2)}; \mathbf{x}_k^{(3)}] \in \mathbb{R}^{216}$ 
2: Standardize:  $\mathbf{z}_k \leftarrow (\mathbf{x}_k - \boldsymbol{\mu}) \oslash \boldsymbol{\sigma}$ 
3: PCA:  $\mathbf{m}_k \leftarrow \mathbf{W}^\top \mathbf{z}_k \in \mathbb{R}^{18}$  ▷ 95% variance retained
4: Record norms:  $r_k^{\text{mag}} \leftarrow \|\mathbf{m}_k\|$ ; normalize  $\hat{\mathbf{m}}_k \leftarrow \mathbf{m}_k / r_k^{\text{mag}}$ 
5: Find  $\beta^*$  via entropy inflection on  $\hat{\mathbf{M}} = [\hat{\mathbf{m}}_1, \dots, \hat{\mathbf{m}}_K]$ 
6: for  $i = 1, \dots, N$  do
7:   Initialize  $\hat{\boldsymbol{\xi}}_0 \sim \text{Uniform}(\mathcal{S}^{17})$  ▷ random point on unit sphere
8:   for  $t = 0, \dots, T - 1$  do ▷ Langevin dynamics
9:      $\hat{\boldsymbol{\xi}}_{t+1} \leftarrow (1 - \alpha)\hat{\boldsymbol{\xi}}_t + \alpha \hat{\mathbf{M}} \text{softmax}(\beta^* \hat{\mathbf{M}}^\top \hat{\boldsymbol{\xi}}_t + \log \mathbf{r}) + \sqrt{2\alpha/\beta^*} \boldsymbol{\varepsilon}_t$ 
10:   end for
11:   Draw magnitude:  $r \sim \{r_k^{\text{mag}}\}$  ▷ condition-specific if weighted
12:   Rescale:  $\mathbf{m}_{\text{synth}} \leftarrow r \cdot \hat{\boldsymbol{\xi}}_T / \|\hat{\boldsymbol{\xi}}_T\|$ 
13:   Inverse PCA + destandardize:  $\mathbf{x}_{\text{synth}} \leftarrow \boldsymbol{\sigma} \odot (\mathbf{W} \mathbf{m}_{\text{synth}} + \bar{\mathbf{z}}) + \boldsymbol{\mu}$ 
14:   Split into visit records:  $\mathbf{x}_{\text{synth}}^{(1)}, \mathbf{x}_{\text{synth}}^{(2)}, \mathbf{x}_{\text{synth}}^{(3)}$ 
15: end for
```

---

coagulation model<sup>6,17</sup>, a system of 58 ODEs with 64 rate constants, in which 5 rate constants were calibrated on Visit 1 real patients at the population level and the remaining 59 were held at literature values. We ran the model on all real and synthetic patients under two conditions (TF-only and TF+TM) and compared predicted-versus-measured thrombin generation features. The primary metric was cloud overlap: the fraction of synthetic predicted/measured ratios falling within the 5th–95th percentile range of real ratios.

## Ethics Statement

This study used de-identified secondary analysis measurements from samples collected longitudinally from women desiring a pregnancy. The primary prospective research study was approved by the Institutional Review Board at the University of Vermont. Written informed consent was obtained from all participants prior to enrollment.

## Data Availability

The generated synthetic datasets and all intermediate validation outputs are available in the project repository at <https://github.com/varnerlab/SA-generation-legacy-dataset-paper>. The real patient data were collected under IRB approval at the University of Vermont; de-identified data are available upon reasonable request to the corresponding author.

## Code Availability

All code for synthetic patient generation, validation, figure reproduction, and mechanistic model calibration is available at <https://github.com/varnerlab/SA-generation-legacy-dataset-paper>.

818 The repository includes Julia source code, experiment scripts, and a Python baseline script for CT-  
819 GAN/TVAE comparison.

## 820 Acknowledgments

821 This work was supported by NIH NHLBI R-33 HL 141787 (PIs Bernstein, Orfeo) and NIH NHLBI  
822 R01 HL 71944 (PI Bernstein).

## 823 Author Contributions

824 J.D.V. conceived the study, developed the SA pipeline and validation framework, performed com-  
825 putational analyses, and wrote the manuscript. M.C.B. and T.O. oversaw the secondary analysis  
826 measurements of coagulation and fibrinolysis markers, dynamic assays, and hormone measurements;  
827 provided domain expertise on coagulation biology and thrombin generation assays; and reviewed  
828 the manuscript. C.M. was involved in the enrollment of participants in the prospective research  
829 study and collection of clinical information of the participants. I.B. designed and oversaw the  
830 primary prospective research study of enrolled participants, provided clinical interpretation, and  
831 reviewed the manuscript. All authors approved the final version.

## 832 Competing Interests

833 The authors declare no competing interests.

## 834 References

- 835 [1] Toy, S. U.S. maternal mortality hits highest level since 1965. *The Wall Street Journal* (2023).  
836 March 16, 2023.
- 837 [2] Harris, E. US maternal mortality continues to worsen. *JAMA* **329**, 1248 (2023).
- 838 [3] Chen, Y. *et al.* Pregnancy-related deaths in the US, 2018–2022. *JAMA Network Open* **8**,  
839 e254325 (2025).
- 840 [4] Mann, K. G. Thrombin generation in hemorrhage control and vascular occlusion. *Circulation*  
841 **124**, 225–235 (2011).
- 842 [5] Butenas, S., Orfeo, T. & Mann, K. G. Tissue factor in coagulation: Which? where? when?  
843 *Arteriosclerosis, Thrombosis, and Vascular Biology* **29**, 1989–1996 (2009).
- 844 [6] Hockin, M. F., Jones, K. C., Everse, S. J. & Mann, K. G. A model for the stoichiometric  
845 regulation of blood coagulation. *Journal of Biological Chemistry* **277**, 18322–18333 (2002).
- 846 [7] Bravo, M. C., Orfeo, T., Mann, K. G. & Everse, S. J. Modeling of human factor Va inactivation  
847 by activated protein C. *BMC Systems Biology* **6**, 45 (2012).
- 848 [8] Hellgren, M. Hemostasis during normal pregnancy and puerperium. *Seminars in Thrombosis*  
849 *and Hemostasis* **29**, 125–130 (2003).

- [9] Grouzi, E., Pouliakis, A., Aktypi, A. *et al.* Pregnancy and thrombosis risk for women without a history of thrombotic events: a retrospective study of the real risks. *Thrombosis Journal* **20** (2022).
- [10] Bremme, K. A. Haemostatic changes in pregnancy. *Best Practice & Research Clinical Haematology* **16**, 153–168 (2003).
- [11] Azziz, R. *et al.* The prevalence and features of the polycystic ovary syndrome in an unselected population. *Journal of Clinical Endocrinology & Metabolism* **89**, 2745–2749 (2004).
- [12] Palomba, S., Santagni, S., Falbo, A. & La Sala, G. B. Complications and challenges associated with polycystic ovary syndrome: current perspectives. *International Journal of Women’s Health* **7**, 745–763 (2015).
- [13] Abalos, E. *et al.* Pre-eclampsia, eclampsia and adverse maternal and perinatal outcomes: a secondary analysis of the World Health Organization multicountry survey on maternal and newborn health. *BJOG: An International Journal of Obstetrics & Gynaecology* **121**, 14–24 (2014).
- [14] Dusse, L. M., Rios, D. R. A., Pinheiro, M. B., Cooper, A. J. & Lwaleed, B. A. Pre-eclampsia: relationship between coagulation, fibrinolysis and inflammation. *Clinica Chimica Acta* **412**, 17–21 (2011).
- [15] Luan, D., Zai, M. & Varner, J. D. Computationally derived points of fragility of a human cascade are consistent with current therapeutic strategies. *PLoS Computational Biology* **3**, e142 (2007).
- [16] Luan, D., Szlam, F., Tanaka, K. A., Barie, P. S. & Varner, J. D. Ensembles of uncertain mathematical models can identify network response to therapeutic interventions. *Molecular BioSystems* **6**, 2272–2286 (2010).
- [17] Brummel-Ziedins, K. E. *et al.* The prothrombotic phenotypes in familial protein c deficiency are differentiated by computational modeling of thrombin generation. *PLoS ONE* **7**, e44378 (2012).
- [18] Hastie, T., Tibshirani, R. & Friedman, J. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (Springer, 2009), 2nd edn.
- [19] Gonzales, A., Guruswamy, G. & Smith, S. R. Synthetic data in health care: A narrative review. *PLOS Digital Health* **2**, e0000082 (2023).
- [20] Kühnel, L. *et al.* Synthetic data generation for a longitudinal cohort study – evaluation, method extension and reproduction of published data analysis results. *Scientific Reports* **14**, 14412 (2024).
- [21] Ledoit, O. & Wolf, M. A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis* **88**, 365–411 (2004).
- [22] Goodfellow, I. *et al.* Generative adversarial nets. In *Advances in Neural Information Processing Systems*, vol. 27, 2672–2680 (2014).
- [23] Kingma, D. P. & Welling, M. Auto-encoding variational Bayes. In *International Conference on Learning Representations (ICLR)* (2014).

- [24] Xu, L., Skoularidou, M., Cuesta-Infante, A. & Veeramachaneni, K. Modeling tabular data using conditional GAN. In *Advances in Neural Information Processing Systems*, vol. 32, 7335–7345 (2019).
- [25] De Mitri, O., Wang, R. & Huber, M. F. Generative adversarial networks with limited data: A survey and benchmarking. *arXiv preprint arXiv:2504.05456* (2025).
- [26] Ramsauer, H. *et al.* Hopfield networks is all you need. In *International Conference on Learning Representations (ICLR)* (2021).
- [27] Varner, J. D. Conditioning protein generation via Hopfield pattern multiplicity. *arXiv preprint arXiv:2603.20115* (2026).
- [28] Bravo, M. *et al.* BSTModelKit.jl: Building biochemical systems theory models of individualized coagulation dynamics during pregnancy. JuliaCon 2023, <https://www.youtube.com/watch?v=IGtUqMLHBvQ> (2023). Conference proceedings talk.
- [29] Varner, J. D. Training-free generation of protein sequences from small family alignments via stochastic attention. *arXiv preprint arXiv:2603.14717* (2026).
- [30] Alswaidan, A. & Varner, J. D. Stochastic attention via Langevin dynamics on the modern Hopfield energy. *arXiv preprint arXiv:2603.06875* (2026).
- [31] Mould, D. R. & Upton, R. N. Basic concepts in population modeling, simulation, and model-based drug development—part 2: introduction to pharmacokinetic modeling methods. *CPT: Pharmacometrics & Systems Pharmacology* **2**, e38 (2013).
- [32] Chase, J. G. *et al.* Physiological modeling, tight glycemic control, and the ICU clinician: what are models and how can they affect practice? *Annals of Intensive Care* **1**, 11 (2011).
- [33] Ribba, B. *et al.* A tumor growth inhibition model for low-grade glioma treated with chemotherapy or radiotherapy. *Clinical Cancer Research* **18**, 5071–5080 (2012).
- [34] Yizhak, K., Chaneton, B., Gottlieb, E. & Ruppin, E. Modeling cancer metabolism on a genome scale. *Molecular Systems Biology* **11**, 817 (2015).
- [35] Shan, M., Dai, D., Vudem, A., Varner, J. D. & Stroock, A. D. Multi-scale computational study of the Warburg effect, reverse Warburg effect and glutamine addiction in solid tumors. *PLoS Computational Biology* **14**, e1006584 (2018).
- [36] McLean, K. C., Bernstein, I. M. & Brummel-Ziedins, K. Tissue factor dependent thrombin generation across pregnancy. *American Journal of Obstetrics and Gynecology* **207**, 135.e1–135.e6 (2012).
- [37] Hale, S. A. *et al.* Coagulation and fibrinolytic system protein profiles in women with normal pregnancies and pregnancies complicated by hypertension. *Pregnancy Hypertension* **2**, 152–157 (2012).
- [38] Psinos, R. B. C., Morris, E. A., McBride, C. A. & Bernstein, I. M. Association of pre-pregnancy subclinical insulin resistance with cardiac dysfunction in healthy nulliparous women. *Pregnancy Hypertension* **26**, 11–16 (2021).
- [39] Bernstein, I. M., Hale, S. A. & Badger, G. J. Differences in cardiovascular function comparing prior preeclampsia with nulliparous controls. *Pregnancy Hypertension* **6**, 320–326 (2016).

Table 1: **Demographics and clinical characteristics of the 23 patients with complete longitudinal data across all three visits.** Values are reported as mean  $\pm$  SD, median (range), or  $n$  (%) as appropriate.

| Characteristic                                                 | Mean $\pm$ SD   | Median (Range) | $n$ (%) |
|----------------------------------------------------------------|-----------------|----------------|---------|
| Age at enrollment (yrs)                                        | 30.2 $\pm$ 5.1  | 29 (20–41)     |         |
| Race                                                           |                 |                |         |
| White                                                          |                 |                | 22 (96) |
| Not disclosed                                                  |                 |                | 1 (4)   |
| Prepregnancy BMI                                               | 26.6 $\pm$ 5.3  | 25 (19.0–37.8) |         |
| Parity                                                         |                 | 0 (0–2)        |         |
| Nulliparous                                                    |                 |                | 16 (70) |
| Parous                                                         |                 |                | 7 (30)  |
| Enrollment condition                                           |                 |                |         |
| Healthy nulliparous                                            |                 |                | 14 (61) |
| PCOS                                                           |                 |                | 3 (13)  |
| Prior PE                                                       |                 |                | 6 (26)  |
| Cycle day at V1                                                | 9.2 $\pm$ 3.6   | 10 (3–14)      |         |
| Gestational age at V2 (days)                                   | 89.0 $\pm$ 5.5  |                |         |
| Gestational age at V3 (days)                                   | 217.0 $\pm$ 4.0 |                |         |
| <i>Enrollment cohort <math>\times</math> pregnancy outcome</i> |                 |                |         |
|                                                                | Uncomplicated   | Developed PE   | Total   |
| Healthy nulliparous                                            | 12              | 2              | 14      |
| Prior PE                                                       | 4               | 2              | 6       |
| PCOS                                                           | 2               | 1              | 3       |
| Total                                                          | 18              | 5              | 23      |

Table 2: Per-visit mean relative error (MRE) between real and SA-generated synthetic patients for representative coagulation features, computed as  $|\bar{x}_{\text{synth}} - \bar{x}_{\text{real}}|/|\bar{x}_{\text{real}}|$ .

| Feature            | Visit 1 | Visit 2 | Visit 3 |
|--------------------|---------|---------|---------|
| Factor II (%)      | 0.018   | 0.017   | 0.007   |
| Factor VIII (%)    | 0.018   | 0.008   | 0.005   |
| AT (%)             | 0.025   | 0.001   | 0.008   |
| Fibrinogen (mg/dL) | 0.005   | 0.009   | <0.001  |
| TF Peak (nM)       | 0.025   | 0.002   | 0.017   |
| TF ETP (nM·min)    | 0.013   | 0.008   | 0.023   |






Figure 1: **Physiological correlations are preserved in synthetic patients.** Scatter plots of coagulation factor levels versus thrombin generation parameters across all three visits. Filled markers are real patients; open markers are SA-generated synthetic patients. Colors encode visit: blue (V1/baseline), green (V2/first trimester), orange (V3/third trimester). Real and synthetic patients occupy the same joint regions across all four pairwise relationships, confirming that biologically meaningful dependencies are preserved: AT inversely correlates with TF-initiated peak thrombin (top left), FVIII positively correlates with peak (top right), and both FII (bottom left) and AT (bottom right) show the expected relationships with ETP. The visit-dependent shift in these relationships, reflecting the pregnancy-driven increase in procoagulant factors, is also preserved in the synthetic cohort.




Figure 2: **Pregnancy-driven longitudinal trajectories are reproduced.** Mean  $\pm$  standard deviation bands for six key coagulation features across visits (BL = baseline, 1st = first trimester, 3rd = third trimester) for real (blue) and SA-generated synthetic (orange) patients. The characteristic pregnancy-driven increases in fibrinogen, Factor VIII, and vWF are captured, as are the stable-to-declining patterns in Factor II and AT.




Figure 3: **Cross-visit correlation structure.** Top row: full  $216 \times 216$  Pearson correlation matrices for Real ( $K=23$ , left), SA ( $N=100$ , center), and MVN ( $N=100$ , right) populations. Each matrix is organized as a  $3 \times 3$  grid of  $72 \times 72$  blocks (delineated by black lines), where the diagonal blocks capture within-visit feature correlations (V1–V1, V2–V2, V3–V3) and the off-diagonal blocks capture cross-visit dependencies (e.g., V1–V3: does a patient’s baseline Factor VIII predict their third-trimester Factor VIII?). SA preserves both the within-visit and cross-visit block structure. Bottom row: pairwise residual matrices (element-wise subtraction), all plotted on the same  $\pm 1$  color scale as the top row. The SA–Real residual shows small, unstructured deviations. The MVN–Real residual is smoother, particularly in the off-diagonal blocks, reflecting the Ledoit–Wolf regularization that systematically shrinks cross-visit correlations toward zero. The SA–MVN residual confirms that the two methods differ most in these off-diagonal blocks, where SA retains longitudinal dependencies that MVN’s rank-deficient covariance estimate cannot represent. A version with amplified residual color scale ( $\pm 0.5$ ) and Frobenius norms is provided in Fig. S1.





Figure 4: **Recovery across data regimes.** (A) On a low-rank Gaussian ground truth with known population covariance, the ratio of multivariate-Gaussian to SA covariance-recovery error (relative Frobenius distance to the population covariance) across sample size  $n$  and dimensionality  $p$ , averaged over latent ranks. Values above one (blue) mark cells where SA recovers the population covariance more accurately than a multivariate Gaussian; the advantage is largest in the rank-deficient  $n < p$  regime that contains the real cohort’s shape ( $n=23$ ,  $p=216$ ) and narrows as  $n$  approaches  $p$ , where the Gaussian is well specified. Because the target covariance is fixed in advance, this measures generalization rather than reproduction of the training draw. (B) On the real cohort, SA’s recovery of the full 23-patient cross-visit correlation structure as the cohort is subsampled to  $n=20, 15, 10, 8$  (mean and standard deviation over ten random subsets, black bars), degrading gracefully toward the full-cohort value (star) rather than collapsing.



**Figure 5: Condition-specific feature preservation.** Grouped bar charts comparing real (dark bars) and SA-generated synthetic (light bars) patient means ( $\pm$  SD error bars) for eight coagulation features, shown separately for each clinical subgroup: Uncomplicated ( $n=18$ , left, blue), PCOS ( $n=3$ , center, orange), and Developed PE ( $n=5$ , right, red). Each bar is the mean of that assay pooled across all three visits (V1–V3); factor activities (FVIII, vWF, FIX, FV,  $\alpha_2$ -AP, FVII, plasminogen) are expressed as % of normal pooled reference plasma, and fibrinogen (Fbgn) in mg/dL. Gray percentages indicate the mean relative error (MRE). Green annotations show the bootstrap Mann–Whitney equivalence result: the fraction of 1,000 bootstrap replicates (synthetic subsampled to  $n_{\text{real}}$ ) in which  $p > 0.05$ , i.e., real and synthetic could not be distinguished. Features achieving  $\geq 90\%$  are shown in green;  $< 90\%$  in red. Across all 24 feature–condition pairs, 20 (83%) were statistically indistinguishable in  $\geq 90\%$  of replicates (median: 98.6%). The four pairs below this threshold were FIX and plasminogen in the PCOS and Developed PE subgroups, where small subgroup sizes and large inter-patient variability reduced test power.





Figure 6: **Mechanistic validation (TF-only).** Top row: BZ2012 ODE-predicted TGA values (vertical axis) versus dataset TGA values (horizontal axis) for five thrombin generation parameters. The dataset values are the TGA measurements from each patient’s record, present for both real and synthetic patients. The ODE-predicted values are computed by running each patient’s coagulation factor levels through the 58-species BZ2012 model. The dashed line indicates where the ODE model would perfectly reproduce the dataset value ( $y=x$ ); systematic deviations from this line (e.g., lagtime underprediction) reflect ODE model bias, not synthetic data failure. Filled markers are real patients, open markers are SA-generated synthetic patients, colored by visit: blue (V1/baseline, training set), green (V2/first trimester), orange (V3/third trimester). Real and synthetic patients occupy the same cloud with the same pattern of ODE model bias. Bottom row: ODE-predicted/dataset ratio distributions for real (blue) and synthetic (orange) patients. If SA had generated biologically implausible factor combinations, the ODE model would process them differently, producing divergent ratio distributions. Instead, the distributions overlap substantially. Each panel is annotated with cloud overlap (OL), the two-sample Kolmogorov-Smirnov statistic ( $D$ ), and  $p$ -value. All five features show  $p > 0.35$ , confirming that the ODE model cannot distinguish real from synthetic patients.





**Figure 7: Downstream utility: synth-calibrated vs. real-calibrated mechanistic model predictions on held-out real patients.** Each panel compares the BZ2012 ODE predictions for one TGA feature when the model is calibrated on real V1 patients (horizontal axis) versus synthetic V1 patients (vertical axis), evaluated on the same held-out real V2 and V3 patients. Points near the  $y=x$  line indicate that the two calibrations produce equivalent predictions. Green: V2 (first trimester); orange: V3 (third trimester). The synth-calibrated model matched or outperformed the real-calibrated model on all five features (overall ratio 0.94 $\times$ ), demonstrating that a mechanistic model trained entirely on synthetic data generalizes to real patients as well as one trained on real data. Both calibrations used 11 random restarts; all starts converged.

## Supplementary Information

### Supplementary Tables

Table S1: **Inverse temperature sensitivity analysis.** SA generation quality as a function of  $\beta/\beta^*$ , where  $\beta^* \approx 2.94$  was identified via entropy inflection in the 18-dimensional PCA space ( $K=23$  stored patterns). Operating at  $\beta = \beta^*$  (bold row) balances generation novelty against fidelity. Med MRE = median relative error of per-feature means; Novelty =  $1 - \max_k \cos(\hat{\xi}, \hat{m}_k)$ .

| $\beta/\beta^*$ | $\beta$     | Noise scale  | PC1 std ratio | Novelty      | Med MRE (%) |
|-----------------|-------------|--------------|---------------|--------------|-------------|
| 0.1             | 0.29        | 0.261        | 0.538         | 0.542        | 0.7         |
| 0.3             | 0.88        | 0.151        | 0.550         | 0.536        | 0.7         |
| 0.5             | 1.47        | 0.117        | 0.562         | 0.528        | 0.7         |
| <b>1.0</b>      | <b>2.94</b> | <b>0.082</b> | <b>0.589</b>  | <b>0.502</b> | <b>0.6</b>  |
| 1.5             | 4.41        | 0.067        | 0.613         | 0.474        | 0.7         |
| 2.0             | 5.88        | 0.058        | 0.638         | 0.446        | 0.7         |
| 3.0             | 8.82        | 0.048        | 0.689         | 0.403        | 0.8         |

Table S2: **Multiplicity weighting parameters for conditional generation.** The multiplicity ratio  $\rho$  was computed to achieve a target effective fraction  $f_{\text{target}} = 0.80$  of softmax attention on the designated patterns.  $K_{\text{eff}}$  is the participation ratio quantifying the effective number of patterns contributing to generation.

| Condition     | $K_{\text{sub}}$ | $K_{\text{bg}}$ | $\rho$ | $f_{\text{eff}}$ | $K_{\text{eff}}$ |
|---------------|------------------|-----------------|--------|------------------|------------------|
| Uncomplicated | 18               | 5               | 1.11   | 0.80             | 23.0             |
| PCOS          | 3                | 20              | 26.67  | 0.80             | 4.6              |
| Developed PE  | 5                | 18              | 14.40  | 0.80             | 7.7              |

Table S3: **SA generation hyperparameters.** The same values were used for unconditioned and conditioned generation except where noted.

| Parameter           | Value | Description                                  |
|---------------------|-------|----------------------------------------------|
| $\alpha$            | 0.01  | Langevin step size (see Table S8)            |
| $T$                 | 2,000 | Langevin iterations per sample               |
| $\beta^*$           | 2.94  | Inverse temperature (entropy inflection)     |
| PCA threshold       | 95%   | Variance retained                            |
| $d_{\text{PCA}}$    | 18    | PCA dimensionality                           |
| $N$                 | 100   | Synthetic patients per cohort                |
| Random seed         | 42    | For reproducibility                          |
| $f_{\text{target}}$ | 0.80  | Target attention fraction (conditioned only) |

Table S4: **BZ2012 ODE model calibration.** Five rate constants were calibrated on Visit 1 real patients ( $n=23$ ) using bounded Nelder–Mead optimization in log-scale-factor space. The cost function was the mean normalized relative squared error across five TGA features under TF-only conditions. Bounds:  $\pm \log(100)$ . Convergence:  $f_{\text{tol}} = 10^{-6}$ , 500 iterations, 11 random restarts.

| Rate constant                   | Role                   | Default               | Calibrated             | Scale         |
|---------------------------------|------------------------|-----------------------|------------------------|---------------|
| prothrombinase $k_{\text{cat}}$ | Thrombin burst         | $63.5 \text{ s}^{-1}$ | $21.94 \text{ s}^{-1}$ | $0.35\times$  |
| intrinsic Xase $k_{\text{cat}}$ | Amplification Xa       | $8.2 \text{ s}^{-1}$  | $0.169 \text{ s}^{-1}$ | $0.021\times$ |
| extrinsic Xase $k_{\text{cat}}$ | Initiation Xa          | $6.0 \text{ s}^{-1}$  | $93.2 \text{ s}^{-1}$  | $15.5\times$  |
| PC activation $k_{\text{cat}}$  | Protein C activation   | $0.41 \text{ s}^{-1}$ | $0.890 \text{ s}^{-1}$ | $2.17\times$  |
| mIIa conversion $k$             | mIIa $\rightarrow$ IIa | $2.3\times 10^8$      | $1.15\times 10^9$      | $5.01\times$  |

Units for mIIa conversion:  $\text{M}^{-1}\text{s}^{-1}$

Table S5: **Summary of mean relative errors across all 72 features.** Per-visit distribution of MREs between real and SA-generated synthetic patient means. The full feature-by-feature table (72 features  $\times$  3 visits = 216 entries) is provided in the code repository as `paper_summary_statistics.csv`.

| Visit                        | Median MRE   | Mean MRE     | 25th pctl    | 75th pctl    | Max MRE      | Below 5%               |
|------------------------------|--------------|--------------|--------------|--------------|--------------|------------------------|
| V1 (baseline)                | 0.022        | 0.029        | 0.008        | 0.039        | 0.158        | 59/72 (81.9%)          |
| V2 (1st tri)                 | 0.009        | 0.021        | 0.004        | 0.022        | 0.224        | 65/72 (90.3%)          |
| V3 (3rd tri)                 | 0.011        | 0.016        | 0.005        | 0.024        | 0.062        | 69/72 (95.8%)          |
| <b>Pooled (all 3 visits)</b> | <b>0.012</b> | <b>0.022</b> | <b>0.005</b> | <b>0.028</b> | <b>0.224</b> | <b>193/216 (89.4%)</b> |

Pooled median MRE 95% bootstrap CI: (0.98%, 1.60%) over 10,000 resamples (seed 42).

Table S6: **Comparison with deep generative baselines.** CTGAN and TVAE were trained on the same 90 per-visit records ( $23 \text{ patients} \times 3 \text{ visits}$ , 72 features) at three epoch counts. SA and MVN results from the main text are shown for reference. CTGAN fails at this sample size (median MRE  $\approx 19\%$ ), consistent with GAN mode collapse on small datasets. TVAE achieves comparable marginal fidelity at high epoch counts but cannot capture cross-visit covariance or perform conditional generation.

| Method | Epochs | Med MRE | Med KS | Corr MAE |
|--------|--------|---------|--------|----------|
| SA     | –      | 0.019   | 0.169  | 0.200    |
| MVN    | –      | 0.024   | 0.161  | 0.090    |
| CTGAN  | 300    | 0.189   | 0.364  | 0.261    |
| CTGAN  | 1,000  | 0.195   | 0.349  | 0.265    |
| CTGAN  | 3,000  | 0.187   | 0.341  | 0.180    |
| TVAE   | 300    | 0.037   | 0.174  | 0.152    |
| TVAE   | 1,000  | 0.026   | 0.248  | 0.082    |
| TVAE   | 3,000  | 0.018   | 0.224  | 0.092    |

Table S7: **PCA variance threshold sensitivity analysis.** SA generation quality is stable across thresholds from 85% to 99%. Median MRE remains between 1.0% and 1.7% across all thresholds. The 95% threshold used in the main text (bold) is not a critical design choice.

| Threshold  | $d_{\text{PCA}}$ | $K/d$       | $\beta^*$   | Novelty      | Med MRE       | Corr MAE     |
|------------|------------------|-------------|-------------|--------------|---------------|--------------|
| 85%        | 13               | 1.77        | 2.94        | 0.420        | 0.0098        | 0.124        |
| 90%        | 15               | 1.53        | 2.94        | 0.477        | 0.0160        | 0.126        |
| <b>95%</b> | <b>18</b>        | <b>1.28</b> | <b>2.94</b> | <b>0.500</b> | <b>0.0124</b> | <b>0.143</b> |
| 97.5%      | 20               | 1.15        | 2.94        | 0.543        | 0.0120        | 0.135        |
| 99%        | 21               | 1.10        | 2.94        | 0.541        | 0.0172        | 0.138        |

Table S8: **ULA vs MALA sampling diagnostics.** Ten chains of 5,000 iterations each were run with both the Unadjusted Langevin Algorithm (ULA) and the Metropolis-Adjusted Langevin Algorithm (MALA) at step size  $\alpha = 0.01$  and  $\beta^* = 2.94$  on the 18-dimensional PCA memory ( $K=23$  patterns). The MALA acceptance rate of 99.9% confirms that virtually every ULA proposal satisfies detailed balance, and the indistinguishable energy and effective sample size statistics confirm negligible discretization bias. [This comparison bounds the discretization error of the reported sampler; the direct reference for the target itself is the analytic sampler of Eq. \(3\), introduced below.](#)

| Metric                     | ULA               | MALA              |
|----------------------------|-------------------|-------------------|
| Acceptance rate (%)        | – (no rejection)  | $99.9 \pm 0.03$   |
| Mean energy (post-burn-in) | $1.932 \pm 0.141$ | $1.886 \pm 0.231$ |
| $\tau_{\text{int}}$        | $109.5 \pm 49.5$  | $106.7 \pm 30.4$  |
| Effective sample size      | $41.8 \pm 14.2$   | $40.0 \pm 10.2$   |

Burn-in: 1,000 iterations discarded.  $\tau_{\text{int}}$ : integrated autocorrelation time estimated with a windowed estimator (threshold 0.05).  $\text{ESS} = (T - T_{\text{burn}})/\tau_{\text{int}}$ .

Table S9: **Memorization diagnostics.** Two complementary checks that SA-generated patients are not memorized copies of the  $K=23$  stored profiles. Novelty score is the angular distance between a synthetic sample and its nearest stored pattern at  $\beta = \beta^*$ . Nearest-neighbor distances are computed in the per-visit standardized 72-dimensional feature space.

| Metric                                                      | Value |
|-------------------------------------------------------------|-------|
| Mean novelty score, $1 - \max_k \cos(\hat{\xi}, \hat{m}_k)$ | 0.50  |
| Median synth–real nearest-neighbor distance (std. units)    | 6.14  |
| Median real–real nearest-neighbor distance (std. units)     | 6.87  |
| Distance ratio (synth–real / real–real)                     | 0.89  |

Table S10: **Bootstrap Mann–Whitney equivalence per feature and condition.** For each of 24 feature–condition pairs, we subsampled the synthetic cohort to match the real sample size, applied a Mann–Whitney U test, and repeated 1,000 times. Frac.  $p > 0.05$  is the fraction of replicates in which the test could not distinguish the synthetic from the real distribution. 20 of 24 pairs achieve  $p > 0.05$  in  $\geq 90\%$  of replicates; the four failing pairs (italicized) all involve high-variance features in the smallest subgroups.

| Condition                | Feature        | MRE          | Frac. $p > 0.05$ |
|--------------------------|----------------|--------------|------------------|
| Uncomplicated ( $n=18$ ) | FVIII          | 0.057        | 0.952            |
| Uncomplicated ( $n=18$ ) | vWF            | 0.022        | 0.999            |
| Uncomplicated ( $n=18$ ) | FIX            | 0.018        | 0.998            |
| Uncomplicated ( $n=18$ ) | FV             | 0.006        | 0.996            |
| Uncomplicated ( $n=18$ ) | $\alpha_2$ -AP | 0.019        | 1.000            |
| Uncomplicated ( $n=18$ ) | FVII           | 0.057        | 0.941            |
| Uncomplicated ( $n=18$ ) | Fbgn           | 0.017        | 0.987            |
| Uncomplicated ( $n=18$ ) | Plgn           | 0.020        | 0.988            |
| PCOS ( $n=3$ )           | FVIII          | 0.109        | 0.985            |
| PCOS ( $n=3$ )           | vWF            | 0.150        | 0.959            |
| PCOS ( $n=3$ )           | <i>FIX</i>     | <i>0.121</i> | <i>0.873</i>     |
| PCOS ( $n=3$ )           | FV             | 0.057        | 0.944            |
| PCOS ( $n=3$ )           | $\alpha_2$ -AP | 0.051        | 0.987            |
| PCOS ( $n=3$ )           | FVII           | 0.029        | 0.999            |
| PCOS ( $n=3$ )           | Fbgn           | 0.043        | 0.992            |
| PCOS ( $n=3$ )           | <i>Plgn</i>    | <i>0.153</i> | <i>0.615</i>     |
| Developed PE ( $n=5$ )   | FVIII          | 0.097        | 0.932            |
| Developed PE ( $n=5$ )   | vWF            | 0.010        | 0.986            |
| Developed PE ( $n=5$ )   | <i>FIX</i>     | <i>0.108</i> | <i>0.780</i>     |
| Developed PE ( $n=5$ )   | FV             | 0.006        | 0.987            |
| Developed PE ( $n=5$ )   | $\alpha_2$ -AP | 0.062        | 0.950            |
| Developed PE ( $n=5$ )   | FVII           | 0.045        | 0.993            |
| Developed PE ( $n=5$ )   | Fbgn           | 0.048        | 0.996            |
| Developed PE ( $n=5$ )   | <i>Plgn</i>    | <i>0.093</i> | <i>0.814</i>     |

Median Frac.  $p > 0.05$  across all 24 pairs: 0.986. Failing pairs ( $< 0.90$ , italicized): PCOS–FIX, PCOS–Plgn, DPE–FIX, DPE–Plgn.

Table S11: **Mechanistic validation diagnostics per feature and condition.** Cloud overlap is the fraction of synthetic ODE-predicted/measured ratios falling within the 5th–95th percentile of the corresponding real distribution. KS  $D$  and KS  $p$  are from a two-sample Kolmogorov–Smirnov test on the same ratio distributions. Real  $n=69$  (23 patients  $\times$  3 visits); synthetic  $n=300$  (100 patients  $\times$  3 visits).

| Condition | TGA Feature | Cloud overlap | KS $D$ | KS $p$ |
|-----------|-------------|---------------|--------|--------|
| TF-only   | Lagtime     | 0.92          | 0.112  | 0.48   |
| TF-only   | Peak        | 0.91          | 0.081  | 0.86   |
| TF-only   | T.Peak      | 0.92          | 0.102  | 0.61   |
| TF-only   | Max Rate    | 0.91          | 0.083  | 0.84   |
| TF-only   | ETP         | 0.83          | 0.123  | 0.36   |
| TF+TM     | Lagtime     | 0.92          | 0.127  | 0.32   |
| TF+TM     | Peak        | 0.88          | 0.079  | 0.88   |
| TF+TM     | T.Peak      | 0.91          | 0.110  | 0.51   |
| TF+TM     | Max Rate    | 0.88          | 0.096  | 0.67   |
| TF+TM     | ETP         | 0.89          | 0.112  | 0.48   |

TF-only cloud overlap range: 0.83–0.92. TF+TM cloud overlap range: 0.88–0.92. Pooled across both conditions and all five features, 0.902. KS  $D$  range across both conditions: 0.079–0.127, all  $p > 0.30$ .

Table S12: **Downstream utility — per-feature held-out V2/V3 performance.** Median relative error of the BZ2012 ODE-predicted TGA values against the held-out real V2 and V3 measurements, for the model calibrated on real V1 patients ( $K=23$ ) and the model calibrated on synthetic V1 patients ( $N=100$ ). The synth-calibrated model achieves slightly lower error on every feature. Synth/Real ratio  $< 1$  favors the synth-calibrated model.

| TGA Feature           | Real-cal MRE | Synth-cal MRE | Synth/Real                     |
|-----------------------|--------------|---------------|--------------------------------|
| Lagtime               | 0.165        | 0.162         | 0.98 $\times$                  |
| Peak                  | 0.097        | 0.087         | 0.90 $\times$                  |
| T.Peak                | 0.330        | 0.306         | 0.93 $\times$                  |
| Max Rate              | 0.286        | 0.259         | 0.91 $\times$                  |
| ETP                   | 0.197        | 0.183         | 0.93 $\times$                  |
| <b>Overall median</b> | <b>0.201</b> | <b>0.189</b>  | <b>0.94<math>\times</math></b> |

Table S13: **PCA-space ablation.** Five generators fit in the identical  $d=18$  PCA subspace and evaluated on the same harness as SA. Median relative error (MRE) and the mechanistic cloud overlap measure fidelity; median novelty measures distance from the stored patterns (Frac. novel is the fraction above the 0.2 novelty threshold); the cross-visit Frobenius error is the deviation of the synthetic correlation matrix from the real one, where a lower value means staying closer to the 23 real patients and therefore rewards memorization rather than generalization. The shared subspace is necessary but not sufficient.  $k$ -nearest-neighbor interpolation collapses novelty and the kernel-density estimator degrades marginal and mechanistic fidelity despite both being fit in it, while the mixture and copula baselines are competitive with SA on the fidelity axes.

| Method     | MRE (%) | Cross-visit Frob. | Mech. overlap | Median novelty | Frac. novel |
|------------|---------|-------------------|---------------|----------------|-------------|
| SA         | 1.24    | 0.641             | 0.902         | 0.514          | 1.00        |
| PCA+GMM    | 1.13    | 0.422             | 0.918         | 0.497          | 0.94        |
| PCA+KDE    | 1.94    | 0.337             | 0.823         | 0.249          | 0.65        |
| PCA+kNN    | 1.63    | 0.587             | 0.938         | 0.051          | 0.05        |
| PCA+Copula | 1.34    | 0.417             | 0.934         | 0.462          | 1.00        |

Table S14: **Convex-hull membership does not discriminate at  $K=23$  in  $d=18$ .** Hull membership and distance for the 100 synthetic patients against the 23 real patients, and for each real patient against the other 22. The radial factor  $t^*$  is where the ray from the hull centroid through a point exits the hull, so  $1/t^*$  is the factor by which a point overshoots the boundary and  $t^* \geq 1$  would mean the point is inside. Distances and overshoots are severity-matched, each point’s nearest vertex having been removed on both sides, because a held-out real patient necessarily loses its own vertex. Matched this way the distance to the hull is not distinguishable between synthetic and real patients (Mann–Whitney  $p = 0.07$ ), while the radial overshoot is smaller for the synthetic ones ( $p < 10^{-4}$ ); both tests treat the 100 synthetic patients as independent, which they are not, so the reported significance is optimistic. The geometry rows show why no point of either kind falls inside. The hull attains the patients’ typical radius only along the 23 vertex directions, and the sampler returns every sample to the unit sphere, where a point inside the hull of unit-norm patterns would have to be one of those patterns.

| Quantity                                                  | Synthetic ( $n=100$ ) | Real, leave-one-out ( $n=23$ ) |
|-----------------------------------------------------------|-----------------------|--------------------------------|
| Fraction inside hull                                      | 0.00                  | 0.00                           |
| Median distance to hull                                   | 11.0                  | 11.9                           |
| Median radial overshoot $1/t^*$                           | 10.5                  | 15.2                           |
| <i>Hull geometry of the real cohort</i>                   |                       |                                |
| Median patient radius                                     |                       | 13.8                           |
| Median hull extent, generic direction                     |                       | 2.0                            |
| Largest $t^*$ attained on the unit sphere                 |                       | 0.231                          |
| <i>Plausibility of the out-of-hull synthetic patients</i> |                       |                                |
| Biological constraints satisfied                          |                       | 81/100                         |
| Mechanistic band, per feature                             |                       | 83 to 92%                      |
| Mechanistic band, all features and visits                 |                       | 41/100                         |



Table S15: **Membership inference recovers memory membership, not unseen patients.** Median nearest-neighbor distances in the standardized 216-dimensional space, and the attack’s discrimination. The attack retraines the generator on 15 training patients and scores each real record by its distance to the resulting synthetic set, treating the 15 training patients as members and the 8 held-out patients as non-members. Every quantity comes from the same protocol, 40 partitions drawn at random with the generation seed held fixed so that only the partition varies, with the distances pooled across all partitions. Non-members sit at the same distance from the synthetic cohort as unrelated real patients sit from one another, so the synthetic cohort carries no proximity signal about patients the generator never saw and the discrimination comes entirely from members sitting closer.

| Quantity                                              | Value |
|-------------------------------------------------------|-------|
| <i>Median nearest-neighbor distance</i>               |       |
| Training member to nearest synthetic                  | 10.7  |
| Held-out non-member to nearest synthetic              | 15.4  |
| Real to nearest real, leave-one-out                   | 15.9  |
| Synthetic to nearest real                             | 13.8  |
| <i>Membership-inference AUC over 40 random splits</i> |       |
| Mean                                                  | 0.97  |
| Standard deviation                                    | 0.02  |
| Median                                                | 0.98  |
| Minimum                                               | 0.90  |
| Maximum                                               | 1.00  |
| Splits reaching perfect separation                    | 9/40  |
| Splits with AUC above 0.9                             | 39/40 |

## 930 Supplementary Figures




Figure S1: **Cross-visit correlation structure (amplified residual scale)**. Same layout as the main-text correlation figure, but with the bottom-row residual panels plotted on a  $\pm 0.5$  color scale (vs.  $\pm 1.0$  in the main text) to reveal fine structure. Frobenius norms of each residual matrix are shown in the panel titles.

## 931 Supplementary Methods: locating the critical inverse temperature

932 The operating point  $\beta^*$  was located from the attention entropy. For a candidate  $\beta$  we computed the  
 933 Shannon entropy of the attention weights  $\text{softmax}(\beta \hat{\mathbf{X}}^\top \boldsymbol{\xi})$  evaluated at 20 probe points placed at  
 934 stored patterns, averaged those entropies, and normalized by  $\log K$  so that the resulting curve runs  
 935 from 0, when one pattern carries all of the weight, to 1, when every pattern carries equal weight.  
 936 Sweeping  $\beta$  logarithmically over 80 values from 0.1 to 1000 traces this curve from the diffuse regime  
 937 into the concentrated one. We took  $\beta^*$  to be the point of most negative second derivative of the  
 938 normalized entropy with respect to  $\log \beta$ , computed by finite differences over the sweep, which is  
 939 the sharpest bend in the curve and marks where attention stops being shared across patterns and  
 940 begins to localize on one. For multiplicity-weighted sampling the same computation was carried  
 941 out with the  $\log r_k$  bias included in the softmax logits, which gives a  $\rho$ -dependent  $\beta^*(\rho)$  because the  
 942 weighting changes how quickly attention concentrates as  $\beta$  grows. For random unit-norm patterns  
 943 the transition is predicted at  $\beta^* = \sqrt{d}$ , the theoretical reference quoted in the main text.




Figure S2: **Covariance eigenvalue spectrum.** Log-scale eigenvalues of the  $216 \times 216$  sample covariance for Real ( $K=23$ , rank 22), SA ( $N=100$ , rank 18), and MVN ( $N=100$ ) populations. Dashed vertical lines mark the SA PCA truncation point (component 18) and the data rank limit (component 22). MVN’s Ledoit–Wolf regularization inflates eigenvalues beyond component 22 (green shaded region), introducing spurious variance in 194 dimensions where the data contain no signal.




Figure S3: **Conditioned generation: PCA projection (pooled)**. Large markers: real patients (blue = Uncomplicated,  $n=18$ ; orange = PCOS,  $n=3$ ; red = Developed PE,  $n=5$ ). Small markers: synthetic patients ( $N=100$  per condition). With only 3 PCOS and 5 Developed PE patients, the subgroups do not form clearly separable clusters in the first two principal components, but the synthetic cohorts concentrate around their respective real counterparts.




Figure S4: **Conditioned generation: PCA projection by visit**. Same as Fig. S3 but separated by visit (V1, V2, V3). Condition-specific clustering is maintained within each individual visit.




Figure S5: **Mechanistic validation (TF+TM condition)**. Same layout as main-text Fig. 6 but under TF+TM conditions. The model systematically underpredicts peak and maximum rate ( $\approx 0.5\times$  measured) for both real and synthetic patients, reflecting overcorrection of the protein C anticoagulant pathway. Cloud overlap remains high at 88–92%, confirming that the model processes real and synthetic patients identically even under imperfect calibration conditions.




Figure S6: **Cross-visit calibration generalization**. Boxplots of the ODE-predicted/dataset ratio by visit for real patients. The model was calibrated on Visit 1 only. Top: TF-only conditions, where ratios near 1.0 at Visits 2 and 3 indicate that the calibration generalizes across pregnancy timepoints. Bottom: TF+TM conditions, where the calibration does not generalize as well, particularly for peak and maximum rate.




Figure S7: **Spearman rank correlations between ODE-predicted and dataset TGA features.** Left: TF-only. Right: TF+TM. ETP is the best-predicted feature ( $\rho \approx 0.6$ – $0.8$  for real patients under TF-only). Weak rank correlations for other features are expected given the 5-parameter population-level calibration.




Figure S8: **PCA projections by visit: SA vs. MVN.** Each panel shows the first two principal components (PC1: 27.7% variance, PC2: 15.4%) computed from standardized per-visit real patient data ( $K=23$  patients per visit, 72 features). Dark markers indicate real patients; lighter markers indicate synthetic patients ( $N=100$ ). Within each visit (column), the SA panel (top) and MVN panel (bottom) share identical axis ranges, so the relative dispersion of the two methods can be read against the same real data cloud. Top row: SA-generated patients (circles) cluster tightly around the real data cloud at all three visits. Bottom row: MVN-generated patients (diamonds) show greater dispersion, particularly at Visits 2 (first trimester) and 3 (third trimester), where MVN patients extend beyond the convex hull of the real data. This excess scatter reflects the spurious variance introduced by Ledoit–Wolf regularization of the rank-deficient covariance (Fig. S2). Colors encode visit: blue (V1, not pregnant), green (V2, first trimester), orange (V3, third trimester).



## 944 Sampling correctness and privacy diagnostics

945 We assessed re-identification risk directly. For every synthetic patient we computed the distance  
 946 to the closest real record (DCR) in the standardized 216-dimensional space and compared the  
 947 synthetic-to-real distribution against the leave-one-out real-to-real distribution. Synthetic patients  
 948 sit slightly closer to the real cohort than the real patients sit to one another, with median DCR  
 949 13.8 against 15.9. A nearest-neighbor membership-inference attack, which scores each real record  
 950 by its proximity to the synthetic set and is evaluated over repeated random train/holdout splits,  
 951 separates the 15 training patients from the 8 held-out patients with a mean AUC of 0.97 (standard  
 952 deviation 0.02 over 40 splits). At the published operating point the generator therefore reproduces  
 953 its training cohort closely enough to be identifiable.

954 To attribute this behavior we asked whether it reflects the target distribution  $\pi(\xi) \propto \exp(-\beta E)$   
 955 or merely the sampler used to draw from it. A mixing diagnostic showed that at  $\beta^*$  the target  
 956 is multimodal, with one component centered on each of the  $K=23$  stored patterns, and that dis-  
 957 cretization bias in the Langevin updates is negligible, the Metropolis acceptance being near 0.999  
 958 (Supplementary Fig. S10). We then constructed a four-rung ladder (Table S16, Fig. S9) that holds  
 959  $\beta$ , the memory, and the decoding fixed and varies only the sampler, from the published anchor-  
 960 initialized single-endpoint scheme to sphere initialization, to burn-in-controlled thinned pooling, to  
 961 Metropolis-adjusted sampling. Across the ladder the membership-inference AUC stays near 0.96  
 962 to 0.97 and the median DCR rises only from 14.0 to 14.3, never reaching the real-to-real baseline of  
 963 15.9, at a modest cost in marginal fidelity. Proper sampling of  $\pi$  does not remove the memorization;  
 964 it is a property of the distribution at this cohort size.

965 The remaining lever is the inverse temperature. Lowering  $\beta$  spreads probability mass away  
 966 from the stored patterns and should improve privacy, at the expense of the memorization that also  
 967 underlies fidelity. Sweeping  $\beta$  across a 40-fold range around  $\beta^*$  (Fig. S9C) confirmed the trade-off  
 968 but bounded it sharply: the median DCR rises monotonically as  $\beta$  falls yet saturates near 14.3,  
 969 remaining well below the real-to-real baseline of 15.9 even at one twentieth of  $\beta^*$ , while the cross-  
 970 visit covariance error grows from 0.59 to 0.69. Marginal and mechanistic fidelity are comparatively  
 971 insensitive to  $\beta$ , so the cost of lowering it is borne by the longitudinal covariance structure. No  
 972 inverse temperature yields a cohort that is at once as private as the real patients are from one  
 973 another and faithful to their joint structure.

974 Neither lever settles the question completely, because every scheme compared so far is a finite-  
 975 time Markov chain. The target itself, however, is available in closed form. Completing the square  
 976 in Eq. (1) gives Eq. (3): for unit-normalized memories  $\pi_r$  is exactly the isotropic Gaussian mixture  
 977  $\sum_k q_k \mathcal{N}(\mathbf{m}_k, \beta^{-1}\mathbf{I})$  with  $q_k = r_k / \sum_j r_j$ . It can therefore be sampled ancestrally, by drawing a com-  
 978 ponent and adding isotropic noise, with no chain, no initialization, no burn-in and no discretization  
 979 error. We generated a protocol-matched exact cohort ( $N=100$  at  $\beta^*$ , reusing the ladder’s seeded  
 980 magnitude draw so that only the direction sampler differs) and evaluated it on the identical harness  
 981 (Table S16). Every metric falls inside the range spanned by the four rungs: median relative error  
 982 1.68%, cross-visit Frobenius error 0.60, mechanistic overlap 0.91, median novelty 0.50, and median  
 983 DCR 13.87. One hundred independently seeded exact cohorts place these quantities in narrow in-  
 984 tervals, with median relative error 1.40% (fifth to ninety-fifth percentile 0.97 to 1.91%) and median  
 985 DCR 13.95 (13.61 to 14.39), so the agreement does not rest on a favorable seed. Evaluated over  
 986 the same 40 train/holdout splits used for the ladder endpoints, the membership-inference AUC  
 987 under exact sampling is 0.98, marginally above the published scheme’s 0.97. The memorization is  
 988 therefore a property of  $\pi_r$  at  $K=23$ , where every stored patient is a component center, rather than  
 989 of any approximation used to sample it.

Table S16: **Sampler-correctness comparison.** Each rung holds  $\beta^*$ , the memory, and the decoding fixed and varies only the direction sampler. Fidelity metrics (median relative error, cross-visit correlation Frobenius error, mechanistic cloud overlap, median novelty) and privacy metrics (median distance to the closest real record, and repeated-split membership-inference AUC where evaluated) are reported per rung. The real-to-real DCR baseline is 15.9. The membership-inference attack was run on the two endpoints of the ladder and on the analytic reference, over the same 40 train/holdout splits. *Exact* is not a rung: it draws ancestrally from Eq. (3) with no chain, and is reported as the analytic reference for the four sampled rungs. Over 100 independently seeded exact cohorts the median relative error is 1.40% (fifth to ninety-fifth percentile 0.97 to 1.91%) and the median DCR is 13.95 (13.61 to 14.39).

| Rung  | Sampler            | MRE (%) | Cross-visit Frob. | Mech. overlap | Novelty | DCR   | MIA AUC |
|-------|--------------------|---------|-------------------|---------------|---------|-------|---------|
| V0    | anchor, endpoint   | 1.31    | 0.65              | 0.89          | 0.51    | 13.97 | 0.97    |
| V1    | sphere init        | 1.53    | 0.57              | 0.89          | 0.50    | 13.96 | n/a     |
| V2    | pooled ULA         | 2.02    | 0.59              | 0.88          | 0.50    | 14.15 | n/a     |
| V3    | pooled MALA        | 1.89    | 0.62              | 0.92          | 0.52    | 14.33 | 0.96    |
| Exact | analytic ancestral | 1.68    | 0.60              | 0.91          | 0.50    | 13.87 | 0.98    |




Figure S9: **Sampling correctness and the memorization/privacy relationship.** (A) Across the V0 to V3 ladder the distance to the closest real record (DCR, left axis) rises only slightly and stays below the real-to-real baseline, while the membership-inference AUC (right axis) remains near 0.96 to 0.97. (B) The fidelity cost of multi-chain pooling, as median relative error. (C) Sweeping the inverse temperature across a 40-fold range around the operating point  $\beta^*$ : the DCR (privacy, left axis) rises as  $\beta$  falls but never reaches the real-to-real baseline, while the cross-visit covariance error (right axis) grows, and the mechanistic overlap stays flat.




Figure S10: **Mixing diagnostic at  $\beta^*$ .** Energy traces, the between-chain  $\hat{R}$  statistic, and the anchor-distance trajectory for sphere-initialized Langevin chains, showing that the target distribution at  $\beta^*$  is multimodal, with one component centered on each stored pattern. Metropolis acceptance near 0.999 indicates negligible discretization bias.