

Validated Synthetic Patient Generation for Small Longitudinal Cohorts: Coagulation Dynamics Across Pregnancy

Jeffrey D. Varner^{1,*}, Maria Cristina Bravo², Carole McBride³,
Thomas Orfeo², Ira Bernstein³

¹Robert F. Smith School of Chemical and Biomolecular Engineering,
Cornell University, Ithaca, NY 14850

²Department of Biochemistry, The Robert Larner, M.D.
College of Medicine, University of Vermont Medical Center, Burlington, VT 05401

³Department of Obstetrics, Gynecology, and Reproductive Sciences,
The Robert Larner, M.D. College of Medicine,
University of Vermont Medical Center, Burlington, VT 05401

*Corresponding author: jdv27@cornell.edu

Abstract

Small longitudinal cohorts, common in maternal health, rare diseases, and early-phase trials, limit computational modeling because enrollment is slow and the data are too sparse to train reliable models. We present multiplicity-weighted Stochastic Attention (SA), a generative framework based on modern Hopfield networks. Stochastic attention stores real patient profiles as memory patterns in a continuous energy landscape. [The resulting distribution is a finite mixture with one component centered on each stored profile. We generated the reported cohort with Langevin dynamics; direct sampling from the same distribution reproduced its fidelity and membership-inference results.](#) Multiplicity weights amplify selected subgroups at inference time without retraining. We applied the method to longitudinal coagulation data from 23 patients, with 72 features measured before pregnancy and during the first and third trimesters. The cohort included three patients with polycystic ovary syndrome and five who developed preeclampsia. Across statistical, structural, and mechanistic tests, including [a separately specified coagulation model that was blind to the generator but calibrated on the same cohort](#), the synthetic profiles closely matched the real profiles under the measures applied at this sample size. A model calibrated on synthetic profiles predicted held-out real-patient TGA measurements as well as one calibrated on real profiles. [These results provide proof of concept for using SA in the tested modeling tasks, including mechanistic calibration and hypothesis generation, with very small longitudinal cohorts.](#)

Keywords: synthetic data generation, stochastic attention, modern Hopfield networks, multiplicity weighting, coagulation, pregnancy, longitudinal data, mechanistic validation

1 Introduction

In 2021, the United States maternal mortality rate reached 32.9 deaths per 100,000 live births, the highest level since 1965^{1,2}. Between 2018 and 2022, Non-Hispanic Black women died at 2.8 times the rate of White women, and disorders related to pregnancy, including hemorrhage and

venous complications, were the leading underlying cause of death, accounting for 17.4% of 6,283 pregnancy-related deaths³. These figures highlight the role of the coagulation system in maternal outcomes. Blood coagulation is a complex enzymatic cascade in which a small tissue factor (TF) stimulus triggers a series of serine protease activation reactions that ultimately produce thrombin, the central effector of hemostasis^{4,5}. Thrombin cleaves fibrinogen into fibrin, activates platelets, and amplifies its own production through positive feedback via Factors V, VIII, and XI, while simultaneously initiating its own shutdown through the thrombomodulin–protein C anticoagulant pathway^{6,7}. Pregnancy alters this hemostatic balance, shifting toward a hypercoagulable state that prepares for delivery^{8–10}. Procoagulant factors including fibrinogen, Factor VIII, and von Willebrand factor increase substantially across gestation, while natural anticoagulants such as antithrombin and protein S decline. Two conditions further shift this balance in ways that are not fully understood. Polycystic ovary syndrome (PCOS), a hormonal and metabolic disorder affecting approximately 6.6% of reproductive-age women¹¹, has been associated with impaired fibrinolysis and a prothrombotic tendency¹². Preeclampsia (PE), a hypertensive disorder that complicates 2–8% of pregnancies worldwide¹³, is characterized by endothelial dysfunction, platelet activation, and altered thrombin generation¹⁴. Mathematical models of the coagulation cascade have provided mechanistic insight into patient-specific thrombin generation^{15,16}. These range from the original Hockin–Mann stoichiometric framework⁶ through extensions incorporating the thrombomodulin–protein C pathway⁷ to the Brummel–Ziedins 2012 (BZ2012) model with detailed activated protein C (APC)-mediated Factor Va inactivation kinetics¹⁷. However, applying these models to pregnancy cohorts requires longitudinal datasets capturing coagulation factors, inhibitors, and thrombin generation assay (TGA) measurements across gestation, data that is rare and expensive to collect.

The core barrier is sample size. Longitudinal pregnancy coagulation studies typically enroll tens of patients with complete follow-up across multiple trimesters, yielding datasets where the number of features p (coagulation factors, TGA parameters, viscoelastic measurements) exceeds the number of patients n . Rare complications compound this problem. Small longitudinal cohorts inevitably contain only a handful of patients with any given condition, far too few for condition-specific statistical analysis. This $n < p$ regime limits conventional approaches because covariance matrices are rank-deficient, cross-validation is unreliable, and overfitting is nearly inevitable¹⁸. Synthetic data can provide dense simulated inputs for mechanistic modeling, workflow testing, and hypothesis generation^{19,20}. **A further motivation was to support the development of deep learning models.** However, existing methods face limitations in the small-cohort longitudinal setting. Sampling from a fitted multivariate normal (MVN) distribution is singular when $n < p$ and must be regularized²¹, introducing bias that distorts the joint distribution. The multivariate-normal model also assumes Gaussian marginals and linear correlations and cannot amplify a specific subpopulation while preserving its signatures. More expressive models such as generative adversarial networks (GANs) and variational autoencoders (VAEs)^{22–24}, including recent longitudinal extensions²⁰, require substantially larger training sets and are prone to mode collapse at small n ²⁵ (we confirm this empirically for Conditional Tabular GAN (CTGAN) at $n=23$ in this work). A generative method is needed that can operate directly on the geometry of a small dataset without **estimating a full covariance model**.

Stochastic Attention (SA) is such a framework. **Rather than fitting a full covariance model or training a neural generator**, SA treats the stored patient profiles themselves as the model. Building on modern Hopfield network theory²⁶, each profile becomes a memory pattern in a continuous energy landscape. **This energy induces a continuous probability distribution that is exactly a finite Gaussian mixture in the reduced space, with one nonzero-variance component centered on each stored pattern. Sampling from this distribution therefore produces continuous variations around the stored patterns rather than exact copies.** Because sampling occurs in a principal component

analysis (PCA)-reduced linear subspace anchored on the observed cohort, SA never forms a full $p \times p$ covariance matrix and so avoids both the rank deficiency that limits MVN and the mode collapse that limits GAN/VAE approaches when $n < p$. We recently introduced a multiplicity-weighted extension of the Hopfield energy that attaches a per-pattern weight to each stored profile²⁷, which allows continuous interpolation between unconditioned generation and targeted amplification of a rare subpopulation at inference time, without retraining. This study asked whether SA could faithfully generate *longitudinal* trajectories from a small cohort. In earlier proceedings work, we paired our mechanistic coagulation pipeline with per-visit MVN sampling to produce synthetic populations²⁸; that approach captured single-visit marginals but, by fitting each visit independently, generated patients whose trajectories across gestation were statistically incoherent. Small longitudinal cohorts require a generator that preserves this joint structure.

In this study, we used multiplicity-weighted SA to generate complete longitudinal patient profiles from a small pregnancy coagulation cohort ($K=23$ patients, 72 features per visit, 3 visits) that includes rare clinical subgroups (PCOS, $n=3$; Developed PE, $n=5$). Our pipeline concatenates each patient’s multi-visit profile into a single vector, applies PCA for dimensionality reduction, and generates new patients via multiplicity-weighted Langevin dynamics on the Hopfield energy landscape, with a direction-magnitude decomposition that preserves the natural variance structure of continuous clinical measurements. We then asked whether the synthetic patients reproduced marginal feature distributions, the cross-visit covariance structure, condition-specific signatures of rare clinical subgroups, and the biological consistency obtained when the profiles were processed by a separately specified, generator-blind ordinary differential equation (ODE) model of the coagulation cascade that was calibrated on the same real cohort¹⁷. As a downstream-utility test, we asked whether a mechanistic model calibrated entirely on synthetic patients could predict held-out real-patient TGA measurements. As a proof of concept, the results show that SA recovered the statistical structure of this cohort and matched its mechanistic plausibility under the tests used here. Generalization to other clinical cohorts remains to be tested.

2 Results

We generated $N=100$ synthetic longitudinal patient profiles using SA and compared them against the $K=23$ real patients across four validation levels. As a baseline, we generated $N=100$ patients from a regularized multivariate normal (MVN) distribution fitted to the same data.

2.1 Marginal Plausibility

We first assessed whether SA-generated patients reproduced per-feature summary statistics across all three visits. The pooled median mean relative error (MRE) across all 216 feature–visit entries was 1.2% (95% bootstrap CI: 1.0–1.6%), with 193 of 216 entries (89%) below 5% (Table S5), indicating that SA captured the central tendencies of the real cohort. Per-visit MREs for representative coagulation features are summarized in Table 2, where Factor II and Factor VIII fell below 2%, antithrombin below 3%, and fibrinogen below 1% across all three visits. Marginal agreement alone could result from profiles that closely copied stored records, so we also quantified similarity to stored profiles with two metrics (Table S9). The mean novelty score ($1 - \max_k \cos(\hat{\xi}, \hat{\mathbf{m}}_k)$) was 0.50 at $\beta = \beta^*$. The median nearest-neighbor distance from each synthetic patient to its closest real patient was 6.14 standardized units, compared with 6.87 for the median real-to-real nearest-neighbor distance, a ratio of 0.89 indicating that synthetic patients were approximately as far from real patients as real patients were from each other. Thus, stochastic attention matched the central tendencies without reproducing stored profiles exactly.

We next examined whether known physiological relationships and longitudinal trajectories were preserved in the synthetic cohort. Pairwise scatter plots of coagulation factor levels against thrombin generation parameters showed that SA-generated patients occupied the same joint regions as real patients (Fig. 1). The expected inverse relationship between antithrombin and TF-initiated peak thrombin was preserved, as were the positive associations between Factor VIII and peak, prothrombin and endogenous thrombin potential (ETP), and the inverse relationship between antithrombin and ETP. Per-visit means and standard deviations for six key coagulation features showed that synthetic patients tracked the real longitudinal trajectories closely, with overlapping confidence bands at all three visits (Fig. 2). Fibrinogen, Factor VIII, and von Willebrand factor showed the expected monotonic increases from baseline through the third trimester in both cohorts, while Factor II increased modestly and antithrombin declined then partially recovered, consistent with the known hypercoagulable shift of pregnancy. The 72 features also included viscoelastic (ROTEM) and fibrinolytic parameters, which had comparable MREs to the coagulation factors (Table S5). [These variables influenced the generated profiles through their loadings on the retained principal components.](#) Together, these results showed that the synthetic cohort preserved the tested physiological relationships and the direction and magnitude of longitudinal change.

To determine whether standard deep tabular models provided a viable baseline at this sample size, we compared SA against Conditional Tabular GAN (CTGAN) and Tabular Variational Autoencoder (TVAE)²⁴ (Table S6). The CTGAN model produced median MREs of approximately 19% across all tested epoch counts (300–3,000), consistent with mode collapse at this sample size. The TVAE model achieved comparable marginal fidelity at high epoch counts (median MRE 1.8% at 3,000 epochs), but operated on per-visit rows ($n=90$) rather than concatenated longitudinal profiles ($n=23$), meaning it could not capture cross-visit covariance structure and had no mechanism for conditional subgroup generation. Thus, neither deep model provided a suitable complete-profile baseline for this 23-patient cohort.

[To examine the role of the shared PCA representation, we compared stochastic attention with a fitted two-component diagonal-covariance Gaussian mixture, a kernel-density estimator, \$k\$ -nearest-neighbor interpolation, and a Gaussian copula in the same \$d=18\$ PCA subspace \(Supplementary Methods, PCA-space ablation; Supplementary Table S13\).](#) Nearest-neighbor samples remained close to stored profiles: median novelty, $1 - \max_k \cos(\hat{\mathbf{x}}, \hat{\mathbf{m}}_k)$, was 0.05, and only 5% of the synthetic patients exceeded the 0.2 novelty cutoff. The kernel-density estimator had marginal error 1.9% and mechanistic overlap 0.82 (the fraction of pooled synthetic predicted-to-recorded TGA ratios within the real 5th–95th percentile range). The two-component mixture and copula matched stochastic attention more closely in marginal error (1.1% and 1.3% versus 1.2%) and mechanistic overlap (0.92 and 0.93 versus 0.90). All profiles from stochastic attention and the copula exceeded the novelty cutoff, compared with 94% from the two-component mixture. Stochastic attention had a larger empirical cross-visit correlation error than the two-component mixture and copula (0.64 versus 0.42 for both). This metric measured agreement with the correlation matrix of the same 23 patients used to construct the stochastic-attention memory matrix or fit the baseline generators; it therefore assessed in-sample reproduction rather than recovery of the unknown population correlation matrix. Holding the PCA representation constant showed that the marginal errors and mechanistic overlaps achieved by the two-component mixture and copula were similar to those of stochastic attention and therefore were not specific to the stochastic-attention generator. In contrast, empirical correlation error and novelty differed across generators. Thus, generator choice affected these aspects of performance, and no generator performed best on every measure.

2.2 Cross-Visit Covariance Structure

Having established that SA reproduces marginal distributions and pairwise relationships, we next asked whether it preserves the higher-order joint structure across visits. We compared the full 216×216 cross-visit correlation matrices for real, SA-generated, and MVN-generated populations. The real data exhibited a characteristic block structure with strong within-visit correlations along the diagonal and structured off-diagonal blocks reflecting cross-visit dependencies; for example, a patient’s Visit 1 Factor X level predicting their Visit 3 Factor X level. We found that SA preserved this block structure, including the off-diagonal cross-visit correlations (Fig. 3). The SA–Real residual matrix showed small, unstructured deviations distributed throughout, whereas the MVN–Real residual was notably smoother, reflecting the regularization-induced shrinkage of off-diagonal correlations toward zero. This difference was most apparent in the off-diagonal blocks, where MVN systematically underestimated cross-visit dependencies that SA retained. We note that Pearson correlation captures linear associations; Spearman rank correlations, used in the mechanistic validation, would additionally capture monotonic nonlinear dependencies. Within the linear analysis used here, stochastic attention retained the observed cross-visit structure more closely than the regularized MVN baseline.

We examined the eigenvalue spectrum of the sample covariance matrix to understand the structural origin of this difference (Fig. S2). The real data had rank 22, reflecting the $K=23$ patient constraint. Stochastic attention and MVN make different trade-offs in handling this rank deficiency: SA truncates via PCA (zeroing variance beyond component 18), while MVN regularizes via Ledoit–Wolf shrinkage (providing a full-rank estimate by inflating eigenvalues in the 194 null dimensions). The consequence of MVN’s approach is that it introduces variance in dimensions where the data contain no signal, which manifested as increased dispersion in PCA projections. Visit-specific PCA projections showed this dispersion gap (Supplementary Fig. S8). Patients generated by stochastic attention occupied the same region of PCA space as real patients across all three visits, while MVN-generated patients showed substantially greater scatter, particularly at Visits 2 and 3. The spectrum and projections therefore linked the MVN baseline’s greater scatter to variance added in directions unsupported by the cohort.

2.3 Robustness to Sample Size, Dimensionality, and Nonlinearity

Results from a single cohort of 23 patients could not show how stochastic attention would perform if the amount or structure of the available data changed. We therefore asked three questions in controlled tests (Methods, Sec. 5.5). First, could generated profiles recover the true covariance? We simulated datasets in which the covariance was known. We varied the size of the simulated datasets from 8 to 80 profiles, the number of variables per profile from 30 to 216, and the complexity of the correlation patterns. From each dataset, we constructed a stochastic-attention memory matrix. For comparison, we estimated a Gaussian generator from the dataset’s sample mean and covariance. Each generator then produced 100 profiles. We compared the covariance of each generated cohort with the known population covariance. Stochastic attention had lower error in 35 of the 39 settings in which the simulated dataset contained fewer profiles than variables (Fig. 4A). For simulations matching the clinical dataset in this study in terms of size, with 23 profiles and 216 variables per profile, the Gaussian generator’s covariance error was approximately 1.5 times the stochastic-attention error after averaging across the three levels of correlation-pattern complexity. Second, we examined how well stochastic attention preserved characteristics of the full 23-patient cohort when fewer patients were available. Using random subsets of 20, 15, 10, or 8 patients, we repeated the random selection ten times at each size, constructed a new memory matrix from

each, generated 100 profiles, and compared every generated cohort with the same full 23-patient reference. As the subset size decreased from 20 to 8, error in the cross-visit correlations increased from 0.63 to 1.11 (Fig. 4B), the median error in variable means increased from 2.0% to 4.9%, and overlap with the mechanistic-model reference remained near 0.90. Third, could stochastic attention and the Gaussian generator reproduce curved relationships? We used a noisy S-curve embedded among either 30 or 120 variables. For each amount of curvature and each number of variables, we created five simulated datasets, each with 23 profiles. From every dataset, we constructed a stochastic-attention memory matrix and estimated a Gaussian generator; each generator produced 100 profiles. As curvature increased, stochastic-attention profiles moved farther from the curve. At curvature 0.5, their mean distance was at least twice the straight-line value after averaging across the five repeats and both numbers of variables. It remained below this threshold at curvature 0.25. Despite this decline, at every nonzero curvature, stochastic-attention profiles remained closer to the curve on average than Gaussian profiles. However, the Gaussian generator reproduced the population covariance more accurately than stochastic attention for the curved data. Together, these tests showed that stochastic attention recovered covariance more accurately than the Gaussian comparator in most of the small linear simulations with more variables than profiles, but its clinical correlations and variable means became less accurate when fewer patients were used. On curved data, each generator performed better on a different measure.

2.4 Interpolation and Extrapolation

The robustness tests examined recovery under controlled changes in the data. We next asked whether profiles generated from the clinical cohort remained within the real cohort or extended beyond it. Every synthetic profile lay outside the convex hull of the 23 real profiles in the $d=18$ PCA space (Supplementary Table S14). However, each real profile also lay outside the hull of the other 22, so binary membership was uninformative in this sparse space. We therefore compared hull distances against 22 real profiles. Each real calculation omitted the profile being evaluated, while each synthetic calculation omitted its nearest real profile. If generation extrapolated unusually far, synthetic distances would exceed leave-one-out real distances. The medians were 11.0 and 11.9, respectively, and a descriptive Mann–Whitney test did not detect a difference ($p=0.07$). While this did not establish equivalence, it provided no evidence that synthetic profiles extended farther from the cohort than real profiles treated as unseen. Geometric similarity did not establish biological plausibility, so we applied two further checks. The first required nonnegative values where appropriate and limited coagulation-factor and TGA values to five real-cohort standard deviations from the mean. Fifty-four of 100 synthetic profiles passed; the other 46 contained a negative estradiol or progesterone value. The unconstrained multivariate-normal baseline produced negative hormones in 43 profiles. Neither method enforced positivity, and positive factors in these draws did not guarantee positivity in others. The second check used the BZ2012 model to test the generated factor combinations. Estradiol and progesterone were not model inputs, and all nine coagulation factor inputs were positive, so all 100 profiles were included. For each TGA feature, 83–92% of profiles had a predicted-to-recorded ratio within the real 5th–95th percentile range. Forty-one met this criterion for every feature and visit, and 24 passed both checks. Positive-variable transformations or constrained decoding could prevent negative values. Together, these checks separated geometric extrapolation from biological plausibility. Hull distance provided no evidence that synthetic profiles extended farther than real profiles treated as unseen, but only 24 of 100 profiles passed both the biological and strict mechanistic criteria. Thus, hull position alone could not determine whether an individual generated profile was biologically plausible.

2.5 Membership-Inference Risk

Geometric proximity raised a separate privacy question: could a synthetic cohort reveal whether an already-known de-identified assay profile had been stored in the memory matrix? To answer this question, we tested whether distance from a candidate profile to the synthetic cohort could distinguish profiles used to construct the memory matrix from profiles withheld from it. In each of 40 random splits of the 23-patient cohort, we constructed the memory matrix from 15 profiles, held out the other 8, and generated a synthetic cohort. A smaller distance to the nearest synthetic profile provided stronger evidence that a candidate had been stored. Stored profiles were closer to the synthetic cohort than held-out profiles (median 10.7 versus 15.4; Supplementary Table S15). We quantified this discrimination using the area under the receiver operating characteristic curve (AUC). The AUC was the probability that the test ranked a randomly chosen stored profile as more likely to have been stored than a randomly chosen held-out profile; 0.5 indicated chance discrimination and 1.0 indicated perfect discrimination. The mean AUC was 0.97, and nine splits reached 1.0. We next asked whether this signal came from the numerical sampling procedure. A Metropolis-adjusted procedure left the mean AUC nearly unchanged at 0.96. Lowering the inverse temperature over a 40-fold range increased the median distance from synthetic profiles to their nearest stored profiles, but cross-visit covariance fidelity declined. Drawing directly from the target distribution, without a Markov chain, produced a mean AUC of 0.98 (Supplementary Table S16, Fig. S9). Thus, the membership signal came from the target distribution in this small-cohort setting, not from the Markov-chain implementation. The test could reveal whether an already-known de-identified profile had been stored, but it could not identify a person or reconstruct an unknown record.

2.6 Conditional Generation of Rare Subgroups

The previous analyses used unconditioned SA, generating from the full cohort. A key clinical application, however, is generating patients from specific subpopulations that are too small to study independently. We therefore tested whether SA could amplify rare clinical subpopulations while preserving their condition-specific signatures. We defined three overlapping subgroups by pregnancy outcome and comorbidity: Uncomplicated ($n=18$), PCOS ($n=3$, cross-cutting; 1 patient also in the PE group), and Developed PE ($n=5$). The PCOS and Developed PE groups were too small for any independent statistical analysis. We used SA’s multiplicity-weighted sampling to generate condition-specific cohorts of 100 patients each by upweighting attention on the relevant stored patterns.

We compared condition-specific feature means between real and SA-generated patients across eight coagulation features that exhibited between-condition variation (Fig. 5). The generated cohorts preserved the condition-specific patterns. Patients with PCOS showed elevated Factor VIII and vWF relative to Healthy patients in both real and synthetic data, PE patients showed elevated α_2 -AP and Factor IX, and the between-condition rank ordering was maintained across all eight features. We then performed a [descriptive bootstrap Mann–Whitney diagnostic](#). For each feature and condition, we subsampled the synthetic cohort to match the real sample size (n_{real}), applied a Mann–Whitney U test, and repeated 1,000 times. We reported the fraction of replicates where $p > 0.05$, meaning that this test did not detect a difference in that replicate. This power-limited fraction was descriptive and did not establish equivalence. Across all 24 feature–condition pairs, 20 (83%) achieved $p > 0.05$ in $\geq 90\%$ of replicates, with a median no-difference-detected fraction of 98.6% (full per-pair results in Table S10). The four pairs that fell below 90% were high-variance features in the smallest subgroups (FIX and plasminogen in PCOS; FIX and plasminogen in Developed PE), where

the real data exhibited large inter-patient variability. No multiple-testing correction was applied because this diagnostic was not used to make simultaneous inferential claims; the per-pair values and the 90% threshold were descriptive reporting summaries. The PCA projections of the conditioned cohorts are provided in Figs. S3–S4. These results showed that SA’s multiplicity-weighted interpolation can amplify rare subgroups in a way that preserves their distinguishing clinical features, useful for exploratory hypothesis generation and simulation-based workflow testing, but not for increasing the effective clinical sample size. The multivariate-normal model cannot do this: fitting a separate MVN to three PCOS patients across 216 features is mathematically impossible (rank 2 covariance), and post-hoc subsetting of unconditional MVN samples does not preserve subgroup-specific structure. Together, these comparisons showed that multiplicity weighting preserved the subgroup patterns represented in the source cohort.

Generating 100 conditioned profiles did not add independent subgroup information. The PCOS and Developed PE signals still came from three and five real patients, and their participation ratios were 4.6 and 7.7 (Table S2). Thus, repeated draws were concentrated around a few stored profiles. Treating them as independent observations would understate uncertainty. They instead provided denser simulations of the distributions suggested by those patients for exploratory modeling and workflow testing.

2.7 Mechanistic Consistency

The preceding validations assessed statistical properties of the synthetic data. We next tested whether factor combinations generated by stochastic attention produced biologically plausible outputs when passed through a separately specified mechanistic model. The model was blind to the generation process but was calibrated on the same cohort. We used the Hockin–Mann BZ2012 model^{6,17}, a system of 58 ordinary differential equations with 64 rate constants that simulates thrombin generation from patient-specific coagulation factor inputs. We calibrated 5 of the 64 rate constants (prothrombinase, intrinsic and extrinsic tenase, protein C activation, and meizothrombin conversion; Table S4) on Visit 1 real patients at the population level and held the remaining 59 at published literature values. We then ran the calibrated model on all 23 real patients and 100 synthetic patients under both TF-only and TF+TM conditions (738 total simulations, 0 failures). For each patient, we computed five ODE-predicted thrombin generation features (lagtime, peak, time-to-peak, maximum rate, and ETP) and compared them against the corresponding values from the patient record. In these comparisons (Fig. 6, top row), the horizontal axis is the TGA value from the patient record, which exists for both real and synthetic patients, while the vertical axis is the value computed by the BZ2012 model from that patient’s factor inputs. Systematic deviations from the $y=x$ line reflect ODE model bias (e.g., lagtime is underpredicted at $\approx 0.77\times$), not a failure of the synthetic data; the key observation is that real and synthetic patients share the same pattern of model bias, occupying the same cloud with the same systematic offset.

We formalized this observation by computing the ODE-predicted/dataset ratio for each patient and comparing the resulting distributions between real and synthetic populations (Fig. 6, bottom row). These ratio distributions capture how the ODE model processes each patient’s factor levels. If SA had generated patients with biologically implausible factor combinations, the model would process them differently, producing a shifted or broadened ratio distribution. Instead, the distributions overlapped substantially, with cloud overlap (fraction of synthetic ratios within the 5th–95th percentile of real ratios) ranging from 0.83 for ETP to 0.92 for T.Peak under TF-only conditions, and two-sample Kolmogorov–Smirnov (KS) tests did not detect a difference between the real and synthetic ratio distributions at this sample size across all five TGA features ($D = 0.081$ – 0.123 , all $p > 0.30$; full diagnostics in Table S11). Under TF+TM conditions, the model systematically

overcorrected the protein C anticoagulant pathway, producing peak and maximum rate predictions approximately $0.5\times$ measured values for both real and synthetic patients (Fig. S5, same layout as the main-text comparison). Cloud overlap remained high (0.88–0.92 across the five features), showing similar model behavior for real and synthetic profiles despite the imperfect calibration.

The model was calibrated only on Visit 1 real patients, but the per-visit median predicted-to-measured ratios under TF-only conditions remained close to 1.0 at Visits 2 and 3 (Fig. S6). This was consistent with rate constants that remained useful beyond the calibration visit. Spearman rank correlations between predicted and measured values were moderate for ETP ($\rho \approx 0.6$ –0.8 for real, $\rho \approx 0.6$ –0.65 for synthetic; Fig. S7) and weak for other features, consistent with a 5-parameter population-level calibration that was not designed to resolve inter-patient variability. The key finding was not patient-level prediction but population-level plausibility. **Applied identically to both groups, the mechanistic model produced overlapping predicted-to-recorded ratio distributions for real and synthetic patients under both experimental conditions. Because the model was calibrated on this same cohort, this was a same-cohort consistency check rather than independent validation.**

2.8 Downstream Utility: Mechanistic Model Calibration

Having established same-cohort mechanistic consistency, we tested whether synthetic profiles could support mechanistic-model calibration. We fit the five BZ2012 rate constants under TF-only conditions to either 100 synthetic V1 profiles or 23 real V1 profiles, using the same bounded Nelder–Mead optimization and 11 random restarts (Table S4). We evaluated both calibrations on held-out real V2 and V3 TGA measurements. Median relative errors for the synthetic calibration were 2–10% lower across the five features, with an overall synthetic-to-real error ratio of 0.94 (Fig. 7; Table S12). The parameter estimates differed, but the two calibrations produced highly correlated predictions. Because the synthetic profiles came from the real cohort, this was not an independent validation. It tested whether the representation retained enough structure for calibration. **The modest reduction in prediction error did not show that synthetic data were superior. It may instead have reflected a more stable calibration objective produced by sampling the represented distribution more densely. Together with the subgroup and factor-level analyses, this result supported exploratory modeling and hypothesis generation. Confirming those hypotheses and establishing clinical utility would require additional real patients or independent experimental data.**

3 Discussion

This study asked whether a geometry-based method could generate synthetic profiles from a very small longitudinal cohort while preserving statistical and modeled biological relationships. Median feature MRE was 1.2%, which the MVN baseline also achieved. Stochastic attention also retained cross-visit covariance, whereas MVN introduced variance in 194 null dimensions. Conditioned generation preserved the subgroup patterns represented by three PCOS and five Developed PE patients. Finally, **a separately specified, generator-blind coagulation model calibrated on the same cohort placed real and synthetic profiles in the same factor-to-thrombin-generation envelope.** Cloud overlap was 83–92%, and the Kolmogorov–Smirnov tests did not detect differences under these conditions ($p > 0.35$ for all five TGA features).

Dense simulations from very small longitudinal datasets could support maternal health research. Rare pregnancy complications such as PCOS and preeclampsia are difficult to study because assembling large, well-characterized longitudinal cohorts requires years of recruitment across multiple clinical sites. Our results support hypothesis generation, computational modeling, and workflow testing with dense simulations of these small cohorts. **Stochastic attention can provide denser**

simulated realizations of the distribution implied by an existing small subgroup. These generated profiles do not add clinical evidence beyond the source cohort: the inferential sample size remains three patients for PCOS and five for Developed PE. Synthetic profiles could also support the development of deep learning models that predict TGA responses from coagulation factors or estimate factor profiles compatible with an observed TGA response. Related modeling work has predicted thrombin dynamics from factor levels and searched for factor changes that move thrombin generation toward a target²⁹. We have not yet evaluated either task.

Understanding why the method worked in this setting required examining the role of the PCA representation as well as the generator. When all methods used the same 18-dimensional subspace, the fitted two-component mixture and copula achieved marginal errors and mechanistic overlaps similar to those of stochastic attention. These results were therefore not unique to stochastic attention, although the comparison did not quantify the contribution of PCA itself. Empirical correlation error and novelty differed across generators. Stochastic attention used the patient profiles in the PCA subspace as its memory matrix and avoided estimating a full 216-dimensional covariance from 23 patients. Its sampler still required a choice of inverse temperature. A natural concern is whether the β^* selection procedure, originally developed for protein sequence generation³⁰, transfers to clinical coagulation data. The entropy transition rule selected the point of greatest downward curvature in attention entropy, a geometric property of K patterns in d dimensions rather than a property of what the features represented. For this dataset it gave $\beta^* \approx 2.94$, with the random-pattern scale $\sqrt{d} \approx 4.24$ serving as a theoretical reference. A sensitivity analysis showed gradual changes in generation quality across $\beta/\beta^* \in [0.1, 3.0]$ (Table S1). Similarly, varying the PCA variance retention threshold from 85% to 99% ($d_{\text{PCA}} = 13$ to 21) had minimal impact on generation quality. Median MRE remained between 1.2% and 1.8% across all thresholds (Table S7), so the 95% threshold was not a critical design choice. Together, these sensitivity analyses showed that performance was not tied to a narrow choice of inverse temperature or PCA threshold within the ranges tested.

Conditional generation raised a different question: how much diversity remained after the selected subgroup was upweighted? Achieving 80% of finite-mixture component probability on the PCOS subgroup required $\rho \approx 26.7$, reducing the participation ratio to $K_{\text{eff}} \approx 4.6$ (multiplicity parameters for all three subgroups in Table S2), which indicated that mixture-component probability was concentrated around a small number of stored profiles across repeated draws. The subgroup-specific signal still rested on only three real PCOS profiles, and the lower-weight background profiles pulled generated means toward the population average; K_{eff} did not represent an independent patient count. Magnitudes were drawn from all 23 empirical PCA-space norms, so conditioning acted through the weighted directions rather than a subgroup-specific magnitude distribution. Thus, weighting concentrated sampling toward the subgroup directions while retaining some influence from the full cohort. Related studies have applied the same geometry-based operating-point rule and multiplicity weighting to protein sequence generation and conditional steering^{27,30,31}.

Among longitudinal clinical generators, VAMBN-MT provides a useful comparison in a larger-data setting²⁰. It combines expert-defined modules, Bayesian-network constraints, and LSTM-augmented variational autoencoders. The authors evaluated it on the DONALD nutritional cohort and the Alzheimer’s Disease Neuroimaging Initiative. DONALD included 1,312 participants and 33 longitudinal variables, compared with 23 patients and 216 features here. Its best variant had cross-visit correlation error near 0.70, and weaker time trends were not reproduced at the tested sample sizes. The methods address different regimes. The VAMBN-MT model uses a trained modular architecture for moderate cohorts. Stochastic attention uses PCA geometry when $n < p$, applies a generic generator-blind mechanistic check instead of domain-specific rules, and uses multiplicity weights for conditional generation. This was a regime comparison, not a head-to-head test.

The mechanistic analysis addressed a different validation question: whether a model of the un-

derlying biology processed real and synthetic profiles similarly. We processed both groups with the same separately specified model, which was blind to the generator but calibrated on the real cohort. A model calibrated on synthetic profiles also predicted held-out real TGA measurements as well as one calibrated on real profiles (error ratio 0.94, with 11 random restarts). The same comparison can be used wherever a validated mechanistic model exists, including pharmacokinetic³², physiological³³, tumor-growth³⁴, and metabolic^{35,36} models. The relevant test is whether the model places real and synthetic profiles in the same predicted envelope.

The average subgroup results did not show whether the conditioned cohorts preserved clinically important distribution tails. In the descriptive bootstrap Mann–Whitney diagnostic, four of 24 feature–condition pairs fell below the 90% no-difference-detected criterion: FIX and plasminogen in the PCOS and Developed PE groups. The relative differences between the real and synthetic subgroup means for these pairs were 9–15%, and variability among the real patients was large. This power-limited diagnostic was not an equivalence test and did not determine whether the lower fractions reflected mean shifts, compressed tails, or limited power. Truncating the PCA representation and drawing from only 23 observed PCA-space norms could nevertheless compress the generated tails. A synthetic cohort that under-represents severe states cannot estimate how often those states occur or how severe they become. Although synthetic profiles extended beyond the real cohort’s convex hull, that result did not show that they reproduced the true tail distribution. We therefore do not recommend using these cohorts to estimate extreme risk without additional real data. The primary mechanistic cloud-overlap measure also focused on the real 5th–95th percentile range and therefore did not validate behavior in the tails. More generally, several findings reflected the limits of generating from only 23 source patients. Hull membership had little resolution, and conditioned cohorts still relied on three PCOS or five Developed PE patients. The modest calibration improvement was consistent with denser sampling of the same patient-derived distribution, not with added biological information. These findings supported simulation and model testing, but they did not increase the clinical evidence supplied by the original cohort.

The data representation imposed a further limit. Concatenating visits allowed the method to preserve cross-visit relationships, but it required every patient to have the same assays at the same aligned visits. The current pipeline could not directly accept irregular timing, missing assays or visits, or variable-length records. Imputation, padding, or a sequence-aware extension could support such data, but each would add assumptions that require validation in a small cohort. The PCA representation imposed a related limit because it captured only linear structure. In the S-curve benchmark, generated profiles moved farther from the known curve as curvature increased. At every nonzero curvature, stochastic attention remained closer to the curve on average than the Gaussian generator, whereas the Gaussian generator recovered covariance more accurately. Neither generator was best on both measures in this benchmark. Nonlinear embeddings could better represent curved or branching relationships, but they would add fitting and decoding choices that are difficult to estimate from 23 patients. The linear representation was therefore a practical choice for this cohort, not evidence that clinical trajectories are generally linear. Thus, the present results did not establish performance on irregular, incomplete, or strongly nonlinear records.

Computational cost and behavior at larger cohort sizes are separate questions. The reported Langevin sampler required work proportional to TKd for each profile. The exact mixture sampler removed the T sequential updates: its component weights could be prepared once, after which each profile required a component draw, a d -dimensional Gaussian draw, and decoding. The principal-component model was estimated once, but inverse transformation and decoding were still performed for every generated profile. The simulation advantage over the sample-covariance Gaussian was concentrated in the $n < p$ settings. We did not test cohorts larger than 23 patients, so stability and fidelity as K grows remain open questions.

This study had two broad limitations: the absence of independent clinical data and the restricted scope of the validation analyses. First, the clinical analysis used one longitudinal cohort of 23 patients. The known-covariance simulations tested selected properties under controlled conditions, and related protein studies applied the sampling framework in another domain, but neither established that the clinical findings would generalize to a different patient population. We did not test the method in a second clinical cohort, which is needed to assess that generalizability. Second, the validation analyses answered narrower questions than full biological or clinical validity. The BZ2012 model was calibrated on the same cohort used for generation. Applying one fixed, generator-blind model to real and synthetic profiles therefore tested consistency with that model, not independent biological validity. The fitted rate constants also required a 50-fold reduction in intrinsic tenase k_{cat} and a 15-fold increase in extrinsic tenase k_{cat} relative to published values (Table S4), likely because the published measurements and the plasma-based assay used different experimental conditions. These fitted values should not be read as new biochemical estimates. The biological scope was also incomplete. The mechanistic test covered thrombin generation but not the fibrinolytic or viscoelastic features. The membership-inference analysis answered a separate, narrow question: whether an already-known de-identified profile had been stored in the memory matrix. It distinguished stored from held-out profiles, but it did not identify a participant or reconstruct an unknown record. Finally, we did not test clinical outcomes. Thus, the results did not establish biological fidelity across the full profile or clinical utility. An independently calibrated model, models covering the other feature classes, and prospective outcome validation could address these gaps. The results establish proof of concept, not generalizability.

4 Conclusion

Multiplicity-weighted Stochastic Attention generated longitudinal profiles from 23 patients with low marginal error, retained cross-visit covariance, and preserved the represented subgroup patterns. The profiles were also consistent with a coagulation ODE model calibrated on the same cohort. A model calibrated on synthetic profiles predicted held-out real TGA measurements as well as one calibrated on real profiles. Direct sampling from the same target distribution reproduced the Langevin sampler’s fidelity and membership-inference results. Together, these findings provide proof of concept for using synthetic profiles in the modeling tasks tested here with very small longitudinal cohorts.

5 Methods

5.1 Population Dataset

We analyzed a longitudinal dataset of coagulation, fibrinolysis, thrombin generation assay (TGA), and rotational thromboelastometry (ROTEM) measurements from women desiring a pregnancy. The study was approved by the Institutional Review Board at the University of Vermont. All participants gave written consent before enrollment. Blood collection, plasma processing, and assay protocols have been described previously^{37–40}. Briefly, citrated platelet-poor plasma was collected without tourniquet use after supine rest at three clinical visits: Visit 1 (pre-pregnancy baseline in the follicular phase), Visit 2 (end of first trimester), and Visit 3 (mid-third trimester). Plasma aliquots were stored at -80°C for subsequent analysis. Measurements of secondary analytes were measured as follows: Stago clotting assays measured coagulation-factor activity, ELISA measured fibrinolysis proteins, and the CLIA-certified laboratory at the University of Vermont Medical Center

measured hormone levels. We performed TGA using 5 pM relipidated tissue factor with fluorogenic substrate on a Synergy4 plate reader as described in McLean et al.³⁷. Clot formation and fibrinolysis dynamics were assessed via viscoelastometry using the ROTEM Delta Instrument (Werfen). Thawed plasma was mixed with or without 4 nM tPA (with respect to plasma volume), recalcified with 15 mM calcium chloride, and incubated for 3 minutes at 37°C. The reaction was initiated by adding the plasma mixture to ROTEM cups containing tissue factor at a 17:1 ratio; the final concentration of TF in the reaction was 5 pM.

The dataset contained 153 total records across approximately 50 subjects. After removing columns with > 30% missing values and retaining only subjects with complete data across all three visits, we obtained $K = 23$ complete longitudinal profiles across $n = 72$ assay measurements per visit. The 72 assay measurements spanned five categories: (i) hormones (estradiol, progesterone), (ii) coagulation factors and inhibitors (Factors II, V, VII, VIII, IX, X, XI, XII; antithrombin (AT), protein C (PC), tissue factor pathway inhibitor (TFPI), von Willebrand factor (vWF), fibrinogen, etc.), (iii) fibrinolytic markers (plasminogen, plasminogen activator inhibitor-1 (PAI-1), plasminogen activator inhibitor-2 (PAI-2), α_2 -antiplasmin (α_2 -AP), thrombin-activatable fibrinolysis inhibitor (TAFI), tissue plasminogen activator (tPA), D-dimer equivalents), (iv) thrombin generation parameters under four initiator conditions (TF at 5 pM, TF+thrombomodulin (TM), No TF, No TF+anti-FXIa), and (v) ROTEM/viscoelastic parameters (clotting time (CT), maximum clot firmness (MCF), alpha angle, lysis onset time (LOT), maximum lysis (ML), lysis time (LT), area under the curve (AUC) under TF Only and TF+tPA conditions).

Demographics and clinical characteristics of the resulting cohort, including age at enrollment, prepregnancy BMI, parity, gestational ages at each visit, and the cross-tabulation of enrollment cohort against pregnancy outcome, are summarized in Table 1. The 23 patients were drawn from three enrollment conditions: Healthy nulliparous women ($n=14$), women with a personal history of preeclampsia from a previous pregnancy (Prior PE, $n=6$), and women with polycystic ovarian syndrome (PCOS, $n=3$; 2 individuals were nulliparous). For conditional generation, we defined three overlapping subgroups based on pregnancy outcome and comorbidity rather than enrollment condition: Uncomplicated ($n=18$, all patients whose study pregnancy did not result in preeclampsia), PCOS ($n=3$, a cross-cutting comorbidity; 1 patient also developed PE), and Developed PE ($n=5$, patients who developed preeclampsia during the study pregnancy, drawn from 2 healthy-enrolled, 2 prior-PE-enrolled, and 1 PCOS-enrolled participants).

5.2 Multiplicity-Weighted Stochastic Attention

We generated synthetic patients using a multiplicity-weighted variant of Stochastic Attention (SA) rooted in modern Hopfield network theory²⁶. We stored K memory patterns as columns of a matrix $\hat{\mathbf{X}} \in \mathbb{R}^{d \times K}$ and defined a weighted Hopfield energy landscape:

$$E_r(\boldsymbol{\xi}) = \frac{1}{2} \|\boldsymbol{\xi}\|^2 - \frac{1}{\beta} \log \sum_{k=1}^K r_k \exp \left(\beta \mathbf{m}_k^\top \boldsymbol{\xi} \right) \quad (1)$$

where $\{\mathbf{m}_k\}_{k=1}^K$ are stored memory patterns, β is an inverse temperature parameter controlling the sharpness of retrieval, and $\{r_k\}_{k=1}^K$ are per-pattern multiplicity weights. When all $r_k = 1$, this reduces to the standard modern Hopfield energy; when $r_k = \rho > 1$ for a designated subset of patterns, the energy landscape is continuously deformed to favor that subset.

This energy defines a probability distribution that can be written exactly. Writing $\pi_r(\boldsymbol{\xi}) \propto$

570 $\exp[-\beta E_r(\boldsymbol{\xi})]$ and completing the square in each exponent,

$$\exp\left[-\frac{\beta}{2}\|\boldsymbol{\xi}\|^2 + \beta \mathbf{m}_k^\top \boldsymbol{\xi}\right] = \exp\left(\frac{\beta}{2}\|\mathbf{m}_k\|^2\right) \exp\left[-\frac{\beta}{2}\|\boldsymbol{\xi} - \mathbf{m}_k\|^2\right], \quad (2)$$

571 so the resulting sampling distribution is exactly a finite isotropic Gaussian mixture,

$$\pi_r(\boldsymbol{\xi}) = \sum_{k=1}^K q_k \mathcal{N}(\mathbf{m}_k, \beta^{-1}\mathbf{I}), \quad q_k \propto r_k \exp\left(\frac{\beta}{2}\|\mathbf{m}_k\|^2\right). \quad (3)$$

572 Each stored patient is the center of one component. The common component variance is β^{-1} ,
 573 and the multiplicity weights enter the mixture probabilities directly. The memories used here
 574 are normalized to unit length, so the exponential factor is shared by every component and $q_k =$
 575 $r_k / \sum_j r_j$ exactly. We could therefore sample this distribution directly: first choose k from q and
 576 then draw $\boldsymbol{\xi} \sim \mathcal{N}(\mathbf{m}_k, \beta^{-1}\mathbf{I})$. This procedure requires no Markov chain. The cohort analyzed in this
 577 work was generated with the Langevin implementation given next, and that is what the reported
 578 results describe. We added the direct sampler for the stochastic-attention target afterward as a
 579 check, using the same normalization, magnitude rescaling, and decoding steps (Table S16).

580 We sampled from this landscape using an Unadjusted Langevin Algorithm (ULA):

$$\boldsymbol{\xi}_{t+1} = (1 - \alpha) \boldsymbol{\xi}_t + \alpha \hat{\mathbf{X}} \text{softmax}\left(\beta \hat{\mathbf{X}}^\top \boldsymbol{\xi}_t + \log \mathbf{r}\right) + \sqrt{\frac{2\alpha}{\beta}} \boldsymbol{\varepsilon}_t \quad (4)$$

581 where α is a step size, $\boldsymbol{\varepsilon}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, and $\log \mathbf{r}$ is the element-wise log of the multiplicity vector. The
 582 deterministic component drove $\boldsymbol{\xi}_t$ toward a multiplicity-weighted combination of stored patterns,
 583 while the stochastic component added variation. The $\log r_k$ bias in the softmax logits interpolated
 584 between unconditioned generation ($\rho = 1$, all patterns equally weighted) and hard subset curation
 585 ($\rho \rightarrow \infty$, only designated patterns contribute). This is an inference-time operation that requires
 586 no change to the stored patterns.

587 The synthetic cohort analyzed in this work was produced by initializing each Langevin chain
 588 near a randomly selected stored pattern and retaining the final state after $T=2000$ steps. Because
 589 the starting point and chain length could affect the output, we compared four sampling procedures.
 590 These used the original initialization near a stored pattern, random initialization on the unit sphere,
 591 pooled samples from multiple chains after burn-in and thinning, and Metropolis-adjusted sampling.
 592 As a separate check, we also drew directly from the exact stochastic-attention mixture in Eq. (3),
 593 without using a Markov chain. The fidelity and membership-inference results were similar across all
 594 four procedures and the direct draws. This agreement indicated that the original implementation
 595 reproduced the evaluated behavior of the sampling distribution at β^* (Supplementary Fig. S9).

596 The inverse temperature β affects Hopfield retrieval and direct sampling in different ways. In
 597 the Hopfield update, small β spreads attention across many stored profiles and tends toward their
 598 weighted average, whereas large β concentrates attention on individual profiles. When sampling
 599 directly from Eq. (3), the stored profile is chosen according to q_k . For the unit-normalized profiles
 600 used here, these probabilities are independent of β and are set by the multiplicity weights. Instead,
 601 β controls the spread around the selected profile through the covariance $\beta^{-1}\mathbf{I}$. Small β therefore
 602 produces broader draws, whereas large β produces tighter draws. The attention-entropy transition
 603 used to define β^* marks the shift from attention shared across profiles to attention concentrated
 604 on individual profiles. We selected β^* from the transition in normalized attention entropy across a
 605 logarithmic sweep of β . Specifically, β^* was the grid point where the entropy curve bent downward
 606 most strongly, calculated as the most negative finite-difference second derivative with respect to

$\log \beta$ (Supplementary Methods). This choice depended on the relative positions of the K profiles in d dimensions, not on what their features represented. For this dataset ($K=23$, $d_{\text{PCA}}=18$), it gave $\beta^* \approx 2.94$; the random-pattern prediction $\sqrt{d} \approx 4.24$ served as a theoretical reference. A sensitivity analysis across $\beta/\beta^* \in [0.1, 3.0]$ showed that the selected value balanced novelty and fidelity (Table S1). For conditioned generation, we repeated the entropy calculation with the multiplicity bias included, giving a ρ -dependent $\beta^*(\rho)$.

We used the participation ratio to measure how strongly multiplicity weighting concentrated repeated draws around a few stored profiles. For unit-normalized memories, $q_k = r_k / \sum_j r_j$ and $K_{\text{eff}} = 1 / \sum_k q_k^2 = (\sum_k r_k)^2 / \sum_k r_k^2$. Thus, K_{eff} is the number of equally weighted components with the same concentration. Equal weights gave $K_{\text{eff}} = 23$; concentrating weight on three profiles moved it toward three. This was not an independent clinical sample size. To assign a target fraction f_{target} of component probability to a subgroup, we used $\rho = f_{\text{target}} K_{\text{bg}} / [K_{\text{des}}(1 - f_{\text{target}})]$, where K_{des} and K_{bg} are the subgroup and background profile counts.

5.3 Pipeline for Continuous Longitudinal Data

Applying SA to continuous longitudinal clinical data required several adaptations beyond the standard Hopfield sampling framework. We first concatenated each patient’s $n=72$ assay measurements across all three visits into a single vector of dimension $d_{\text{concat}} = 3n = 216$, so that each stored memory pattern encoded a complete longitudinal profile and generated synthetic patients would have internally consistent multi-visit trajectories. We then standardized each feature to zero mean and unit variance and applied PCA retaining 95% of the variance, reducing dimensionality from 216 to $d_{\text{PCA}}=18$. This compression served two purposes: it removed collinearity among the 216 features and yielded a memory-to-dimension ratio of $K/d_{\text{PCA}} \approx 1.28$, placing SA in a favorable operating regime where the number of stored patterns exceeded the dimensionality of the representation.

Standard SA operates on unit-norm patterns, which is appropriate for discrete data (e.g., one-hot protein sequences) but destroys the anisotropic variance structure of continuous clinical measurements, collapsing dispersion across all principal components to a unit sphere. We addressed this with a direction-magnitude decomposition. Before normalization, we recorded each pattern’s PCA-space norm $s_k = \|\mathbf{m}_k\|$; we then ran SA on unit-norm patterns to obtain a directional sample $\hat{\xi}$, drew a magnitude s from the empirical distribution of $\{s_k\}$, and reconstructed the rescaled sample as $\xi_{\text{rescaled}} = s \cdot \hat{\xi}$. This preserved the directional structure learned by the Hopfield energy while restoring the natural scale of variation. The rescaled PCA vector was mapped back to the original feature space via inverse PCA and destandardization, then split into three 72-dimensional visit records.

To generate condition-specific cohorts (e.g., PCOS-only), we set the multiplicity $r_k = \rho$ for patterns belonging to the target subgroup and $r_k = 1$ for all others, and sampled using the weighted Langevin update (Eq. 4). The multiplicity ρ was computed to place a target fraction $f_{\text{target}} = 0.80$ of the mixture-component probability on the designated subset, while retaining 20% probability on background patterns, via $\rho = f_{\text{target}} \cdot K_{\text{bg}} / (K_{\text{des}} \cdot (1 - f_{\text{target}}))$. Multiplicity parameters for all three subgroups are listed in Table S2. For the PCOS subgroup ($K_{\text{des}}=3$), this required $\rho \approx 26.7$, reducing the participation ratio to $K_{\text{eff}} \approx 4.6$. **For both unconditioned and conditioned generation, magnitudes were drawn from the empirical distribution of all 23 PCA-space norms. Subgroup weighting affected directional sampling, not the distribution of magnitudes.**

The complete generation pipeline is summarized in Algorithm 1. The full set of SA hyperparameters used for all experiments is listed in Table S3, with key operating values $\alpha = 0.01$, $T = 2,000$ Langevin iterations per sample, and random seed 42 for reproducibility. A control experiment comparing ULA with the Metropolis-Adjusted Langevin Algorithm (MALA) showed that

the discretization bias is negligible at this step size. Across 10 chains of 5,000 iterations the MALA acceptance rate was $99.9 \pm 0.03\%$, and mean post-burn-in energies and effective sample sizes were indistinguishable between the two samplers (Supplementary Table S8). The pipeline was implemented in Julia (v1.12) using the packages listed in the code repository. Generating 100 synthetic patients required approximately 0.1 seconds on a standard laptop (Apple M-series), comparable to MVN sampling (≈ 0.15 s), because SA operates in the 18-dimensional PCA space rather than the full 216-dimensional feature space. Each synthetic patient was generated by an independent Langevin chain (no shared state between patients); generation quality was stable across $T \in [500, 4000]$ iterations (median MRE 1.0–1.6%), indicating that $T=2,000$ was sufficient for convergence.

Algorithm 1 Multiplicity-Weighted SA for Longitudinal Patient Generation

Require: Patient data $\{\mathbf{x}_k^{(v)}\}$ for $k = 1, \dots, K$ patients, $v = 1, 2, 3$ visits

Require: Multiplicity vector $\mathbf{r}$ (all ones for unconditioned; $r_k = \rho$ for designated subset)

Ensure: N synthetic longitudinal patient profiles

```

1: Concatenate:  $\mathbf{x}_k \leftarrow [\mathbf{x}_k^{(1)}; \mathbf{x}_k^{(2)}; \mathbf{x}_k^{(3)}] \in \mathbb{R}^{216}$ 
2: Standardize:  $\mathbf{z}_k \leftarrow (\mathbf{x}_k - \boldsymbol{\mu}) \oslash \boldsymbol{\sigma}$ 
3: PCA:  $\mathbf{m}_k \leftarrow \mathbf{W}^\top \mathbf{z}_k \in \mathbb{R}^{18}$  ▷ 95% variance retained
4: Record norms:  $s_k \leftarrow \|\mathbf{m}_k\|$ ; normalize  $\hat{\mathbf{m}}_k \leftarrow \mathbf{m}_k / s_k$ 
5: Select  $\beta^*$  where the attention-entropy curve bends downward most strongly, including log  $\mathbf{r}$  when conditioned
6: for  $i = 1, \dots, N$  do
7:   Select anchor  $k_0$  from all patterns, or from the designated subset when conditioned
8:    $\hat{\boldsymbol{\xi}}_0 \leftarrow \hat{\mathbf{m}}_{k_0} + 0.01\boldsymbol{\varepsilon}_0$ ; normalize  $\hat{\boldsymbol{\xi}}_0$ 
9:   for  $t = 0, \dots, T - 1$  do ▷ Langevin dynamics
10:     $\hat{\boldsymbol{\xi}}_{t+1} \leftarrow (1 - \alpha)\hat{\boldsymbol{\xi}}_t + \alpha \hat{\mathbf{M}} \text{softmax}(\beta^* \hat{\mathbf{M}}^\top \hat{\boldsymbol{\xi}}_t + \log \mathbf{r}) + \sqrt{2\alpha/\beta^*} \boldsymbol{\varepsilon}_t$ 
11:   end for
12:   Draw magnitude:  $s \sim \{s_k\}_{k=1}^K$ 
13:   Rescale:  $\mathbf{m}_{\text{synth}} \leftarrow s \cdot \hat{\boldsymbol{\xi}}_T / \|\hat{\boldsymbol{\xi}}_T\|$ 
14:   Inverse PCA + destandardize:  $\mathbf{x}_{\text{synth}} \leftarrow \boldsymbol{\sigma} \odot (\mathbf{W} \mathbf{m}_{\text{synth}} + \bar{\mathbf{z}}) + \boldsymbol{\mu}$ 
15:   Split into visit records:  $\mathbf{x}_{\text{synth}}^{(1)}, \mathbf{x}_{\text{synth}}^{(2)}, \mathbf{x}_{\text{synth}}^{(3)}$ 
16: end for
```

5.4 Baseline and Validation Framework

As a baseline, we generated synthetic patients by fitting a single multivariate normal (MVN) distribution to the same $K=23$ concatenated 216-dimensional patient profiles. Both methods operated on the same concatenated representation. Each patient’s three visits were joined into one vector before generation, so that both methods had the opportunity to capture cross-visit dependencies. For SA, this structure was preserved through the Hopfield energy landscape operating on 18-dimensional PCA projections of the concatenated vectors. For MVN, the 216×216 sample covariance matrix had rank at most 22 and was regularized using Ledoit–Wolf shrinkage²¹, which inflated eigenvalues in the 194 null dimensions. To examine the role of the shared PCA representation, we also evaluated a fitted two-component diagonal-covariance Gaussian mixture, a kernel-density estimator, k -nearest-neighbor interpolation, and a Gaussian copula in the same 18-dimensional PCA space. Their implementations and shared evaluation pipeline are described in Supplementary Methods (PCA-space ablation). We then evaluated both SA and MVN synthetic patients through four pro-

gressively stringent validation levels. **Level 1 (Marginal Plausibility)** tested whether individual feature distributions matched, using per-feature means, standard deviations, and known physiological relationships. **Level 2 (Joint Structure)** tested whether the cross-visit covariance was preserved, by comparing full 216×216 correlation matrices, eigenvalue spectra, and PCA projections between real, SA, and MVN populations. **Level 3 (Conditional Structure)** tested whether rare subgroups could be amplified while preserving condition-specific signatures, by generating cohorts of 100 patients each from the Uncomplicated ($n=18$), PCOS ($n=3$), and Developed PE ($n=5$) subgroups using multiplicity-weighted sampling. **Level 4 (Mechanistic Consistency)** tested whether synthetic patients produced biologically plausible outputs under a *separately specified, generator-blind mechanistic model calibrated on the same real cohort*. We used the Hockin–Mann BZ2012 coagulation model^{6,17}, a system of 58 ODEs with 64 rate constants, in which 5 rate constants were calibrated on Visit 1 real patients at the population level and the remaining 59 were held at literature values. We ran the model on all real and synthetic patients under two conditions (TF-only and TF+TM) and compared predicted-versus-measured thrombin generation features. The primary metric was cloud overlap: the fraction of synthetic predicted/measured ratios falling within the 5th–95th percentile range of real ratios.

For the membership-inference analysis, we drew 40 random partitions of the 23 real profiles. For each partition, we repeated the preprocessing using 15 profiles, constructed the memory matrix, and generated 100 synthetic profiles; the other 8 real profiles were held out. We assigned each real profile a membership score equal to the negative Euclidean distance to its nearest synthetic profile in standardized 216-dimensional space, so a higher score indicated stronger evidence that the profile had been stored. A receiver operating characteristic curve compared the fraction of stored profiles correctly identified with the fraction of held-out profiles incorrectly identified as the score threshold changed. We calculated the area under this curve using the equivalent Mann–Whitney rank statistic, with ties contributing one half, and reported the mean across the 40 partitions. An AUC of 0.5 indicated chance discrimination, whereas 1.0 indicated perfect discrimination.

5.5 Robustness Tests

Results from one clinical cohort could not show how stochastic-attention generation behaved as training-set size, dimension, and nonlinear structure changed. We therefore used three tests. The covariance test compared stochastic attention with a sample-covariance Gaussian generator on low-rank Gaussian populations with known covariance. We generated one training cohort for every tested combination from the factor model $\mathbf{X} = \mathbf{Z}\mathbf{L}^\top + \epsilon$, where $\mathbf{Z} \in \mathbb{R}^{n \times r}$ and $\mathbf{L} \in \mathbb{R}^{p \times r}$ had independent standard-normal entries and ϵ had independent Gaussian entries with variance 0.1. This gave the known population covariance $\Sigma_{\text{pop}} = \mathbf{L}\mathbf{L}^\top + 0.1\mathbf{I}$. We tested every combination of $n \in \{8, 15, 23, 40, 80\}$, $p \in \{30, 120, 216\}$, and $r \in \{3, 5, 10\}$. For each training cohort, we standardized the features, retained enough principal components to explain 95% of the variance, constructed the stochastic-attention memory matrix from the resulting scores, and selected β separately using the same attention-entropy criterion as in the clinical analysis. Stochastic attention generated 100 profiles using ULA with $\alpha=0.01$, $T=2,000$, and empirical magnitude rescaling. The Gaussian generator used the training mean and sample covariance. For numerical sampling, diagonal jitter began at $10^{-8}\mathbf{I}$ and increased tenfold as needed through $10^{-2}\mathbf{I}$; if Cholesky sampling still failed, we sampled in the nonnegative eigenspace of the sample covariance. We applied no covariance shrinkage. The Gaussian generator also produced 100 profiles. Covariance error was $\|\hat{\Sigma}_{\text{gen}} - \Sigma_{\text{pop}}\|_F / \|\Sigma_{\text{pop}}\|_F$. Figure 4A reports Gaussian error divided by stochastic-attention error, averaged across the three latent ranks. A ratio above one favored stochastic attention.

The remaining tests asked how stochastic attention changed with less clinical information and

how it compared with the Gaussian generator on a nonlinear population. For the sample-size test, we drew ten subsets without replacement from the 23 complete patients at each $n \in \{20, 15, 10, 8\}$. For every subset, we repeated the standardization, PCA, and β selection, constructed a new stochastic-attention memory matrix, generated 100 profiles, and evaluated the generated cohort against the same 23-patient reference. The outcomes were relative cross-visit correlation error, pooled median marginal relative error, and TF-only mechanistic cloud overlap. Figure 4B shows the mean and standard deviation of cross-visit error across the ten subsets. For the nonlinearity test, we drew t uniformly from $[-1, 1]$ and defined an S-curve by $x_1 = t$ and $x_2 = c \sin(\pi t)$. We mapped the curve into $p=30$ or 120 dimensions with a new random orthonormal embedding for each replicate and added isotropic Gaussian noise with standard deviation 0.05. We used $n=23$ and curvature $c \in \{0, 0.25, 0.5, 1, 1.5, 2\}$, with five replicates per setting. From every training cohort, we constructed a stochastic-attention memory matrix and estimated the same sample-covariance Gaussian comparator; each generator produced 100 profiles. We projected generated profiles onto the known two-dimensional embedding plane and measured their mean distance from the curve, thereby excluding noise orthogonal to that plane. We also measured relative covariance error against the covariance of an independent 50,000-profile population draw. For the reported curvature threshold, we averaged stochastic-attention curve distance across the five replicates and both dimensions at each curvature and identified the first nonzero value at which this mean was at least twice its value for the straight-line population. All simulation cells used deterministic, cell-specific random-number streams derived from seed 42.

Ethics Statement

This study used de-identified secondary analysis measurements from samples collected longitudinally from women desiring a pregnancy. The primary prospective research study was approved by the Institutional Review Board at the University of Vermont. Written informed consent was obtained from all participants prior to enrollment.

Data Availability

The generated synthetic datasets and all intermediate validation outputs are available in the project repository at project GitHub repository. The real patient data were collected under IRB approval at the University of Vermont; de-identified data are available upon reasonable request to the corresponding author.

Code Availability

All code for synthetic patient generation, validation, figure reproduction, and mechanistic model calibration is available at project GitHub repository. The repository includes Julia source code, experiment scripts, and a Python baseline script for CTGAN/TVAE comparison.

Acknowledgments

This work was supported by NIH NHLBI R-33 HL 141787 (PIs Bernstein, Orfeo) and NIH NHLBI R01 HL 71944 (PI Bernstein).

Author Contributions

J.D.V. conceived the study, developed the SA pipeline and validation framework, performed computational analyses, and wrote the manuscript. M.C.B. and T.O. oversaw the secondary analysis measurements of coagulation and fibrinolysis markers, dynamic assays, and hormone measurements; provided domain expertise on coagulation biology and thrombin generation assays; and reviewed the manuscript. C.M. was involved in the enrollment of participants in the prospective research study and collection of clinical information of the participants. I.B. designed and oversaw the primary prospective research study of enrolled participants, provided clinical interpretation, and reviewed the manuscript. All authors approved the final version.

Competing Interests

The authors declare no competing interests.

References

- [1] Toy, S. U.S. maternal mortality hits highest level since 1965. *The Wall Street Journal* (2023). March 16, 2023.
- [2] Harris, E. US maternal mortality continues to worsen. *JAMA* **329**, 1248 (2023).
- [3] Chen, Y. *et al.* Pregnancy-related deaths in the US, 2018–2022. *JAMA Network Open* **8**, e254325 (2025).
- [4] Mann, K. G. Thrombin generation in hemorrhage control and vascular occlusion. *Circulation* **124**, 225–235 (2011).
- [5] Butenas, S., Orfeo, T. & Mann, K. G. Tissue factor in coagulation: Which? where? when? *Arteriosclerosis, Thrombosis, and Vascular Biology* **29**, 1989–1996 (2009).
- [6] Hockin, M. F., Jones, K. C., Everse, S. J. & Mann, K. G. A model for the stoichiometric regulation of blood coagulation. *Journal of Biological Chemistry* **277**, 18322–18333 (2002).
- [7] Bravo, M. C., Orfeo, T., Mann, K. G. & Everse, S. J. Modeling of human factor Va inactivation by activated protein C. *BMC Systems Biology* **6**, 45 (2012).
- [8] Hellgren, M. Hemostasis during normal pregnancy and puerperium. *Seminars in Thrombosis and Hemostasis* **29**, 125–130 (2003).
- [9] Grouzi, E., Pouliakis, A., Aktypi, A. *et al.* Pregnancy and thrombosis risk for women without a history of thrombotic events: a retrospective study of the real risks. *Thrombosis Journal* **20** (2022).
- [10] Bremme, K. A. Haemostatic changes in pregnancy. *Best Practice & Research Clinical Haematology* **16**, 153–168 (2003).
- [11] Azziz, R. *et al.* The prevalence and features of the polycystic ovary syndrome in an unselected population. *Journal of Clinical Endocrinology & Metabolism* **89**, 2745–2749 (2004).

- [12] Palomba, S., Santagni, S., Falbo, A. & La Sala, G. B. Complications and challenges associated with polycystic ovary syndrome: current perspectives. *International Journal of Women's Health* **7**, 745–763 (2015).
- [13] Abalos, E. *et al.* Pre-eclampsia, eclampsia and adverse maternal and perinatal outcomes: a secondary analysis of the World Health Organization multicountry survey on maternal and newborn health. *BJOG: An International Journal of Obstetrics & Gynaecology* **121**, 14–24 (2014).
- [14] Dusse, L. M., Rios, D. R. A., Pinheiro, M. B., Cooper, A. J. & Lwaleed, B. A. Pre-eclampsia: relationship between coagulation, fibrinolysis and inflammation. *Clinica Chimica Acta* **412**, 17–21 (2011).
- [15] Luan, D., Zai, M. & Varner, J. D. Computationally derived points of fragility of a human cascade are consistent with current therapeutic strategies. *PLoS Computational Biology* **3**, e142 (2007).
- [16] Luan, D., Szlam, F., Tanaka, K. A., Barie, P. S. & Varner, J. D. Ensembles of uncertain mathematical models can identify network response to therapeutic interventions. *Molecular BioSystems* **6**, 2272–2286 (2010).
- [17] Brummel-Ziedins, K. E. *et al.* The prothrombotic phenotypes in familial protein c deficiency are differentiated by computational modeling of thrombin generation. *PLoS ONE* **7**, e44378 (2012).
- [18] Hastie, T., Tibshirani, R. & Friedman, J. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (Springer, 2009), 2nd edn.
- [19] Gonzales, A., Guruswamy, G. & Smith, S. R. Synthetic data in health care: A narrative review. *PLOS Digital Health* **2**, e0000082 (2023).
- [20] Kühnel, L. *et al.* Synthetic data generation for a longitudinal cohort study – evaluation, method extension and reproduction of published data analysis results. *Scientific Reports* **14**, 14412 (2024).
- [21] Ledoit, O. & Wolf, M. A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis* **88**, 365–411 (2004).
- [22] Goodfellow, I. *et al.* Generative adversarial nets. In *Advances in Neural Information Processing Systems*, vol. 27, 2672–2680 (2014).
- [23] Kingma, D. P. & Welling, M. Auto-encoding variational Bayes. In *International Conference on Learning Representations (ICLR)* (2014).
- [24] Xu, L., Skoularidou, M., Cuesta-Infante, A. & Veeramachaneni, K. Modeling tabular data using conditional GAN. In *Advances in Neural Information Processing Systems*, vol. 32, 7335–7345 (2019).
- [25] De Mitri, O., Wang, R. & Huber, M. F. Generative adversarial networks with limited data: A survey and benchmarking. *arXiv preprint arXiv:2504.05456* (2025).
- [26] Ramsauer, H. *et al.* Hopfield networks is all you need. In *International Conference on Learning Representations (ICLR)* (2021).

- [27] Varner, J. D. Conditioning protein generation via Hopfield pattern multiplicity. *arXiv preprint arXiv:2603.20115* (2026).
- [28] Bravo, M. *et al.* BSTModelKit.jl: Building biochemical systems theory models of individualized coagulation dynamics during pregnancy. JuliaCon 2023, <https://www.youtube.com/watch?v=IGtUqMLHBvQ> (2023). Conference proceedings talk.
- [29] Ghetmiri, D. E., Cohen, M. J. & Menezes, A. A. Personalized modulation of coagulation factors using a thrombin dynamics model to treat trauma-induced coagulopathy. *npj Systems Biology and Applications* **7**, 44 (2021).
- [30] Varner, J. D. Training-free generation of protein sequences from small family alignments via stochastic attention. *arXiv preprint arXiv:2603.14717* (2026).
- [31] Alswaidan, A. & Varner, J. D. Stochastic attention via Langevin dynamics on the modern Hopfield energy. *arXiv preprint arXiv:2603.06875* (2026).
- [32] Mould, D. R. & Upton, R. N. Basic concepts in population modeling, simulation, and model-based drug development—part 2: introduction to pharmacokinetic modeling methods. *CPT: Pharmacometrics & Systems Pharmacology* **2**, e38 (2013).
- [33] Chase, J. G. *et al.* Physiological modeling, tight glycemic control, and the ICU clinician: what are models and how can they affect practice? *Annals of Intensive Care* **1**, 11 (2011).
- [34] Ribba, B. *et al.* A tumor growth inhibition model for low-grade glioma treated with chemotherapy or radiotherapy. *Clinical Cancer Research* **18**, 5071–5080 (2012).
- [35] Yizhak, K., Chaneton, B., Gottlieb, E. & Ruppin, E. Modeling cancer metabolism on a genome scale. *Molecular Systems Biology* **11**, 817 (2015).
- [36] Shan, M., Dai, D., Vudem, A., Varner, J. D. & Stroock, A. D. Multi-scale computational study of the Warburg effect, reverse Warburg effect and glutamine addiction in solid tumors. *PLoS Computational Biology* **14**, e1006584 (2018).
- [37] McLean, K. C., Bernstein, I. M. & Brummel-Ziedins, K. Tissue factor dependent thrombin generation across pregnancy. *American Journal of Obstetrics and Gynecology* **207**, 135.e1–135.e6 (2012).
- [38] Hale, S. A. *et al.* Coagulation and fibrinolytic system protein profiles in women with normal pregnancies and pregnancies complicated by hypertension. *Pregnancy Hypertension* **2**, 152–157 (2012).
- [39] Psinos, R. B. C., Morris, E. A., McBride, C. A. & Bernstein, I. M. Association of pre-pregnancy subclinical insulin resistance with cardiac dysfunction in healthy nulliparous women. *Pregnancy Hypertension* **26**, 11–16 (2021).
- [40] Bernstein, I. M., Hale, S. A. & Badger, G. J. Differences in cardiovascular function comparing prior preeclampsia with nulliparous controls. *Pregnancy Hypertension* **6**, 320–326 (2016).

Table 1: **Demographics and clinical characteristics of the 23 patients with complete longitudinal data across all three visits.** Values are reported as mean $\pm$ SD, median (range), or n (%) as appropriate.

Characteristic	Mean $\pm$ SD	Median (Range)	n (%)
Age at enrollment (yrs)	30.2 $\pm$ 5.1	29 (20–41)	
Race			
White			22 (96)
Not disclosed			1 (4)
Prepregnancy BMI	26.6 $\pm$ 5.3	25 (19.0–37.8)	
Parity		0 (0–2)	
Nulliparous			16 (70)
Parous			7 (30)
Enrollment condition			
Healthy nulliparous			14 (61)
PCOS			3 (13)
Prior PE			6 (26)
Cycle day at V1	9.2 $\pm$ 3.6	10 (3–14)	
Gestational age at V2 (days)	89.0 $\pm$ 5.5		
Gestational age at V3 (days)	217.0 $\pm$ 4.0		
<i>Enrollment cohort $\times$ pregnancy outcome</i>			
	Uncomplicated	Developed PE	Total
Healthy nulliparous	12	2	14
Prior PE	4	2	6
PCOS	2	1	3
Total	18	5	23

Table 2: Per-visit mean relative error (MRE) between real and SA-generated synthetic patients for representative coagulation features, computed as $|\bar{x}_{\text{synth}} - \bar{x}_{\text{real}}|/|\bar{x}_{\text{real}}|$.

Feature	Visit 1	Visit 2	Visit 3
Factor II (%)	0.018	0.017	0.007
Factor VIII (%)	0.018	0.008	0.005
AT (%)	0.025	0.001	0.008
Fibrinogen (mg/dL)	0.005	0.009	<0.001
TF Peak (nM)	0.025	0.002	0.017
TF ETP (nM·min)	0.013	0.008	0.023




Figure 1: **Physiological correlations.** Scatter plots of coagulation factor levels versus thrombin generation parameters across all three visits. Filled markers are real patients; open markers are SA-generated synthetic patients. Colors encode visit: blue (V1/baseline), green (V2/first trimester), orange (V3/third trimester). Real and synthetic patients occupy the same joint regions across all four pairwise relationships. AT inversely correlates with TF-initiated peak thrombin (top left), FVIII positively correlates with peak (top right), and both FII (bottom left) and AT (bottom right) show the expected relationships with ETP. The visit-dependent shift in these relationships, reflecting the pregnancy-driven increase in procoagulant factors, is also preserved in the synthetic cohort.




Figure 2: **Pregnancy-related longitudinal trajectories.** Mean $\pm$ standard deviation bands for six key coagulation features across visits (BL = baseline, 1st = first trimester, 3rd = third trimester) for real (blue) and SA-generated synthetic (orange) patients. The characteristic pregnancy-driven increases in fibrinogen, Factor VIII, and vWF are captured, as are the stable-to-declining patterns in Factor II and AT.




Figure 3: **Cross-visit correlation structure.** Top row: full 216×216 Pearson correlation matrices for Real ($K=23$, left), SA ($N=100$, center), and MVN ($N=100$, right) populations. Each matrix is organized as a 3×3 grid of 72×72 blocks (delineated by black lines), where the diagonal blocks capture within-visit feature correlations (V1–V1, V2–V2, V3–V3) and the off-diagonal blocks capture cross-visit dependencies (e.g., V1–V3 correlations between baseline and third-trimester Factor VIII). Stochastic attention preserves both the within-visit and cross-visit block structure. Bottom row: pairwise residual matrices (element-wise subtraction), all plotted on the same ± 1 color scale as the top row. The SA–Real residual shows small, unstructured deviations. The MVN–Real residual is smoother, particularly in the off-diagonal blocks, reflecting the Ledoit–Wolf regularization that systematically shrinks cross-visit correlations toward zero. The SA–MVN residual shows the largest differences in these off-diagonal blocks, where SA retains longitudinal dependencies that MVN’s rank-deficient covariance estimate cannot represent. A version with amplified residual color scale (± 0.5) and Frobenius norms is provided in Fig. S1.




Figure 4: **Population recovery and clinical-cohort subsampling.** (A) For each low-rank Gaussian setting, the stochastic-attention memory matrix was constructed and a sample-covariance Gaussian generator was estimated from the same simulated training cohort; both generators produced 100 profiles. Color represented the ratio of their relative Frobenius errors against the known population covariance (Gaussian error divided by stochastic-attention error), averaged over latent ranks 3, 5, and 10. Values above one (blue) indicated lower error for stochastic attention. Because the target covariance was fixed before the training cohort was drawn, this measured population recovery rather than reproduction of the training data. (B) A stochastic-attention memory matrix was constructed from each of ten random subsets of the clinical cohort at $n=20, 15, 10$, and 8 , and each matrix generated 100 profiles. Points and bars represented the mean and standard deviation of the relative cross-visit correlation error against the same full 23-patient reference; the star represented generation using all 23 patients.




Figure 5: Condition-specific feature preservation. Grouped bar charts comparing real (dark bars) and SA-generated synthetic (light bars) patient means ($\pm$ SD error bars) for eight coagulation features, shown separately for each clinical subgroup: Uncomplicated ($n=18$, left, blue), PCOS ($n=3$, center, orange), and Developed PE ($n=5$, right, red). Each bar is the mean of that assay pooled across all three visits (V1–V3); factor activities (FVIII, vWF, FIX, FV, α_2 -AP, FVII, plasminogen) are expressed as % of normal pooled reference plasma, and fibrinogen (Fbgn) in mg/dL. Gray percentages indicate the mean relative error (MRE). Green annotations show a descriptive bootstrap Mann–Whitney diagnostic: the fraction of 1,000 bootstrap replicates (synthetic subsampled to n_{real}) in which $p > 0.05$, meaning that this test detected no difference at the available sample size. This is not an equivalence test. Features achieving $\geq 90\%$ are shown in green; $< 90\%$ in red. Across all 24 feature–condition pairs, 20 (83%) had no difference detected in $\geq 90\%$ of replicates (median: 98.6%). The four pairs below this threshold were FIX and plasminogen in the PCOS and Developed PE subgroups, where small subgroup sizes and large inter-patient variability reduced test power.




Figure 6: **Mechanistic validation (TF-only)**. Top row: BZ2012 ODE-predicted TGA values (vertical axis) versus dataset TGA values (horizontal axis) for five thrombin generation parameters. The dataset values are the TGA measurements from each patient’s record, present for both real and synthetic patients. The ODE-predicted values are computed by running each patient’s coagulation factor levels through the 58-species BZ2012 model. The dashed line indicates where the ODE model would perfectly reproduce the dataset value ($y=x$); systematic deviations from this line (e.g., lagtime underprediction) reflect ODE model bias, not synthetic data failure. Filled markers are real patients, open markers are SA-generated synthetic patients, colored by visit: blue (V1/baseline, training set), green (V2/first trimester), orange (V3/third trimester). Real and synthetic patients occupy the same cloud with the same pattern of ODE model bias. Bottom row: ODE-predicted/dataset ratio distributions for real (blue) and synthetic (orange) patients. The distributions overlap substantially. Each panel is annotated with cloud overlap (OL), the two-sample Kolmogorov–Smirnov statistic (D), and p -value. All five features had $p > 0.35$; these tests did not detect a difference between real and synthetic ratio distributions.




Figure 7: **Downstream utility: synthetic and real calibrations on held-out real profiles.** Each panel compares the BZ2012 ODE predictions for one TGA feature when the model is calibrated on real V1 patients (horizontal axis) versus synthetic V1 patients (vertical axis), evaluated on the same held-out real V2 and V3 patients. Points near the $y=x$ line indicate similar predictions from the two calibrations. Green: V2 (first trimester); orange: V3 (third trimester). The synth-calibrated model had slightly lower median relative error on all five features (overall synthetic-to-real error ratio 0.94). Both calibrations used 11 random restarts; all starts converged.

Supplementary Information

Supplementary Methods

PCA-space ablation

We compared the canonical stochastic-attention cohort with four baseline generators operating on the same unnormalized 18-dimensional PCA coordinates. Each baseline generated 100 profiles with a separate random-number stream initialized using seed 42. The fitted two-component Gaussian-mixture baseline used diagonal covariance matrices, 100 expectation-maximization iterations, and a variance floor of 10^{-3} . The kernel-density sampler chose a stored PCA point uniformly and added Gaussian noise with the empirical PCA covariance and Scott bandwidth $h = K^{-1/(d+4)}$; $10^{-8}\mathbf{I}$ was added for numerical factorization. The nearest-neighbor sampler chose a stored point, selected one of its three nearest neighbors uniformly, and interpolated between them with a coefficient drawn from $\text{Uniform}(0, 1)$. The Gaussian copula transformed empirical marginal ranks to normal scores, estimated their correlation matrix with $10^{-6}\mathbf{I}$ added for factorization, and mapped sampled scores back through the empirical marginal quantiles. All samples were passed through the same inverse-PCA and destandardization pipeline and evaluated with the same marginal, cross-visit, mechanistic, and novelty calculations.

Selecting the inverse-temperature operating point

For each candidate β , we computed the Shannon entropy of $\text{softmax}(\beta \hat{\mathbf{X}}^\top \boldsymbol{\xi})$ at 20 stored-pattern probes, averaged the values, and divided by $\log K$. Zero means one pattern carries all attention; one means all patterns have equal attention. We evaluated 80 logarithmically spaced values from 0.1 to 1000 and selected β^* at the most negative finite-difference second derivative with respect to $\log \beta$. This marked the shift from shared to concentrated attention. For conditioned sampling, we repeated the calculation with the $\log r_k$ bias, producing $\beta^*(\rho)$. The random-pattern scale $\sqrt{d}$ was the theoretical reference.

Supplementary Tables

Table S1: **Inverse-temperature sensitivity.** Generation metrics across β/β^* , where $\beta^* \approx 2.94$ was selected from the attention-entropy transition in the 18-dimensional PCA space. The bold row marks the reported value. Med MRE is the median relative error of feature means; novelty is $1 - \max_k \cos(\hat{\xi}, \hat{m}_k)$.

β/β^*	β	Noise scale	PC1 std ratio	Novelty	Med MRE (%)
0.1	0.29	0.261	0.538	0.542	0.7
0.3	0.88	0.151	0.550	0.536	0.7
0.5	1.47	0.117	0.562	0.528	0.7
1.0	2.94	0.082	0.589	0.502	0.6
1.5	4.41	0.067	0.613	0.474	0.7
2.0	5.88	0.058	0.638	0.446	0.7
3.0	8.82	0.048	0.689	0.403	0.8

Table S2: **Multiplicity weighting parameters for conditional generation.** For the unit-normalized memories used here, the multiplicity ratio ρ was computed to place a target fraction $f_{\text{target}} = 0.80$ of finite-mixture component probability on the designated patterns. K_{eff} is the participation ratio of those component probabilities and describes their concentration across repeated draws.

Condition	K_{sub}	K_{bg}	ρ	f_{eff}	K_{eff}
Uncomplicated	18	5	1.11	0.80	23.0
PCOS	3	20	26.67	0.80	4.6
Developed PE	5	18	14.40	0.80	7.7

Table S3: **SA generation hyperparameters.** The same values were used for unconditioned and conditioned generation except that β^* was recomputed with the multiplicity bias for conditioned cohorts.

Parameter	Value	Description
α	0.01	Langevin step size (see Table S8)
T	2,000	Langevin iterations per sample
β^*	2.94	Entropy-transition point (unconditioned)
PCA threshold	95%	Variance retained
d_{PCA}	18	PCA dimensionality
N	100	Synthetic patients per cohort
Random seed	42	For reproducibility
f_{target}	0.80	Designated-component probability (conditioned)

Table S4: **BZ2012 ODE model calibration.** Five rate constants were calibrated on Visit 1 real patients ($n=23$) using bounded Nelder–Mead optimization in log-scale-factor space. The cost function was the mean normalized relative squared error across five TGA features under TF-only conditions. Bounds: $\pm \log(100)$. Convergence: $f_{\text{tol}} = 10^{-6}$, 500 iterations, 11 random restarts.

Rate constant	Role	Default	Calibrated	Scale
prothrombinase k_{cat}	Thrombin burst	63.5 s^{-1}	21.94 s^{-1}	$0.35 \times$
intrinsic Xase k_{cat}	Amplification Xa	8.2 s^{-1}	0.169 s^{-1}	$0.021 \times$
extrinsic Xase k_{cat}	Initiation Xa	6.0 s^{-1}	93.2 s^{-1}	$15.5 \times$
PC activation k_{cat}	Protein C activation	0.41 s^{-1}	0.890 s^{-1}	$2.17 \times$
mIIa conversion k	mIIa $\rightarrow$ IIa	2.3×10^8	1.15×10^9	$5.01 \times$

Units for mIIa conversion: $\text{M}^{-1}\text{s}^{-1}$

Table S5: **Summary of mean relative errors across all 72 features.** Per-visit distribution of MREs between real and SA-generated synthetic patient means. The full feature-by-feature table (72 features $\times$ 3 visits = 216 entries) is provided in the code repository as `paper_summary_statistics.csv`.

Visit	Median	Mean	25th pct.	75th pct.	Maximum	MRE below 5%
V1 (baseline)	0.022	0.029	0.008	0.039	0.158	59/72 (81.9%)
V2 (1st tri)	0.009	0.021	0.004	0.022	0.224	65/72 (90.3%)
V3 (3rd tri)	0.011	0.016	0.005	0.024	0.062	69/72 (95.8%)
Pooled (all 3 visits)	0.012	0.022	0.005	0.028	0.224	193/216 (89.4%)

Pooled median MRE 95% bootstrap CI: (0.98%, 1.60%) over 10,000 resamples (seed 42).

Table S6: **Comparison with deep generative baselines.** CTGAN and TVAE were trained on the same 90 per-visit records (23 patients $\times$ 3 visits, 72 features) at three epoch counts. Results for SA and MVN are shown for reference. The CTGAN model had median MRE near 19%, consistent with mode collapse at this sample size. The TVAE model matched marginal fidelity at high epoch counts but could not capture cross-visit covariance or perform conditional generation.

Method	Epochs	Med MRE	Med KS	Corr MAE
SA	–	0.019	0.169	0.200
MVN	–	0.024	0.161	0.090
CTGAN	300	0.189	0.364	0.261
CTGAN	1,000	0.195	0.349	0.265
CTGAN	3,000	0.187	0.341	0.180
TVAE	300	0.037	0.174	0.152
TVAE	1,000	0.026	0.248	0.082
TVAE	3,000	0.018	0.224	0.092

Table S7: **PCA variance threshold sensitivity analysis.** SA generation quality is stable across thresholds from 85% to 99%. Median MRE remains between 1.0% and 1.7% across all thresholds. The 95% threshold used in the main text (bold) is not a critical design choice.

Threshold	d_{PCA}	K/d	β^*	Novelty	Med MRE	Corr MAE
85%	13	1.77	2.94	0.420	0.0098	0.124
90%	15	1.53	2.94	0.477	0.0160	0.126
95%	18	1.28	2.94	0.500	0.0124	0.143
97.5%	20	1.15	2.94	0.543	0.0120	0.135
99%	21	1.10	2.94	0.541	0.0172	0.138

Table S8: **ULA and MALA sampling diagnostics.** Ten chains of 5,000 iterations each were run with both the Unadjusted Langevin Algorithm (ULA) and the Metropolis-Adjusted Langevin Algorithm (MALA) at step size $\alpha = 0.01$ and $\beta^* = 2.94$ on the 18-dimensional PCA memory ($K=23$ patterns). The MALA acceptance rate was 99.9%, and the energy and effective sample-size estimates were similar between methods, indicating little discretization bias. [This comparison bounded the discretization error of the reported sampler. The direct reference for the target itself was the analytic sampler of Eq. \(3\), introduced below.](#)

Metric	ULA	MALA
Acceptance rate (%)	– (no rejection)	99.9 ± 0.03
Mean energy (post-burn-in)	1.932 ± 0.141	1.886 ± 0.231
τ_{int}	109.5 ± 49.5	106.7 ± 30.4
Effective sample size	41.8 ± 14.2	40.0 ± 10.2
Burn-in: 1,000 iterations discarded. τ_{int} : integrated autocorrelation time estimated with a windowed estimator (threshold 0.05). $\text{ESS} = (T - T_{\text{burn}})/\tau_{\text{int}}$.		

Table S9: **Proximity to stored profiles.** Two measures compared synthetic and stored profiles. Novelty is the angular distance between a synthetic profile and its nearest stored pattern at $\beta = \beta^*$. Nearest-neighbor distances are computed in the per-visit standardized 72-dimensional feature space.

Metric	Value
Mean novelty score, $1 - \max_k \cos(\hat{\xi}, \hat{m}_k)$	0.50
Median synth–real nearest-neighbor distance (std. units)	6.14
Median real–real nearest-neighbor distance (std. units)	6.87
Distance ratio (synth–real / real–real)	0.89

Table S10: **Descriptive bootstrap Mann–Whitney diagnostic per feature and condition.** For each of 24 feature–condition pairs, we subsampled the synthetic cohort to match the real sample size, applied a Mann–Whitney U test, and repeated 1,000 times. Frac. $p > 0.05$ is the fraction of replicates in which the test detected no difference at the available sample size. This power-limited diagnostic is not an equivalence test. Twenty of 24 pairs had $p > 0.05$ in $\geq 90\%$ of replicates; the four below-threshold pairs (italicized) involved high-variance features in the smallest subgroups.

Condition	Feature	MRE	Frac. $p > 0.05$
Uncomplicated ($n=18$)	FVIII	0.057	0.952
Uncomplicated ($n=18$)	vWF	0.022	0.999
Uncomplicated ($n=18$)	FIX	0.018	0.998
Uncomplicated ($n=18$)	FV	0.006	0.996
Uncomplicated ($n=18$)	α_2 -AP	0.019	1.000
Uncomplicated ($n=18$)	FVII	0.057	0.941
Uncomplicated ($n=18$)	Fbgn	0.017	0.987
Uncomplicated ($n=18$)	Plgn	0.020	0.988
PCOS ($n=3$)	FVIII	0.109	0.985
PCOS ($n=3$)	vWF	0.150	0.959
PCOS ($n=3$)	<i>FIX</i>	<i>0.121</i>	<i>0.873</i>
PCOS ($n=3$)	FV	0.057	0.944
PCOS ($n=3$)	α_2 -AP	0.051	0.987
PCOS ($n=3$)	FVII	0.029	0.999
PCOS ($n=3$)	Fbgn	0.043	0.992
PCOS ($n=3$)	<i>Plgn</i>	<i>0.153</i>	<i>0.615</i>
Developed PE ($n=5$)	FVIII	0.097	0.932
Developed PE ($n=5$)	vWF	0.010	0.986
Developed PE ($n=5$)	<i>FIX</i>	<i>0.108</i>	<i>0.780</i>
Developed PE ($n=5$)	FV	0.006	0.987
Developed PE ($n=5$)	α_2 -AP	0.062	0.950
Developed PE ($n=5$)	FVII	0.045	0.993
Developed PE ($n=5$)	Fbgn	0.048	0.996
Developed PE ($n=5$)	<i>Plgn</i>	<i>0.093</i>	<i>0.814</i>

Median Frac. $p > 0.05$ across all 24 pairs: 0.986. Failing pairs (< 0.90 , italicized): PCOS–FIX, PCOS–Plgn, DPE–FIX, DPE–Plgn.

Table S11: **Mechanistic validation diagnostics per feature and condition.** Cloud overlap is the fraction of synthetic ODE-predicted/measured ratios falling within the 5th–95th percentile of the corresponding real distribution. KS D and KS p are from a two-sample Kolmogorov–Smirnov test on the same ratio distributions. Real $n=69$ (23 patients $\times$ 3 visits); synthetic $n=300$ (100 patients $\times$ 3 visits).

Condition	TGA Feature	Cloud overlap	KS D	KS p
TF-only	Lagtime	0.92	0.112	0.48
TF-only	Peak	0.91	0.081	0.86
TF-only	T.Peak	0.92	0.102	0.61
TF-only	Max Rate	0.91	0.083	0.84
TF-only	ETP	0.83	0.123	0.36
TF+TM	Lagtime	0.92	0.127	0.32
TF+TM	Peak	0.88	0.079	0.88
TF+TM	T.Peak	0.91	0.110	0.51
TF+TM	Max Rate	0.88	0.096	0.67
TF+TM	ETP	0.89	0.112	0.48

Under TF-only conditions, the cloud overlap range was 0.83–0.92. Under TF+TM conditions, the range was 0.88–0.92. Pooled across both conditions and all five features, the overlap was 0.902. KS D range across both conditions: 0.079–0.127, all $p > 0.30$.

Table S12: **Downstream utility: per-feature held-out V2/V3 performance.** Median relative error of the BZ2012 ODE-predicted TGA values against the held-out real V2 and V3 measurements, for the model calibrated on real V1 patients ($K=23$) and the model calibrated on synthetic V1 patients ($N=100$). The synth-calibrated model achieves slightly lower error on every feature. Synth/Real ratio < 1 favors the synth-calibrated model.

TGA Feature	Real-cal MRE	Synth-cal MRE	Synth/Real
Lagtime	0.165	0.162	0.98 $\times$
Peak	0.097	0.087	0.90 $\times$
T.Peak	0.330	0.306	0.93 $\times$
Max Rate	0.286	0.259	0.91 $\times$
ETP	0.197	0.183	0.93 $\times$
Overall median	0.201	0.189	0.94$\times$

Table S13: **PCA-space ablation.** Five generators used the same $d=18$ PCA subspace and evaluation. PCA+GMM is the fitted two-component diagonal-covariance baseline, not the exact stochastic-attention mixture. MRE measures marginal fidelity. Mechanistic overlap is the fraction of pooled synthetic predicted-to-recorded TGA ratios within the real 5th–95th percentile range. Median novelty is $1 - \max_k \cos(\hat{\xi}, \hat{\mathbf{m}}_k)$; Frac. novel is the fraction above 0.2. Cross-visit Frobenius error measures deviation from the real correlation matrix. Lower error indicates closer in-sample reproduction and can reward memorization.

Method	MRE (%)	Cross-visit Frob.	Mech. overlap	Median novelty	Frac. novel
SA	1.24	0.641	0.902	0.514	1.00
PCA+GMM	1.13	0.422	0.918	0.497	0.94
PCA+KDE	1.94	0.337	0.823	0.249	0.65
PCA+kNN	1.63	0.587	0.938	0.051	0.05
PCA+Copula	1.34	0.417	0.934	0.462	1.00

Table S14: **Convex-hull geometry and plausibility at $K=23$ in $d=18$.** We compared 100 synthetic profiles with leave-one-out real profiles using hulls of 22 real points. For each synthetic profile, we omitted its nearest real profile. The radial overshoot $1/t^*$ used the point where a ray from the hull centroid exited the hull. All profiles lay outside their comparison hulls. Median hull distance was 11.0 for synthetic profiles and 11.9 for real profiles ($p=0.07$); radial overshoot was smaller for synthetic profiles ($p < 10^{-4}$). These descriptive tests treated the synthetic profiles as independent, so their significance was optimistic. Hull membership had little resolution because every real profile formed an outer corner, and a unit-normalized generated direction could remain inside only by matching a stored direction. The biological check required nonnegative values and limited coagulation-factor and TGA values to five real-cohort standard deviations from the mean. Fifty-four synthetic profiles passed. The other 46 contained negative estradiol or progesterone values: estradiol was negative in 19 profiles and progesterone in 37, with 10 overlapping; the minima were -2109.8 and -75.8 . No variable was constrained to be positive, and linear decoding had unbounded support. Hormones were closer to zero relative to their observed variation than the coagulation factors, which explained why only hormones crossed zero in these draws but did not guarantee positive factors in other draws. All nine factor inputs to BZ2012 were positive, so all 100 profiles entered the mechanistic check. The unconstrained MVN baseline produced negative hormones in 43 profiles, compared with 46 for stochastic attention.

Quantity	Synthetic ($n=100$)	Real, leave-one-out ($n=23$)
Fraction inside hull	0.00	0.00
Median distance to hull	11.0	11.9
Median radial overshoot $1/t^*$	10.5	15.2
<i>Hull geometry of the real cohort</i>		
Median patient radius		13.8
Median hull extent, generic direction		2.0
Largest t^* attained on the unit sphere		0.231
<i>Plausibility of the out-of-hull synthetic patients</i>		
Biological constraints satisfied		54/100
Mechanistic band, per feature		83 to 92%
Mechanistic band, all features and visits		41/100
Both biological and strict mechanistic criteria		24/100

Table S15: **Membership inference recovered memory membership, not unseen patients.** Median nearest-neighbor distances in the standardized 216-dimensional space, and the test’s discrimination. For each split, we constructed the memory matrix from 15 patients and scored each real record by its distance to the resulting synthetic set, treating the 15 included patients as members and the 8 held-out patients as non-members. Every quantity came from the same protocol. We drew 40 partitions at random with the generation seed held fixed, so that only the partition varied, and pooled the distances across all partitions. Non-members sat at the same distance from the synthetic cohort as unrelated real patients sat from one another. The synthetic cohort carried no proximity signal about patients the generator never saw, and the discrimination came entirely from members sitting closer. The area under the receiver operating characteristic curve (AUC), calculated using the equivalent Mann–Whitney rank statistic, measured how often the test assigned stronger membership evidence to a randomly chosen member than to a randomly chosen non-member; 0.5 indicated no discrimination and 1.0 indicated perfect discrimination.

Quantity	Value
<i>Median nearest-neighbor distance</i>	
Memory-matrix member to nearest synthetic	10.7
Held-out non-member to nearest synthetic	15.4
Real to nearest real, leave-one-out	15.9
Synthetic to nearest real	13.8
<i>Membership-inference AUC over 40 random splits</i>	
Mean	0.97
Standard deviation	0.02
Median	0.98
Minimum	0.90
Maximum	1.00
Splits reaching perfect separation	9/40
Splits with AUC above 0.9	39/40

891 **Supplementary Figures**




Figure S1: **Cross-visit correlation structure (amplified residual scale)**. Same layout as the main-text correlation figure, but with the bottom-row residual panels plotted on a ± 0.5 color scale (vs. ± 1.0 in the main text) to reveal fine structure. Frobenius norms of each residual matrix are shown in the panel titles.




Figure S2: **Covariance eigenvalue spectrum.** Log-scale eigenvalues of the 216×216 sample covariance for real ($K=23$, rank 22), SA ($N=100$, rank 18), and MVN ($N=100$) populations. Dashed vertical lines mark the SA PCA truncation point (component 18) and the data rank limit (component 22). Ledoit–Wolf regularization for MVN inflates eigenvalues beyond component 22 (green shaded region), introducing spurious variance in 194 dimensions where the data contain no signal.




Figure S3: **Conditioned generation: PCA projection (pooled)**. Large markers: real patients (blue = Uncomplicated, $n=18$; orange = PCOS, $n=3$; red = Developed PE, $n=5$). Small markers: synthetic patients ($N=100$ per condition). With only 3 PCOS and 5 Developed PE patients, the subgroups do not form clearly separable clusters in the first two principal components, but the synthetic cohorts concentrate around their respective real counterparts.




Figure S4: **Conditioned generation: PCA projection by visit**. Same as Fig. S3 but separated by visit (V1, V2, V3). Condition-specific clustering is maintained within each individual visit.




Figure S5: **Mechanistic validation (TF+TM condition).** Same layout as main-text Fig. 6 but under TF+TM conditions. The model systematically underpredicts peak and maximum rate ($\approx 0.5\times$ measured) for both real and synthetic patients, reflecting overcorrection of the protein C anticoagulant pathway. Cloud overlap was 88–92%, showing similar model behavior for real and synthetic profiles despite imperfect calibration.




Figure S6: **Cross-visit calibration generalization.** Boxplots of the ODE-predicted/dataset ratio by visit for real patients. The model was calibrated on Visit 1 only. Top: under TF-only conditions, ratios remained near 1.0 at Visits 2 and 3. Bottom: under TF+TM conditions, the ratios differed more across visits, particularly for peak and maximum rate.




Figure S7: **Spearman rank correlations between ODE-predicted and dataset TGA features.** Left: TF-only. Right: TF+TM. ETP is the best-predicted feature ($\rho \approx 0.6$ – 0.8 for real patients under TF-only). Weak rank correlations for other features are expected given the 5-parameter population-level calibration.




Figure S8: **PCA projections by visit: SA and MVN.** Each panel shows the first two principal components (PC1: 27.7% variance, PC2: 15.4%) computed from standardized per-visit real patient data ($K=23$ patients per visit, 72 features). Dark markers indicate real patients; lighter markers indicate synthetic patients ($N=100$). Within each visit (column), the SA panel (top) and MVN panel (bottom) share identical axis ranges, so the relative dispersion of the two methods can be read against the same real data cloud. Top row: SA-generated patients (circles) cluster tightly around the real data cloud at all three visits. Bottom row: MVN-generated patients (diamonds) show greater dispersion, particularly at Visits 2 (first trimester) and 3 (third trimester), where MVN patients extend beyond the convex hull of the real data. This excess scatter reflects the spurious variance introduced by Ledoit–Wolf regularization of the rank-deficient covariance (Fig. S2). Colors encode visit: blue (V1, not pregnant), green (V2, first trimester), orange (V3, third trimester).

Sampling correctness and privacy diagnostics

We asked whether a synthetic cohort could reveal that an already-known de-identified assay profile was stored in the memory matrix. The source data had no direct identifiers or links to named individuals, so this was not a test of personal identification. In standardized 216-dimensional space, the median distance to the closest real record (DCR) was 13.8 for synthetic profiles, compared with a leave-one-out real-to-real distance of 15.9. We then tested 40 random splits with 15 stored and 8 held-out profiles. Distance to the synthetic cohort distinguished the groups with mean AUC 0.97 (standard deviation 0.02; 0.5 meant no discrimination and 1.0 perfect discrimination). To test whether the numerical sampler caused this signal, we compared the reported sampler with sphere initialization, pooled samples after burn-in and thinning, and Metropolis-adjusted sampling (Table S16, Fig. S9). Metropolis acceptance was 0.999, and the AUC remained 0.96–0.97. Median DCR changed only from 14.0 to 14.3 and remained below 15.9. Changing the Markov-chain procedure therefore did not remove the membership signal.

We also varied the inverse temperature and sampled the target distribution directly. Across a 40-fold sweep around β^* , lower β increased median DCR but it plateaued near 14.3; cross-visit covariance error increased from 0.59 to 0.69 (Fig. S9C). The closed-form stochastic-attention target in Eq. (3) is the exact mixture $\sum_k q_k \mathcal{N}(\mathbf{m}_k, \beta^{-1} \mathbf{I})$, with $q_k = r_k / \sum_j r_j$ for unit-normalized memories. Its direct sampler requires only a component draw and isotropic Gaussian noise. It has no initialization, burn-in, or discretization error. A matched direct-sampler cohort gave median relative error 1.68%, cross-visit error 0.60, mechanistic overlap 0.91, median novelty 0.50, and median DCR 13.87, all within the range of the four Markov-chain procedures. Across 100 direct-sampler cohorts, median relative error was 1.40% (5th–95th percentile 0.97–1.91%) and median DCR was 13.95 (13.61–14.39). Exact sampling gave AUC 0.98 over the same 40 splits. This showed that the membership signal came from the target distribution, not from the numerical sampler. At this cohort size, the test could therefore distinguish whether an already-known de-identified profile was stored regardless of how the target distribution was sampled.

Table S16: **Sampler comparison.** All procedures used the same β^* , memory, and decoder. We report median relative error (MRE), cross-visit correlation error, mechanistic overlap, novelty, DCR, and membership-inference AUC where evaluated. The real-to-real DCR baseline was 15.9. The AUC used the same 40 splits for V0, V3, and the direct stochastic-attention sampler. *Exact* drew directly from Eq. (3) and was not a Markov-chain procedure. Across 100 direct-sampler cohorts, median MRE was 1.40% (5th–95th percentile 0.97–1.91%) and median DCR was 13.95 (13.61–14.39).

Rung	Sampler	MRE (%)	Cross-visit Frob.	Mech. overlap	Novelty	DCR	MIA AUC
V0	anchor, endpoint	1.31	0.65	0.89	0.51	13.97	0.97
V1	sphere init	1.53	0.57	0.89	0.50	13.96	n/a
V2	pooled ULA	2.02	0.59	0.88	0.50	14.15	n/a
V3	pooled MALA	1.89	0.62	0.92	0.52	14.33	0.96
Exact	analytic ancestral	1.68	0.60	0.91	0.50	13.87	0.98




Figure S9: **Sampling correctness and the membership signal.** (A) Across V0–V3, DCR rose slightly but stayed below the real-to-real baseline, while membership-inference AUC remained 0.96–0.97. (B) Median relative error across the four procedures. (C) Across a 40-fold inverse-temperature sweep around β^* , DCR increased as β fell but remained below the real-to-real baseline. Cross-visit covariance error increased, while mechanistic overlap changed little.




Figure S10: **Mixing diagnostic at β^* .** Energy traces, the between-chain $\hat{R}$ statistic, and the anchor-distance trajectory for sphere-initialized Langevin chains. The target distribution at β^* was multimodal, with one component centered on each stored pattern. Metropolis acceptance near 0.999 indicated negligible discretization bias.