

# Supplementary Information

## Supplementary Tables

Table S1: **Inverse temperature sensitivity analysis.** SA generation quality as a function of  $\beta/\beta^*$ , where  $\beta^* \approx 2.94$  was identified via entropy inflection in the 18-dimensional PCA space ( $K=23$  stored patterns). Operating at  $\beta = \beta^*$  (bold row) balances generation novelty against fidelity. Med MRE = median relative error of per-feature means; Novelty =  $1 - \max_k \cos(\hat{\xi}, \hat{m}_k)$ .

| $\beta/\beta^*$ | $\beta$     | Noise scale  | PC1 std ratio | Novelty      | Med MRE (%) |
|-----------------|-------------|--------------|---------------|--------------|-------------|
| 0.1             | 0.29        | 0.261        | 0.538         | 0.542        | 0.7         |
| 0.3             | 0.88        | 0.151        | 0.550         | 0.536        | 0.7         |
| 0.5             | 1.47        | 0.117        | 0.562         | 0.528        | 0.7         |
| <b>1.0</b>      | <b>2.94</b> | <b>0.082</b> | <b>0.589</b>  | <b>0.502</b> | <b>0.6</b>  |
| 1.5             | 4.41        | 0.067        | 0.613         | 0.474        | 0.7         |
| 2.0             | 5.88        | 0.058        | 0.638         | 0.446        | 0.7         |
| 3.0             | 8.82        | 0.048        | 0.689         | 0.403        | 0.8         |

Table S2: **Multiplicity weighting parameters for conditional generation.** The multiplicity ratio  $\rho$  was computed to achieve a target effective fraction  $f_{\text{target}} = 0.80$  of softmax attention on the designated patterns.  $K_{\text{eff}}$  is the participation ratio quantifying the effective number of patterns contributing to generation.

| Condition     | $K_{\text{sub}}$ | $K_{\text{bg}}$ | $\rho$ | $f_{\text{eff}}$ | $K_{\text{eff}}$ |
|---------------|------------------|-----------------|--------|------------------|------------------|
| Uncomplicated | 18               | 5               | 1.11   | 0.80             | 23.0             |
| PCOS          | 3                | 20              | 26.67  | 0.80             | 4.6              |
| Developed PE  | 5                | 18              | 14.40  | 0.80             | 7.7              |

Table S3: **SA generation hyperparameters.** The same values were used for unconditioned and conditioned generation except where noted.

| Parameter           | Value | Description                                  |
|---------------------|-------|----------------------------------------------|
| $\alpha$            | 0.01  | Langevin step size (see Table S8)            |
| $T$                 | 2,000 | Langevin iterations per sample               |
| $\beta^*$           | 2.94  | Inverse temperature (entropy inflection)     |
| PCA threshold       | 95%   | Variance retained                            |
| $d_{\text{PCA}}$    | 18    | PCA dimensionality                           |
| $N$                 | 100   | Synthetic patients per cohort                |
| Random seed         | 42    | For reproducibility                          |
| $f_{\text{target}}$ | 0.80  | Target attention fraction (conditioned only) |

Table S4: **BZ2012 ODE model calibration.** Five rate constants were calibrated on Visit 1 real patients ( $n=23$ ) using bounded Nelder–Mead optimization in log-scale-factor space. The cost function was the mean normalized relative squared error across five TGA features under TF-only conditions. Bounds:  $\pm \log(100)$ . Convergence:  $f_{\text{tol}} = 10^{-6}$ , 500 iterations, 11 random restarts.

| Rate constant                   | Role                   | Default               | Calibrated             | Scale         |
|---------------------------------|------------------------|-----------------------|------------------------|---------------|
| prothrombinase $k_{\text{cat}}$ | Thrombin burst         | $63.5 \text{ s}^{-1}$ | $21.94 \text{ s}^{-1}$ | $0.35\times$  |
| intrinsic Xase $k_{\text{cat}}$ | Amplification Xa       | $8.2 \text{ s}^{-1}$  | $0.169 \text{ s}^{-1}$ | $0.021\times$ |
| extrinsic Xase $k_{\text{cat}}$ | Initiation Xa          | $6.0 \text{ s}^{-1}$  | $93.2 \text{ s}^{-1}$  | $15.5\times$  |
| PC activation $k_{\text{cat}}$  | Protein C activation   | $0.41 \text{ s}^{-1}$ | $0.890 \text{ s}^{-1}$ | $2.17\times$  |
| mIIa conversion $k$             | mIIa $\rightarrow$ IIa | $2.3\times 10^8$      | $1.15\times 10^9$      | $5.01\times$  |

Units for mIIa conversion:  $\text{M}^{-1}\text{s}^{-1}$

Table S5: **Summary of mean relative errors across all 72 features.** Per-visit distribution of MREs between real and SA-generated synthetic patient means. The full feature-by-feature table (72 features  $\times$  3 visits = 216 entries) is provided in the code repository as `paper_summary_statistics.csv`.

| Visit                        | Median MRE   | Mean MRE     | 25th pctl    | 75th pctl    | Max MRE      | Below 5%               |
|------------------------------|--------------|--------------|--------------|--------------|--------------|------------------------|
| V1 (baseline)                | 0.022        | 0.029        | 0.008        | 0.039        | 0.158        | 59/72 (81.9%)          |
| V2 (1st tri)                 | 0.009        | 0.021        | 0.004        | 0.022        | 0.224        | 65/72 (90.3%)          |
| V3 (3rd tri)                 | 0.011        | 0.016        | 0.005        | 0.024        | 0.062        | 69/72 (95.8%)          |
| <b>Pooled (all 3 visits)</b> | <b>0.012</b> | <b>0.022</b> | <b>0.005</b> | <b>0.028</b> | <b>0.224</b> | <b>193/216 (89.4%)</b> |

Pooled median MRE 95% bootstrap CI: (0.98%, 1.60%) over 10,000 resamples (seed 42).

Table S6: **Comparison with deep generative baselines.** CTGAN and TVAE were trained on the same 90 per-visit records ( $23 \text{ patients} \times 3 \text{ visits}$ , 72 features) at three epoch counts. SA and MVN results from the main text are shown for reference. CTGAN fails at this sample size (median MRE  $\approx 19\%$ ), consistent with GAN mode collapse on small datasets. TVAE achieves comparable marginal fidelity at high epoch counts but cannot capture cross-visit covariance or perform conditional generation.

| Method | Epochs | Med MRE | Med KS | Corr MAE |
|--------|--------|---------|--------|----------|
| SA     | –      | 0.019   | 0.169  | 0.200    |
| MVN    | –      | 0.024   | 0.161  | 0.090    |
| CTGAN  | 300    | 0.189   | 0.364  | 0.261    |
| CTGAN  | 1,000  | 0.195   | 0.349  | 0.265    |
| CTGAN  | 3,000  | 0.187   | 0.341  | 0.180    |
| TVAE   | 300    | 0.037   | 0.174  | 0.152    |
| TVAE   | 1,000  | 0.026   | 0.248  | 0.082    |
| TVAE   | 3,000  | 0.018   | 0.224  | 0.092    |

Table S7: **PCA variance threshold sensitivity analysis.** SA generation quality is stable across thresholds from 85% to 99%. Median MRE remains between 1.0% and 1.7% across all thresholds. The 95% threshold used in the main text (bold) is not a critical design choice.

| Threshold  | $d_{\text{PCA}}$ | $K/d$       | $\beta^*$   | Novelty      | Med MRE       | Corr MAE     |
|------------|------------------|-------------|-------------|--------------|---------------|--------------|
| 85%        | 13               | 1.77        | 2.94        | 0.420        | 0.0098        | 0.124        |
| 90%        | 15               | 1.53        | 2.94        | 0.477        | 0.0160        | 0.126        |
| <b>95%</b> | <b>18</b>        | <b>1.28</b> | <b>2.94</b> | <b>0.500</b> | <b>0.0124</b> | <b>0.143</b> |
| 97.5%      | 20               | 1.15        | 2.94        | 0.543        | 0.0120        | 0.135        |
| 99%        | 21               | 1.10        | 2.94        | 0.541        | 0.0172        | 0.138        |

Table S8: **ULA vs MALA sampling diagnostics.** Ten chains of 5,000 iterations each were run with both the Unadjusted Langevin Algorithm (ULA) and the Metropolis-Adjusted Langevin Algorithm (MALA) at step size  $\alpha = 0.01$  and  $\beta^* = 2.94$  on the 18-dimensional PCA memory ( $K=23$  patterns). The MALA acceptance rate of 99.9% confirms that virtually every ULA proposal satisfies detailed balance, and the indistinguishable energy and effective sample size statistics confirm negligible discretization bias. [This comparison bounds the discretization error of the reported sampler.](#) [The direct reference for the target itself is the analytic sampler of Eq. \(3\), introduced below.](#)

| Metric                     | ULA               | MALA              |
|----------------------------|-------------------|-------------------|
| Acceptance rate (%)        | – (no rejection)  | $99.9 \pm 0.03$   |
| Mean energy (post-burn-in) | $1.932 \pm 0.141$ | $1.886 \pm 0.231$ |
| $\tau_{\text{int}}$        | $109.5 \pm 49.5$  | $106.7 \pm 30.4$  |
| Effective sample size      | $41.8 \pm 14.2$   | $40.0 \pm 10.2$   |

Burn-in: 1,000 iterations discarded.  $\tau_{\text{int}}$ : integrated autocorrelation time estimated with a windowed estimator (threshold 0.05).  $\text{ESS} = (T - T_{\text{burn}})/\tau_{\text{int}}$ .

Table S9: **Memorization diagnostics.** Two complementary checks that SA-generated patients are not memorized copies of the  $K=23$  stored profiles. Novelty score is the angular distance between a synthetic sample and its nearest stored pattern at  $\beta = \beta^*$ . Nearest-neighbor distances are computed in the per-visit standardized 72-dimensional feature space.

| Metric                                                      | Value |
|-------------------------------------------------------------|-------|
| Mean novelty score, $1 - \max_k \cos(\hat{\xi}, \hat{m}_k)$ | 0.50  |
| Median synth–real nearest-neighbor distance (std. units)    | 6.14  |
| Median real–real nearest-neighbor distance (std. units)     | 6.87  |
| Distance ratio (synth–real / real–real)                     | 0.89  |

Table S10: **Bootstrap Mann–Whitney equivalence per feature and condition.** For each of 24 feature–condition pairs, we subsampled the synthetic cohort to match the real sample size, applied a Mann–Whitney U test, and repeated 1,000 times. Frac.  $p > 0.05$  is the fraction of replicates in which the test could not distinguish the synthetic from the real distribution. 20 of 24 pairs achieve  $p > 0.05$  in  $\geq 90\%$  of replicates; the four failing pairs (italicized) all involve high-variance features in the smallest subgroups.

| Condition                | Feature        | MRE          | Frac. $p > 0.05$ |
|--------------------------|----------------|--------------|------------------|
| Uncomplicated ( $n=18$ ) | FVIII          | 0.057        | 0.952            |
| Uncomplicated ( $n=18$ ) | vWF            | 0.022        | 0.999            |
| Uncomplicated ( $n=18$ ) | FIX            | 0.018        | 0.998            |
| Uncomplicated ( $n=18$ ) | FV             | 0.006        | 0.996            |
| Uncomplicated ( $n=18$ ) | $\alpha_2$ -AP | 0.019        | 1.000            |
| Uncomplicated ( $n=18$ ) | FVII           | 0.057        | 0.941            |
| Uncomplicated ( $n=18$ ) | Fbgn           | 0.017        | 0.987            |
| Uncomplicated ( $n=18$ ) | Plgn           | 0.020        | 0.988            |
| PCOS ( $n=3$ )           | FVIII          | 0.109        | 0.985            |
| PCOS ( $n=3$ )           | vWF            | 0.150        | 0.959            |
| PCOS ( $n=3$ )           | <i>FIX</i>     | <i>0.121</i> | <i>0.873</i>     |
| PCOS ( $n=3$ )           | FV             | 0.057        | 0.944            |
| PCOS ( $n=3$ )           | $\alpha_2$ -AP | 0.051        | 0.987            |
| PCOS ( $n=3$ )           | FVII           | 0.029        | 0.999            |
| PCOS ( $n=3$ )           | Fbgn           | 0.043        | 0.992            |
| PCOS ( $n=3$ )           | <i>Plgn</i>    | <i>0.153</i> | <i>0.615</i>     |
| Developed PE ( $n=5$ )   | FVIII          | 0.097        | 0.932            |
| Developed PE ( $n=5$ )   | vWF            | 0.010        | 0.986            |
| Developed PE ( $n=5$ )   | <i>FIX</i>     | <i>0.108</i> | <i>0.780</i>     |
| Developed PE ( $n=5$ )   | FV             | 0.006        | 0.987            |
| Developed PE ( $n=5$ )   | $\alpha_2$ -AP | 0.062        | 0.950            |
| Developed PE ( $n=5$ )   | FVII           | 0.045        | 0.993            |
| Developed PE ( $n=5$ )   | Fbgn           | 0.048        | 0.996            |
| Developed PE ( $n=5$ )   | <i>Plgn</i>    | <i>0.093</i> | <i>0.814</i>     |

Median Frac.  $p > 0.05$  across all 24 pairs: 0.986. Failing pairs ( $< 0.90$ , italicized): PCOS–FIX, PCOS–Plgn, DPE–FIX, DPE–Plgn.

Table S11: **Mechanistic validation diagnostics per feature and condition.** Cloud overlap is the fraction of synthetic ODE-predicted/measured ratios falling within the 5th–95th percentile of the corresponding real distribution. KS  $D$  and KS  $p$  are from a two-sample Kolmogorov–Smirnov test on the same ratio distributions. Real  $n=69$  (23 patients  $\times$  3 visits); synthetic  $n=300$  (100 patients  $\times$  3 visits).

| Condition | TGA Feature | Cloud overlap | KS $D$ | KS $p$ |
|-----------|-------------|---------------|--------|--------|
| TF-only   | Lagtime     | 0.92          | 0.112  | 0.48   |
| TF-only   | Peak        | 0.91          | 0.081  | 0.86   |
| TF-only   | T.Peak      | 0.92          | 0.102  | 0.61   |
| TF-only   | Max Rate    | 0.91          | 0.083  | 0.84   |
| TF-only   | ETP         | 0.83          | 0.123  | 0.36   |
| TF+TM     | Lagtime     | 0.92          | 0.127  | 0.32   |
| TF+TM     | Peak        | 0.88          | 0.079  | 0.88   |
| TF+TM     | T.Peak      | 0.91          | 0.110  | 0.51   |
| TF+TM     | Max Rate    | 0.88          | 0.096  | 0.67   |
| TF+TM     | ETP         | 0.89          | 0.112  | 0.48   |

TF-only cloud overlap range: 0.83–0.92. TF+TM cloud overlap range: 0.88–0.92. Pooled across both conditions and all five features, 0.902. KS  $D$  range across both conditions: 0.079–0.127, all  $p > 0.30$ .

Table S12: **Downstream utility — per-feature held-out V2/V3 performance.** Median relative error of the BZ2012 ODE-predicted TGA values against the held-out real V2 and V3 measurements, for the model calibrated on real V1 patients ( $K=23$ ) and the model calibrated on synthetic V1 patients ( $N=100$ ). The synth-calibrated model achieves slightly lower error on every feature. Synth/Real ratio  $< 1$  favors the synth-calibrated model.

| TGA Feature           | Real-cal MRE | Synth-cal MRE | Synth/Real                     |
|-----------------------|--------------|---------------|--------------------------------|
| Lagtime               | 0.165        | 0.162         | 0.98 $\times$                  |
| Peak                  | 0.097        | 0.087         | 0.90 $\times$                  |
| T.Peak                | 0.330        | 0.306         | 0.93 $\times$                  |
| Max Rate              | 0.286        | 0.259         | 0.91 $\times$                  |
| ETP                   | 0.197        | 0.183         | 0.93 $\times$                  |
| <b>Overall median</b> | <b>0.201</b> | <b>0.189</b>  | <b>0.94<math>\times</math></b> |

Table S13: **PCA-space ablation.** Five generators fit in the identical  $d=18$  PCA subspace and evaluated on the same harness as SA. Median relative error (MRE) and the mechanistic cloud overlap measure fidelity. Median novelty measures distance from the stored patterns, and Frac. novel is the fraction above the 0.2 novelty threshold. The cross-visit Frobenius error is the deviation of the synthetic correlation matrix from the real one. A lower value means staying closer to the 23 real patients, so this metric rewards memorization. The shared subspace is necessary but not sufficient.  $k$ -nearest-neighbor interpolation collapses novelty, and the kernel-density estimator degrades marginal and mechanistic fidelity, despite both being fit in it. The mixture and copula baselines match SA on the fidelity axes.

| Method     | MRE (%) | Cross-visit Frob. | Mech. overlap | Median novelty | Frac. novel |
|------------|---------|-------------------|---------------|----------------|-------------|
| SA         | 1.24    | 0.641             | 0.902         | 0.514          | 1.00        |
| PCA+GMM    | 1.13    | 0.422             | 0.918         | 0.497          | 0.94        |
| PCA+KDE    | 1.94    | 0.337             | 0.823         | 0.249          | 0.65        |
| PCA+kNN    | 1.63    | 0.587             | 0.938         | 0.051          | 0.05        |
| PCA+Copula | 1.34    | 0.417             | 0.934         | 0.462          | 1.00        |

Table S14: **Convex-hull membership does not discriminate at  $K=23$  in  $d=18$ .** Hull membership and distance for the 100 synthetic patients against the 23 real patients, and for each real patient against the other 22. The radial factor  $t^*$  is where the ray from the hull centroid through a point exits the hull, so  $1/t^*$  is the factor by which a point overshoots the boundary and  $t^* \geq 1$  would mean the point is inside. Distances and overshoots are severity-matched, with each point's nearest vertex removed on both sides, because a held-out real patient necessarily loses its own vertex. Matched this way, the distance to the hull is not distinguishable between synthetic and real patients (Mann–Whitney  $p = 0.07$ ), while the radial overshoot is smaller for the synthetic ones ( $p < 10^{-4}$ ). Both tests treat the 100 synthetic patients as independent, which they are not, so the reported significance is optimistic. The geometry rows show why no point of either kind falls inside. The hull attains the patients' typical radius only along the 23 vertex directions, and the sampler returns every sample to the unit sphere. A point inside the hull of unit-norm patterns would have to be one of those patterns.

| Quantity                                                  | Synthetic ( $n=100$ ) | Real, leave-one-out ( $n=23$ ) |
|-----------------------------------------------------------|-----------------------|--------------------------------|
| Fraction inside hull                                      | 0.00                  | 0.00                           |
| Median distance to hull                                   | 11.0                  | 11.9                           |
| Median radial overshoot $1/t^*$                           | 10.5                  | 15.2                           |
| <i>Hull geometry of the real cohort</i>                   |                       |                                |
| Median patient radius                                     |                       | 13.8                           |
| Median hull extent, generic direction                     |                       | 2.0                            |
| Largest $t^*$ attained on the unit sphere                 |                       | 0.231                          |
| <i>Plausibility of the out-of-hull synthetic patients</i> |                       |                                |
| Biological constraints satisfied                          |                       | 81/100                         |
| Mechanistic band, per feature                             |                       | 83 to 92%                      |
| Mechanistic band, all features and visits                 |                       | 41/100                         |

Table S15: **Membership inference recovers memory membership, not unseen patients.** Median nearest-neighbor distances in the standardized 216-dimensional space, and the attack’s discrimination. The attack retrains the generator on 15 training patients and scores each real record by its distance to the resulting synthetic set, treating the 15 training patients as members and the 8 held-out patients as non-members. Every quantity comes from the same protocol. We drew 40 partitions at random with the generation seed held fixed, so that only the partition varies, and pooled the distances across all partitions. Non-members sit at the same distance from the synthetic cohort as unrelated real patients sit from one another. The synthetic cohort carries no proximity signal about patients the generator never saw, and the discrimination comes entirely from members sitting closer.

| Quantity                                              | Value |
|-------------------------------------------------------|-------|
| <i>Median nearest-neighbor distance</i>               |       |
| Training member to nearest synthetic                  | 10.7  |
| Held-out non-member to nearest synthetic              | 15.4  |
| Real to nearest real, leave-one-out                   | 15.9  |
| Synthetic to nearest real                             | 13.8  |
| <i>Membership-inference AUC over 40 random splits</i> |       |
| Mean                                                  | 0.97  |
| Standard deviation                                    | 0.02  |
| Median                                                | 0.98  |
| Minimum                                               | 0.90  |
| Maximum                                               | 1.00  |
| Splits reaching perfect separation                    | 9/40  |
| Splits with AUC above 0.9                             | 39/40 |

## Supplementary Figures




Figure S1: **Cross-visit correlation structure (amplified residual scale)**. Same layout as the main-text correlation figure, but with the bottom-row residual panels plotted on a  $\pm 0.5$  color scale (vs.  $\pm 1.0$  in the main text) to reveal fine structure. Frobenius norms of each residual matrix are shown in the panel titles.

## Supplementary Methods: locating the critical inverse temperature

The operating point  $\beta^*$  was located from the attention entropy. For a candidate  $\beta$  we computed the Shannon entropy of the attention weights  $\text{softmax}(\beta \hat{\mathbf{X}}^\top \boldsymbol{\xi})$  at 20 probe points placed at stored patterns, averaged those entropies, and normalized by  $\log K$ . The normalized curve runs from 0 to 1. A value of 0 means one pattern carries all the weight, and a value of 1 means every pattern carries equal weight. Sweeping  $\beta$  logarithmically over 80 values from 0.1 to 1000 traces this curve from the diffuse regime into the concentrated one. We took  $\beta^*$  to be the point of most negative second derivative of the normalized entropy with respect to  $\log \beta$ , computed by finite differences over the sweep. That point is the sharpest bend in the curve, and it marks where attention stops being shared across patterns and begins to localize on one. For multiplicity-weighted sampling we repeated the same computation with the  $\log r_k$  bias included in the softmax logits. This gives a  $\rho$ -dependent  $\beta^*(\rho)$ , because the weighting changes how quickly attention concentrates as  $\beta$  grows. For random unit-norm patterns the transition is predicted at  $\beta^* = \sqrt{d}$ , the theoretical reference quoted in the main text.






Figure S2: **Covariance eigenvalue spectrum.** Log-scale eigenvalues of the  $216 \times 216$  sample covariance for Real ( $K=23$ , rank 22), SA ( $N=100$ , rank 18), and MVN ( $N=100$ ) populations. Dashed vertical lines mark the SA PCA truncation point (component 18) and the data rank limit (component 22). MVN’s Ledoit–Wolf regularization inflates eigenvalues beyond component 22 (green shaded region), introducing spurious variance in 194 dimensions where the data contain no signal.




Figure S3: **Conditioned generation: PCA projection (pooled).** Large markers: real patients (blue = Uncomplicated,  $n=18$ ; orange = PCOS,  $n=3$ ; red = Developed PE,  $n=5$ ). Small markers: synthetic patients ( $N=100$  per condition). With only 3 PCOS and 5 Developed PE patients, the subgroups do not form clearly separable clusters in the first two principal components, but the synthetic cohorts concentrate around their respective real counterparts.




Figure S4: **Conditioned generation: PCA projection by visit.** Same as Fig. S3 but separated by visit (V1, V2, V3). Condition-specific clustering is maintained within each individual visit.




Figure S5: **Mechanistic validation (TF+TM condition)**. Same layout as main-text Fig. 6 but under TF+TM conditions. The model systematically underpredicts peak and maximum rate ( $\approx 0.5\times$  measured) for both real and synthetic patients, reflecting overcorrection of the protein C anticoagulant pathway. Cloud overlap remains high at 88–92%, confirming that the model processes real and synthetic patients identically even under imperfect calibration conditions.




Figure S6: **Cross-visit calibration generalization**. Boxplots of the ODE-predicted/dataset ratio by visit for real patients. The model was calibrated on Visit 1 only. Top: TF-only conditions, where ratios near 1.0 at Visits 2 and 3 indicate that the calibration generalizes across pregnancy timepoints. Bottom: TF+TM conditions, where the calibration does not generalize as well, particularly for peak and maximum rate.




Figure S7: **Spearman rank correlations between ODE-predicted and dataset TGA features.** Left: TF-only. Right: TF+TM. ETP is the best-predicted feature ( $\rho \approx 0.6$ – $0.8$  for real patients under TF-only). Weak rank correlations for other features are expected given the 5-parameter population-level calibration.




Figure S8: **PCA projections by visit: SA vs. MVN.** Each panel shows the first two principal components (PC1: 27.7% variance, PC2: 15.4%) computed from standardized per-visit real patient data ( $K=23$  patients per visit, 72 features). Dark markers indicate real patients; lighter markers indicate synthetic patients ( $N=100$ ). Within each visit (column), the SA panel (top) and MVN panel (bottom) share identical axis ranges, so the relative dispersion of the two methods can be read against the same real data cloud. Top row: SA-generated patients (circles) cluster tightly around the real data cloud at all three visits. Bottom row: MVN-generated patients (diamonds) show greater dispersion, particularly at Visits 2 (first trimester) and 3 (third trimester), where MVN patients extend beyond the convex hull of the real data. This excess scatter reflects the spurious variance introduced by Ledoit–Wolf regularization of the rank-deficient covariance (Fig. S2). Colors encode visit: blue (V1, not pregnant), green (V2, first trimester), orange (V3, third trimester).

## Sampling correctness and privacy diagnostics

We assessed re-identification risk directly. For every synthetic patient we computed the distance to the closest real record (DCR) in the standardized 216-dimensional space and compared the synthetic-to-real distribution against the leave-one-out real-to-real distribution. Synthetic patients sit slightly closer to the real cohort than the real patients sit to one another, with median DCR 13.8 against 15.9. We also ran a nearest-neighbor membership-inference attack, which scores each real record by its proximity to the synthetic set over repeated random train/holdout splits. It separated the 15 training patients from the 8 held-out patients with a mean AUC of 0.97 (standard deviation 0.02 over 40 splits). At the published operating point the generator therefore reproduces its training cohort closely enough to be identifiable.

To attribute this behavior we asked whether it reflects the target distribution  $\pi(\xi) \propto \exp(-\beta E)$  or merely the sampler used to draw from it. A mixing diagnostic showed that at  $\beta^*$  the target is multimodal, with one component centered on each of the  $K=23$  stored patterns. It also showed that discretization bias in the Langevin updates is negligible, with Metropolis acceptance near 0.999 (Supplementary Fig. S10). We then constructed a four-rung ladder (Table S16, Fig. S9) that holds  $\beta$ , the memory, and the decoding fixed and varies only the sampler. The rungs run from the published anchor-initialized single-endpoint scheme, to sphere initialization, to burn-in-controlled thinned pooling, to Metropolis-adjusted sampling. Across the ladder the membership-inference AUC stays near 0.96 to 0.97. The median DCR rises only from 14.0 to 14.3 and never reaches the real-to-real baseline of 15.9, at a modest cost in marginal fidelity. Proper sampling of  $\pi$  does not remove the memorization. At  $K=23$  the memorization is a property of the distribution.

The remaining lever is the inverse temperature. Lowering  $\beta$  spreads probability mass away from the stored patterns and should improve privacy, at the expense of the memorization that also underlies fidelity. We swept  $\beta$  across a 40-fold range around  $\beta^*$  (Fig. S9C). The sweep confirmed the trade-off and bounded it sharply. The median DCR rises monotonically as  $\beta$  falls, but it saturates near 14.3 and stays below the real-to-real baseline of 15.9 even at one twentieth of  $\beta^*$ . Over the same range the cross-visit covariance error grows from 0.59 to 0.69. Marginal and mechanistic fidelity are insensitive to  $\beta$ , so the longitudinal covariance structure bears the cost of lowering it. No inverse temperature yields a cohort that is both as private as the real patients are from one another and faithful to their joint structure.

Neither lever settles the question completely, because every scheme compared so far is a finite-time Markov chain. The target itself is available in closed form. Completing the square in Eq. (1) gives Eq. (3). For unit-normalized memories  $\pi_r$  is exactly the isotropic Gaussian mixture  $\sum_k q_k \mathcal{N}(\mathbf{m}_k, \beta^{-1}\mathbf{I})$  with  $q_k = r_k / \sum_j r_j$ . It can therefore be sampled ancestrally, by drawing a component and adding isotropic noise. There is no chain, no initialization, no burn-in and no discretization error. We generated a protocol-matched exact cohort ( $N=100$  at  $\beta^*$ , reusing the ladder’s seeded magnitude draw so that only the direction sampler differs) and evaluated it on the identical harness (Table S16). Every metric falls inside the range spanned by the four rungs: median relative error 1.68%, cross-visit Frobenius error 0.60, mechanistic overlap 0.91, median novelty 0.50, and median DCR 13.87. One hundred independently seeded exact cohorts place these quantities in narrow intervals, with median relative error 1.40% (fifth to ninety-fifth percentile 0.97 to 1.91%) and median DCR 13.95 (13.61 to 14.39). The agreement does not rest on a favorable seed. Over the same 40 train/holdout splits used for the ladder endpoints, the membership-inference AUC under exact sampling is 0.98, marginally above the published scheme’s 0.97. The memorization is therefore a property of  $\pi_r$  at  $K=23$ , where every stored patient is a component center. It is not a property of any approximation used to sample it.

Table S16: **Sampler-correctness comparison.** Each rung holds  $\beta^*$ , the memory, and the decoding fixed and varies only the direction sampler. Fidelity metrics (median relative error, cross-visit correlation Frobenius error, mechanistic cloud overlap, median novelty) and privacy metrics (median distance to the closest real record, and repeated-split membership-inference AUC where evaluated) are reported per rung. The real-to-real DCR baseline is 15.9. The membership-inference attack was run on the two endpoints of the ladder and on the analytic reference, over the same 40 train/holdout splits. *Exact* is not a rung. It draws ancestrally from Eq. (3) with no chain, and serves as the analytic reference for the four sampled rungs. Over 100 independently seeded exact cohorts the median relative error is 1.40% (fifth to ninety-fifth percentile 0.97 to 1.91%) and the median DCR is 13.95 (13.61 to 14.39).

| Rung  | Sampler            | MRE (%) | Cross-visit Frob. | Mech. overlap | Novelty | DCR   | MIA AUC |
|-------|--------------------|---------|-------------------|---------------|---------|-------|---------|
| V0    | anchor, endpoint   | 1.31    | 0.65              | 0.89          | 0.51    | 13.97 | 0.97    |
| V1    | sphere init        | 1.53    | 0.57              | 0.89          | 0.50    | 13.96 | n/a     |
| V2    | pooled ULA         | 2.02    | 0.59              | 0.88          | 0.50    | 14.15 | n/a     |
| V3    | pooled MALA        | 1.89    | 0.62              | 0.92          | 0.52    | 14.33 | 0.96    |
| Exact | analytic ancestral | 1.68    | 0.60              | 0.91          | 0.50    | 13.87 | 0.98    |




Figure S9: **Sampling correctness and the memorization/privacy relationship.** (A) Across the V0 to V3 ladder the distance to the closest real record (DCR, left axis) rises only slightly and stays below the real-to-real baseline, while the membership-inference AUC (right axis) remains near 0.96 to 0.97. (B) The fidelity cost of multi-chain pooling, as median relative error. (C) The inverse temperature swept across a 40-fold range around the operating point  $\beta^*$ . The DCR (privacy, left axis) rises as  $\beta$  falls but never reaches the real-to-real baseline. The cross-visit covariance error (right axis) grows, and the mechanistic overlap stays flat.




Figure S10: **Mixing diagnostic at  $\beta^*$ .** Energy traces, the between-chain  $\hat{R}$  statistic, and the anchor-distance trajectory for sphere-initialized Langevin chains. The target distribution at  $\beta^*$  is multimodal, with one component centered on each stored pattern. Metropolis acceptance near 0.999 indicates negligible discretization bias.