Supplementary Information

Supplementary Tables

Table S1: **Inverse temperature sensitivity analysis.** SA generation quality as a function of β/β^*, where $\beta^* \approx 2.94$ was identified via entropy inflection in the 18-dimensional PCA space (K=23 stored patterns). Operating at $\beta = \beta^*$ (bold row) balances generation novelty against fidelity. Med MRE = median relative error of per-feature means; Novelty = $1 - \max_k \cos(\hat{\xi}, \hat{m}_k)$.

β/β^*	β	Noise scale	PC1 std ratio	Novelty	Med MRE (%)
0.1	0.29	0.261	0.538	0.542	0.7
0.3	0.88	0.151	0.550	0.536	0.7
0.5	1.47	0.117	0.562	0.528	0.7
1.0	**2.94**	**0.082**	**0.589**	**0.502**	**0.6**
1.5	4.41	0.067	0.613	0.474	0.7
2.0	5.88	0.058	0.638	0.446	0.7
3.0	8.82	0.048	0.689	0.403	0.8

Table S2: **Multiplicity weighting parameters for conditional generation.** The multiplicity ratio ρ was computed to achieve a target effective fraction $f_{\text{target}} = 0.80$ of softmax attention on the designated patterns. K_{eff} is the participation ratio quantifying the effective number of patterns contributing to generation.

Condition	K_{sub}	K_{bg}	ρ	f_{eff}	K_{eff}
Uncomplicated	18	5	1.11	0.80	23.0
PCOS	3	20	26.67	0.80	4.6
Developed PE	5	18	14.40	0.80	7.7

Table S3: **SA generation hyperparameters.** The same values were used for unconditioned and conditioned generation except where noted.

Parameter	Value	Description
α	0.01	Langevin step size (ULA valid without MH correction)
T	2,000	Langevin iterations per sample
β^*	2.94	Inverse temperature (entropy inflection)
PCA threshold	95%	Variance retained
d_{PCA}	18	PCA dimensionality
N	100	Synthetic patients per cohort
Random seed	42	For reproducibility
f_{target}	0.80	Target attention fraction (conditioned only)

Table S4: **BZ2012 ODE model calibration.** Five rate constants were calibrated on Visit 1 real patients (n=23) using bounded Nelder–Mead optimization in log-scale-factor space. The cost function was the mean normalized relative squared error across five TGA features under TF-only conditions. Bounds: $\pm\log(100)$. Convergence: $f_{\text{tol}} = 10^{-6}$, 500 iterations, 11 random restarts.

Rate constant	Role	Default	Calibrated	Scale
prothrombinase k_{cat}	Thrombin burst	63.5 s^{-1}	21.94 s^{-1}	0.35×
intrinsic Xase k_{cat}	Amplification Xa	8.2 s^{-1}	0.169 s^{-1}	0.021×
extrinsic Xase k_{cat}	Initiation Xa	6.0 s^{-1}	93.2 s^{-1}	15.5×
PC activation k_{cat}	Protein C activation	0.41 s^{-1}	0.890 s^{-1}	2.17×
mIIa conversion k	mIIa→IIa	2.3×10^8	1.15×10^9	5.01×

Units for mIIa conversion: $\text{M}^{-1}\text{s}^{-1}$

Table S5: **Summary of mean relative errors across all 72 features.** Per-visit distribution of MREs between real and SA-generated synthetic patient means. The full feature-by-feature table (72 features × 3 visits = 216 entries) is provided in the code repository as `paper_summary_statistics.csv`.

Visit	Median MRE	Mean MRE	25th pctl	75th pctl	Max MRE
V1 (baseline)	0.019	0.027	0.010	0.033	0.157
V2 (1st tri)	0.019	0.027	0.009	0.034	0.124
V3 (3rd tri)	0.016	0.025	0.006	0.040	0.110

Table S6: **Comparison with deep generative baselines.** CTGAN and TVAE were trained on the same 90 per-visit records (23 patients × 3 visits, 72 features) at three epoch counts. SA and MVN results from the main text are shown for reference. CTGAN fails at this sample size (median MRE $\approx$ 19%), consistent with GAN mode collapse on small datasets. TVAE achieves comparable marginal fidelity at high epoch counts but cannot capture cross-visit covariance or perform conditional generation.

Method	Epochs	Med MRE	Med KS	Corr MAE
SA	–	0.019	0.169	0.200
MVN	–	0.024	0.161	0.090
CTGAN	300	0.189	0.364	0.261
CTGAN	1,000	0.195	0.349	0.265
CTGAN	3,000	0.187	0.341	0.180
TVAE	300	0.037	0.174	0.152
TVAE	1,000	0.026	0.248	0.082
TVAE	3,000	0.018	0.224	0.092

Table S7: **PCA variance threshold sensitivity analysis.** SA generation quality is stable across thresholds from 85% to 99%: median MRE remains between 1.2% and 1.8%. The 95% threshold used in the main text (bold) is not a critical design choice.

Threshold	d_{PCA}	K/d	β^*	Novelty	Med MRE	Corr MAE
85%	13	1.77	2.94	0.413	0.012	0.116
90%	15	1.53	2.94	0.462	0.018	0.131
95%	**18**	**1.28**	**2.94**	**0.516**	**0.012**	**0.148**
98%	20	1.15	2.94	0.538	0.013	0.127
99%	21	1.10	2.94	0.540	0.013	0.133

Supplementary Figures

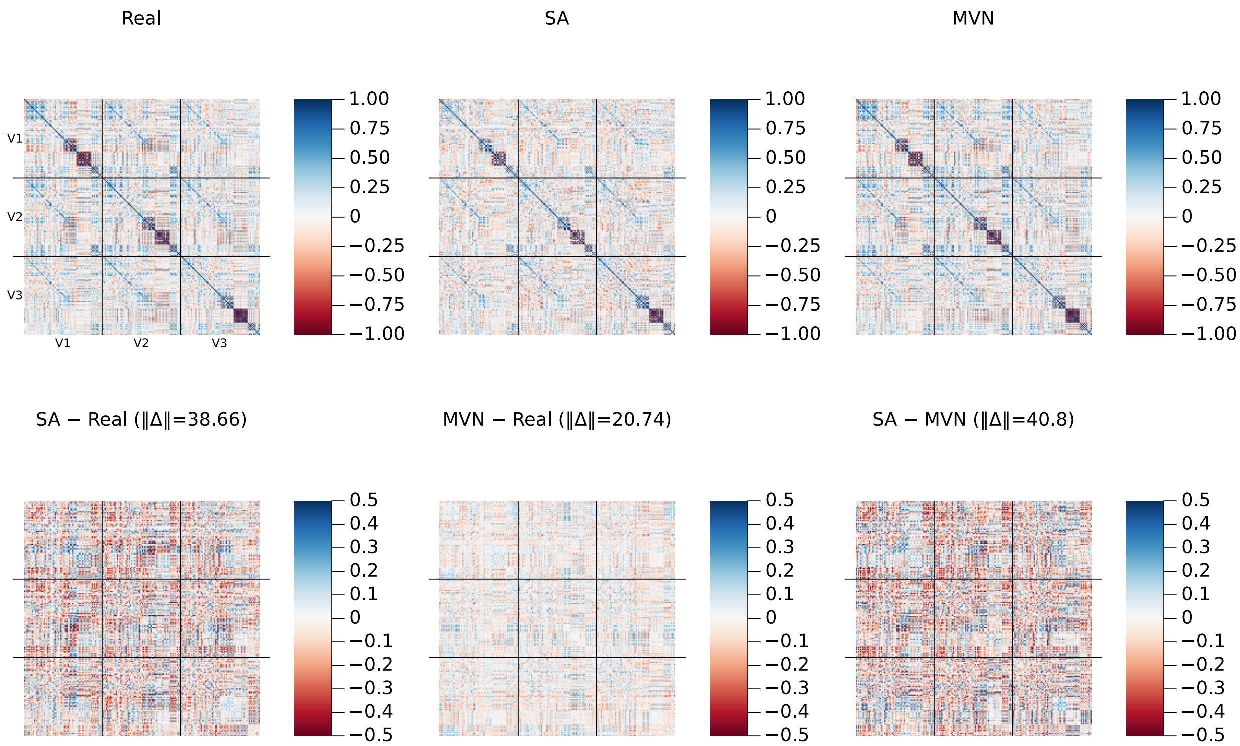

Figure S1: **Cross-visit correlation structure (amplified residual scale).** Same layout as the main-text correlation figure, but with the bottom-row residual panels plotted on a ± 0.5 color scale (vs. ± 1.0 in the main text) to reveal fine structure. Frobenius norms of each residual matrix are shown in the panel titles.

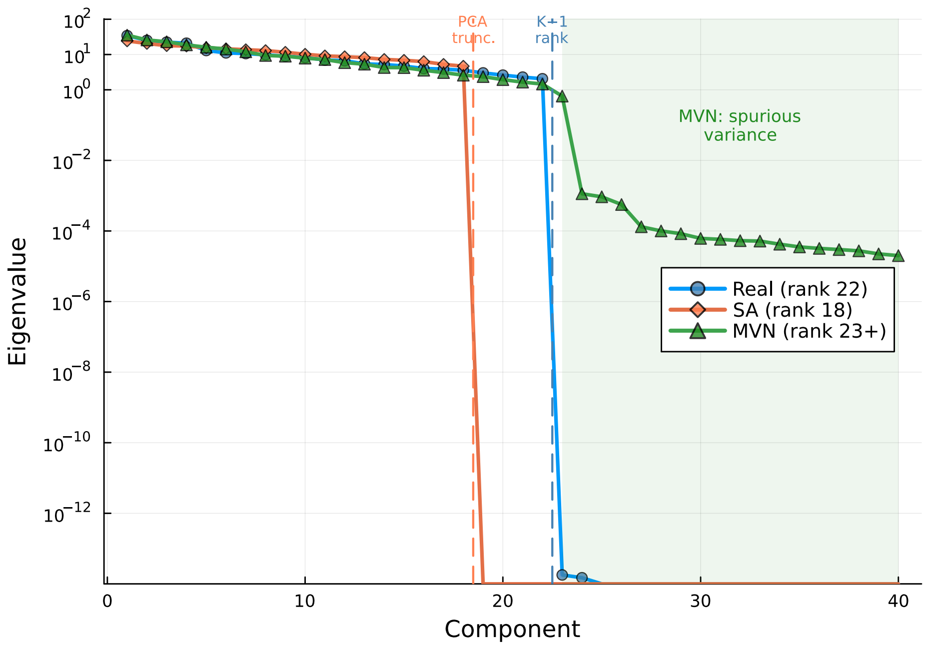

Figure S2: **Covariance eigenvalue spectrum.** Log-scale eigenvalues of the 216×216 sample covariance for Real (K=23, rank 22), SA (N=100, rank 18), and MVN (N=100) populations. Dashed vertical lines mark the SA PCA truncation point (component 18) and the data rank limit (component 22). MVN's Ledoit–Wolf regularization inflates eigenvalues beyond component 22 (green shaded region), introducing spurious variance in 194 dimensions where the data contain no signal.

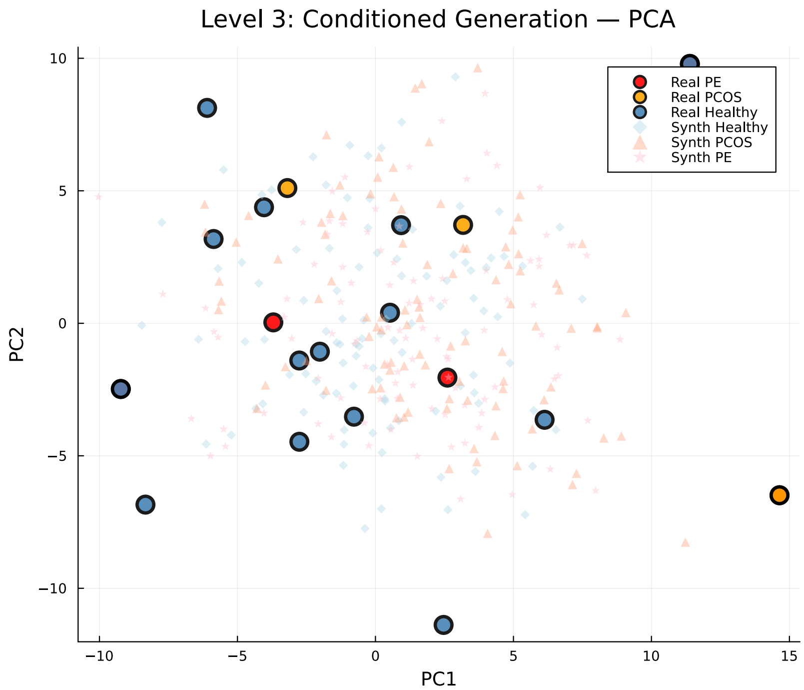

Figure S3: **Conditioned generation: PCA projection (pooled).** Large markers: real patients (blue = Uncomplicated, n=18; orange = PCOS, n=3; red = Developed PE, n=5). Small markers: synthetic patients (N=100 per condition). With only 3 PCOS and 5 Developed PE patients, the subgroups do not form clearly separable clusters in the first two principal components, but the synthetic cohorts concentrate around their respective real counterparts.

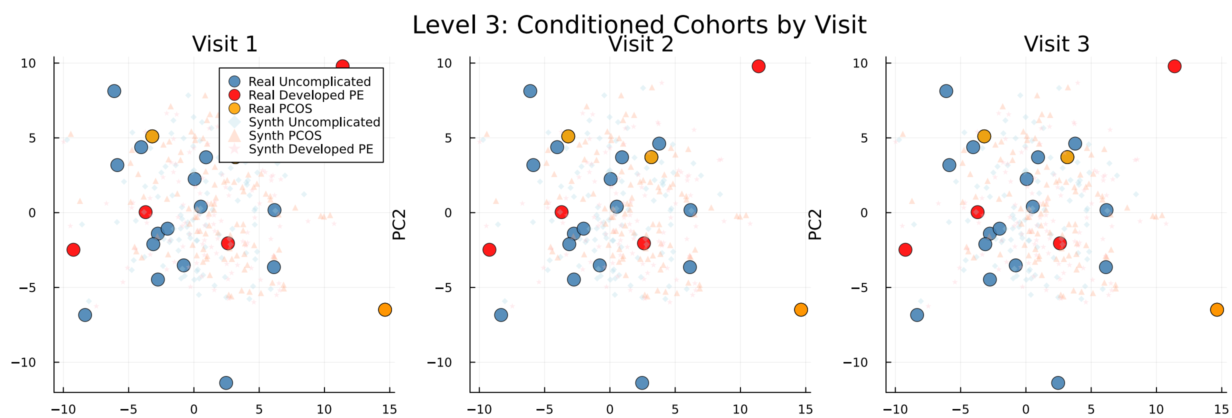

Figure S4: **Conditioned generation: PCA projection by visit.** Same as Supplementary Fig. S3 but separated by visit (V1, V2, V3). Condition-specific clustering is maintained within each individual visit.

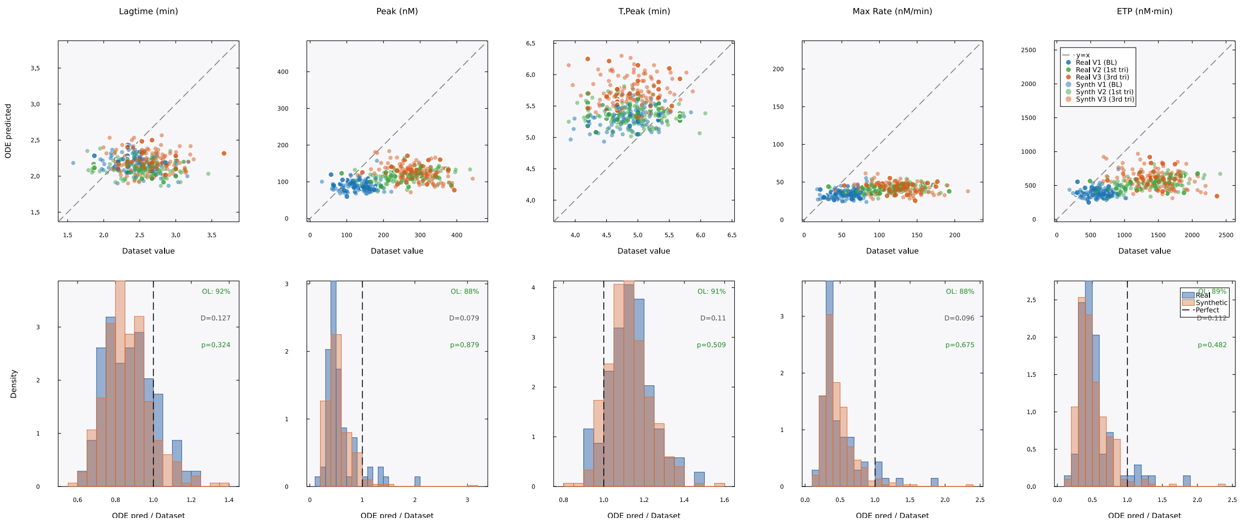

Figure S5: **Mechanistic validation (TF+TM condition).** Same layout as main-text Fig. 6 but under TF+TM conditions. The model systematically underpredicts peak and maximum rate ($\approx 0.5\times$ measured) for both real and synthetic patients, reflecting overcorrection of the protein C anticoagulant pathway. Cloud overlap remains high at 88–92%, confirming that the model processes real and synthetic patients identically even under imperfect calibration conditions.

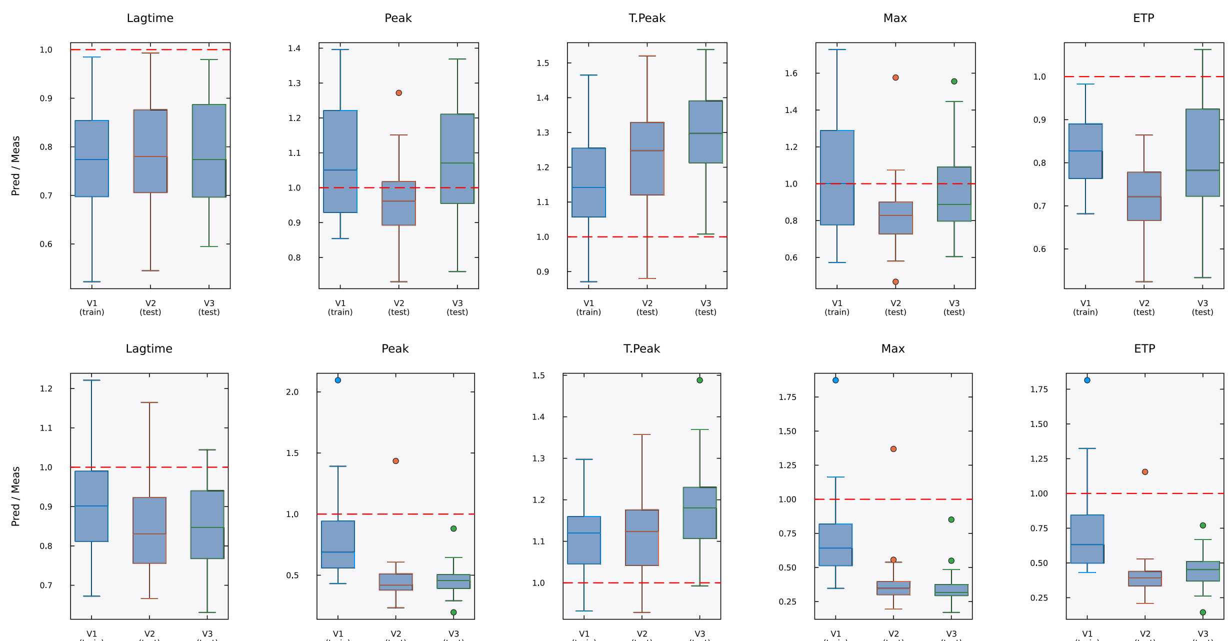

Figure S6: **Cross-visit calibration generalization.** Boxplots of the ODE-predicted/dataset ratio by visit for real patients. The model was calibrated on Visit 1 only. Top: TF-only conditions, where ratios near 1.0 at Visits 2 and 3 indicate that the calibration generalizes across pregnancy timepoints. Bottom: TF+TM conditions, where the calibration does not generalize as well, particularly for peak and maximum rate.

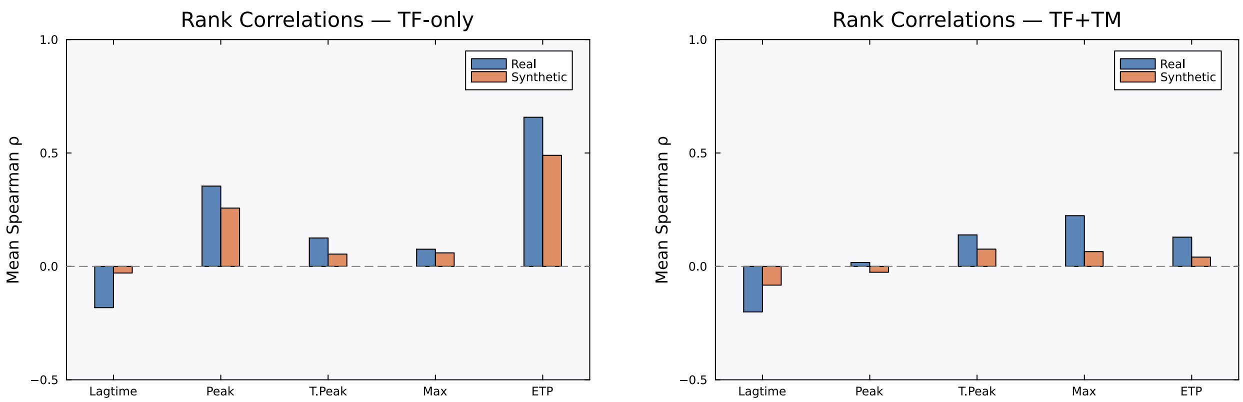

Figure S7: **Spearman rank correlations between ODE-predicted and dataset TGA features.** Left: TF-only. Right: TF+TM. ETP is the best-predicted feature ($\rho \approx 0.6$–0.8 for real patients under TF-only). Weak rank correlations for other features are expected given the 5-parameter population-level calibration.